

# Desenvolvimento de Aplicações com Arquitetura Baseada em Microservices

Prof. Vinicius Cardoso Garcia  
[vcg@cin.ufpe.br](mailto:vcg@cin.ufpe.br) :: [@vinicius3w](https://twitter.com/vinicius3w) :: [assertlab.com](http://assertlab.com)

[IF1007] - Tópicos Avançados em SI 4  
<https://github.com/vinicius3w/if1007-Microservices>

# Licença do material

Este Trabalho foi licenciado com uma Licença

Creative Commons - Atribuição-NãoComercial-Compartilhagual  
3.0 Não Adaptada

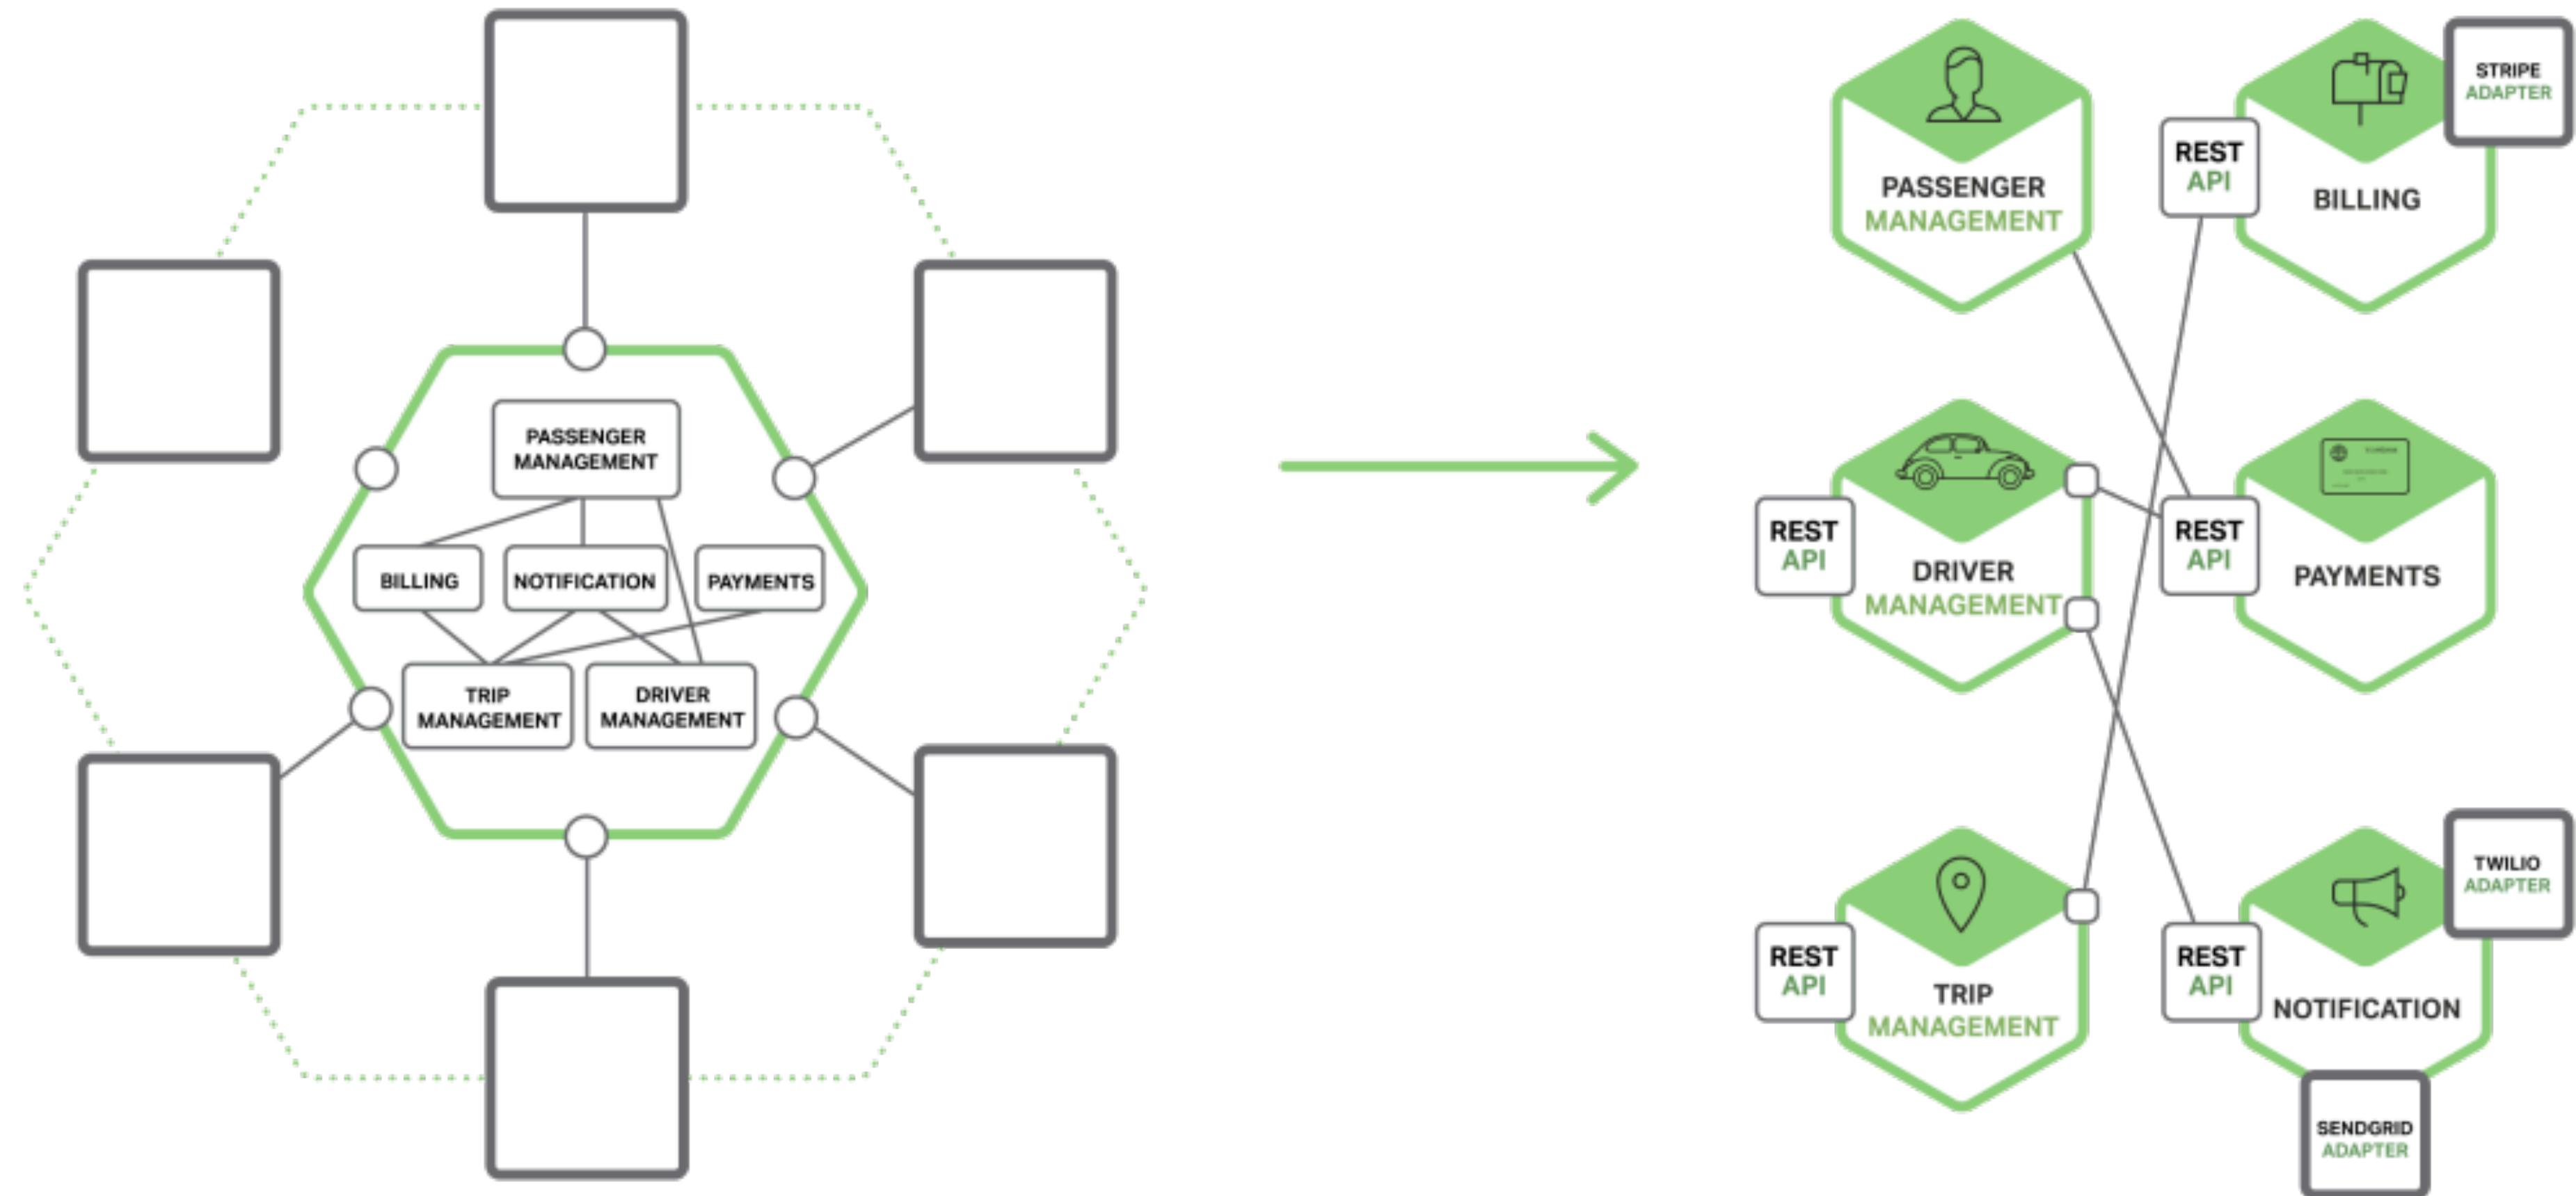


Mais informações visite

<http://creativecommons.org/licenses/by-nc-sa/3.0/deed.pt>

# Resources

- There is no textbook required. However, the following are some books that may be recommended:
  - [Building Microservices: Designing Fine-Grained Systems](#)
  - [Spring Microservices](#)
  - [Spring Boot: Acelere o desenvolvimento de microserviços](#)
  - [Microservices for Java Developers A Hands-on Introduction to Frameworks and Containers](#)
  - [Migrating to Cloud-Native Application Architectures](#)
  - [Continuous Integration](#)
  - [Getting started guides from spring.io](#)



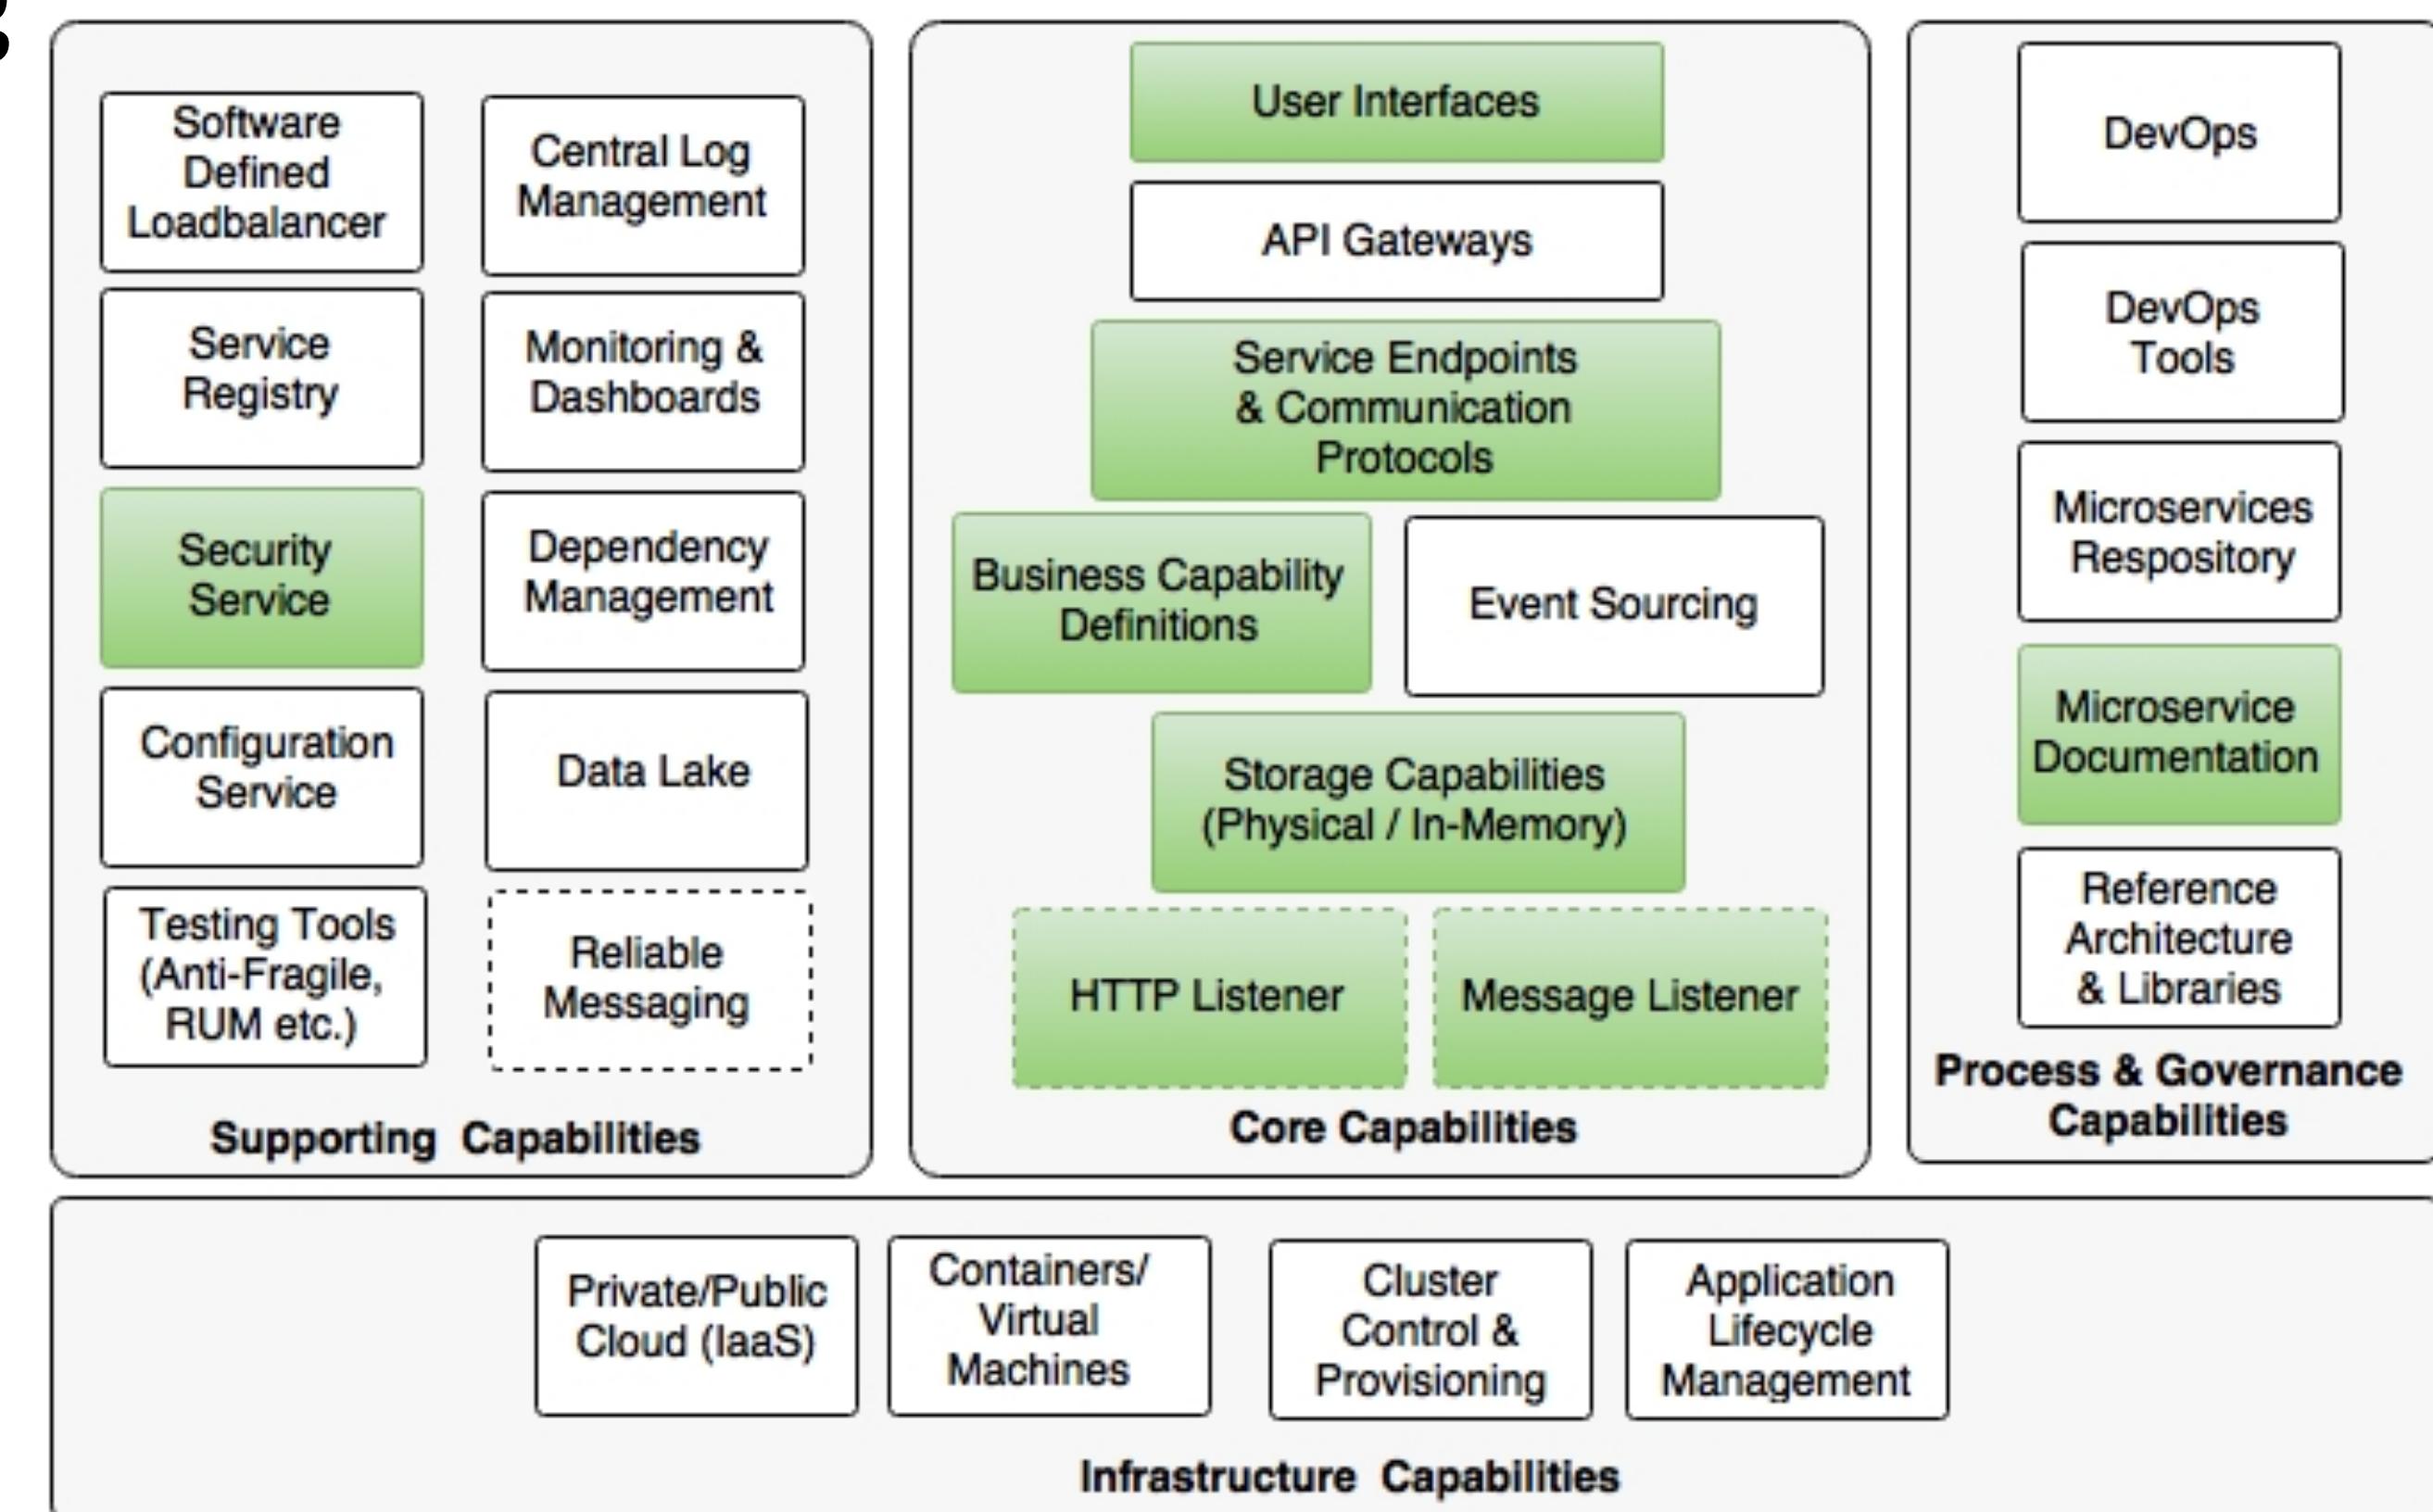
# Microservices Evolution - A Case Study

# Context

- Like SOA, a microservices architecture can be **interpreted** differently by different organizations, based on the **problem** in hand
- BrownField Airline (BF), a fictitious budget airline, and their **journey** from a **monolithic** Passenger Sales and Service (PSS) application to a next generation **microservices architecture**
- The intention of this case study is to get us **as close as possible** to a **live scenario** so that the architecture concepts can be set in stone

# Reviewing the microservices capability model

- The examples in this lecture explore the following microservices capabilities from the microservices capability model discussed in last lecture
  - HTTP Listener
  - Message Listener
  - Storage Capabilities (Physical/In-Memory)
  - Business Capability Definitions
  - Service Endpoints & Communication Protocols
  - User Interfaces
  - Security Service
  - Microservice Documentation

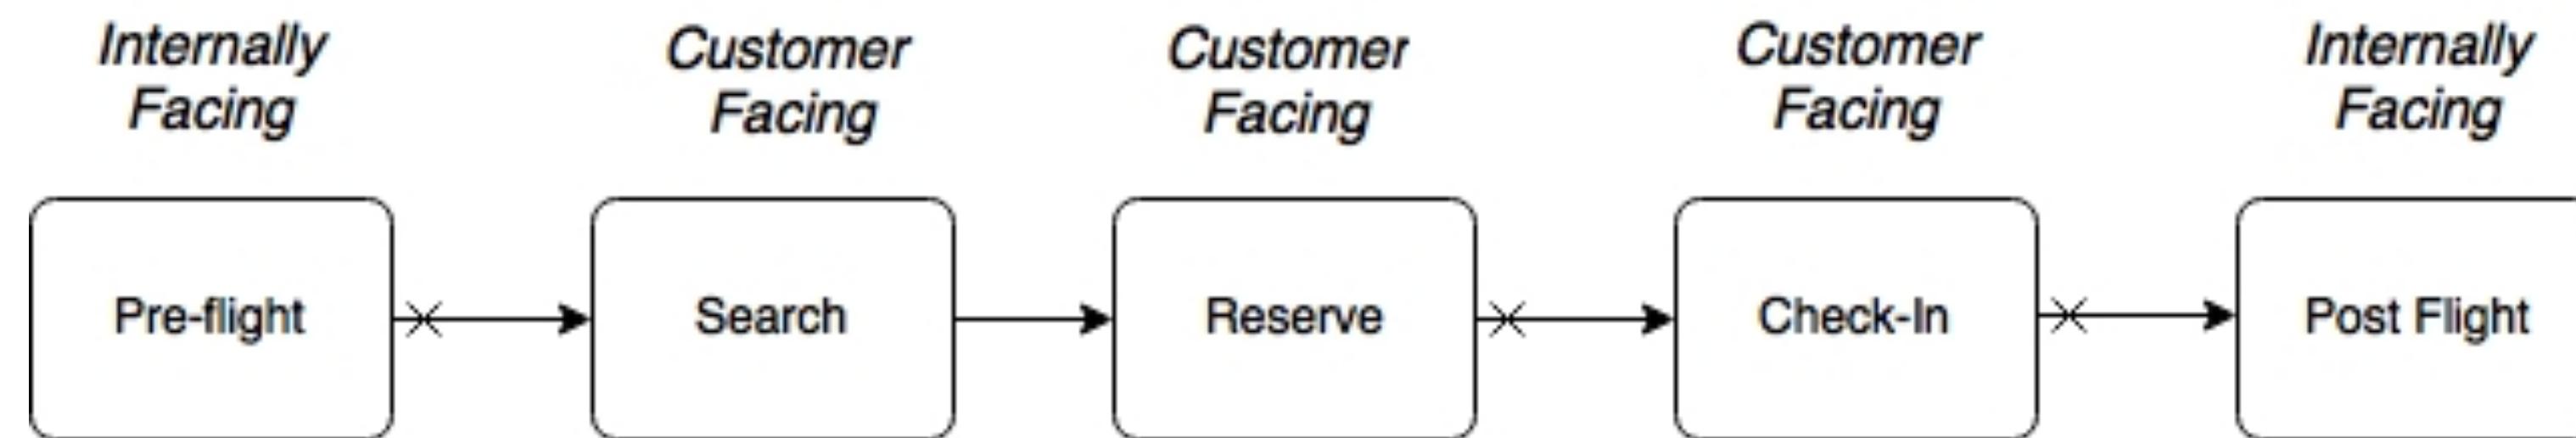


# Understanding the PSS application

- BrownField Airline is one of the fastest growing low-cost, regional airlines, flying directly to more than 100 destinations from its hub
- As a start-up airline, BrownField Airline started its operations with few destinations and few aircrafts
- BrownField developed its home-grown PSS application to handle their passenger sales and services



# Business process view



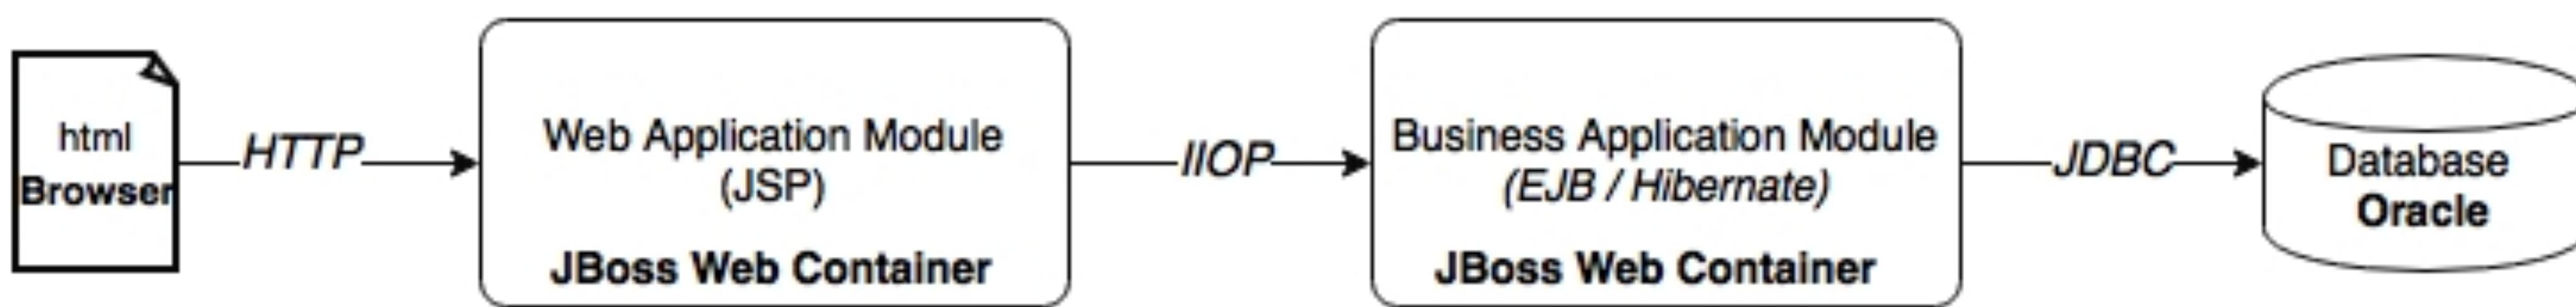
- There are two internally facing functions, **Pre-flight** and **Post-flight**.
  - **Pre-flight** functions include the planning phase, used for preparing flight schedules, plans, aircrafts, and so on
  - **Post-flight** functions are used by the back office for revenue management, accounting, and so on
- The **Search** and **Reserve** functions are part of the **online seat reservation process**, and the **Check-in** function is the process of accepting passengers at the airport
  - The **Check-in** function is also accessible to the end users over the Internet for **online check-in**

# Functional view

<b>Search Functions</b>	<b>Search</b> <i>Flight availability between cities for a given date</i>	<b>Flight</b> <i>Flight routes, aircraft type and schedules</i>	<b>Fare</b> <i>Fares between cities for each flight &amp; date</i>	
<b>Reservation Functions</b>	<b>Book</b> <i>Book passengers on a selected flight &amp; date</i>	<b>Inventory</b> <i>Number of seats available on a flight &amp; date</i>	<b>Payment</b> <i>Payment gateway for online payments</i>	
<b>Check In Functions</b>	<b>Check In</b> <i>Accept a passenger on a flight on the day of travel</i>	<b>Boarding</b> <i>Mark passenger as boarded on the airplane</i>	<b>Seating</b> <i>Allocate passenger a seat based on rules</i>	<b>Baggage</b> <i>Accept passenger baggage and print bag tag</i>
<b>Back Office Functions</b>	<b>CRM</b> <i>Customer relationship management</i>	<b>Data Analysis</b> <i>Business intelligence analysis and reporting</i>	<b>Revenue Management</b> <i>Fare calculations based on forecasts</i>	<b>Loyalty</b> <i>Update passengers loyalty points</i>
<b>Data Management Functions</b>	<b>Reference Data</b> <i>Country, City, Aircrafts, Currency etc.</i>	<b>Customer</b> <i>Manage customers</i>		<b>Accounting</b> <i>Invoicing and Billing</i>
<b>Cross Cutting Functions</b>	<b>User Management</b> <i>Manage user, roles, privileges</i>	<b>Notification</b> <i>Send SMS and e-mails to customers</i>		

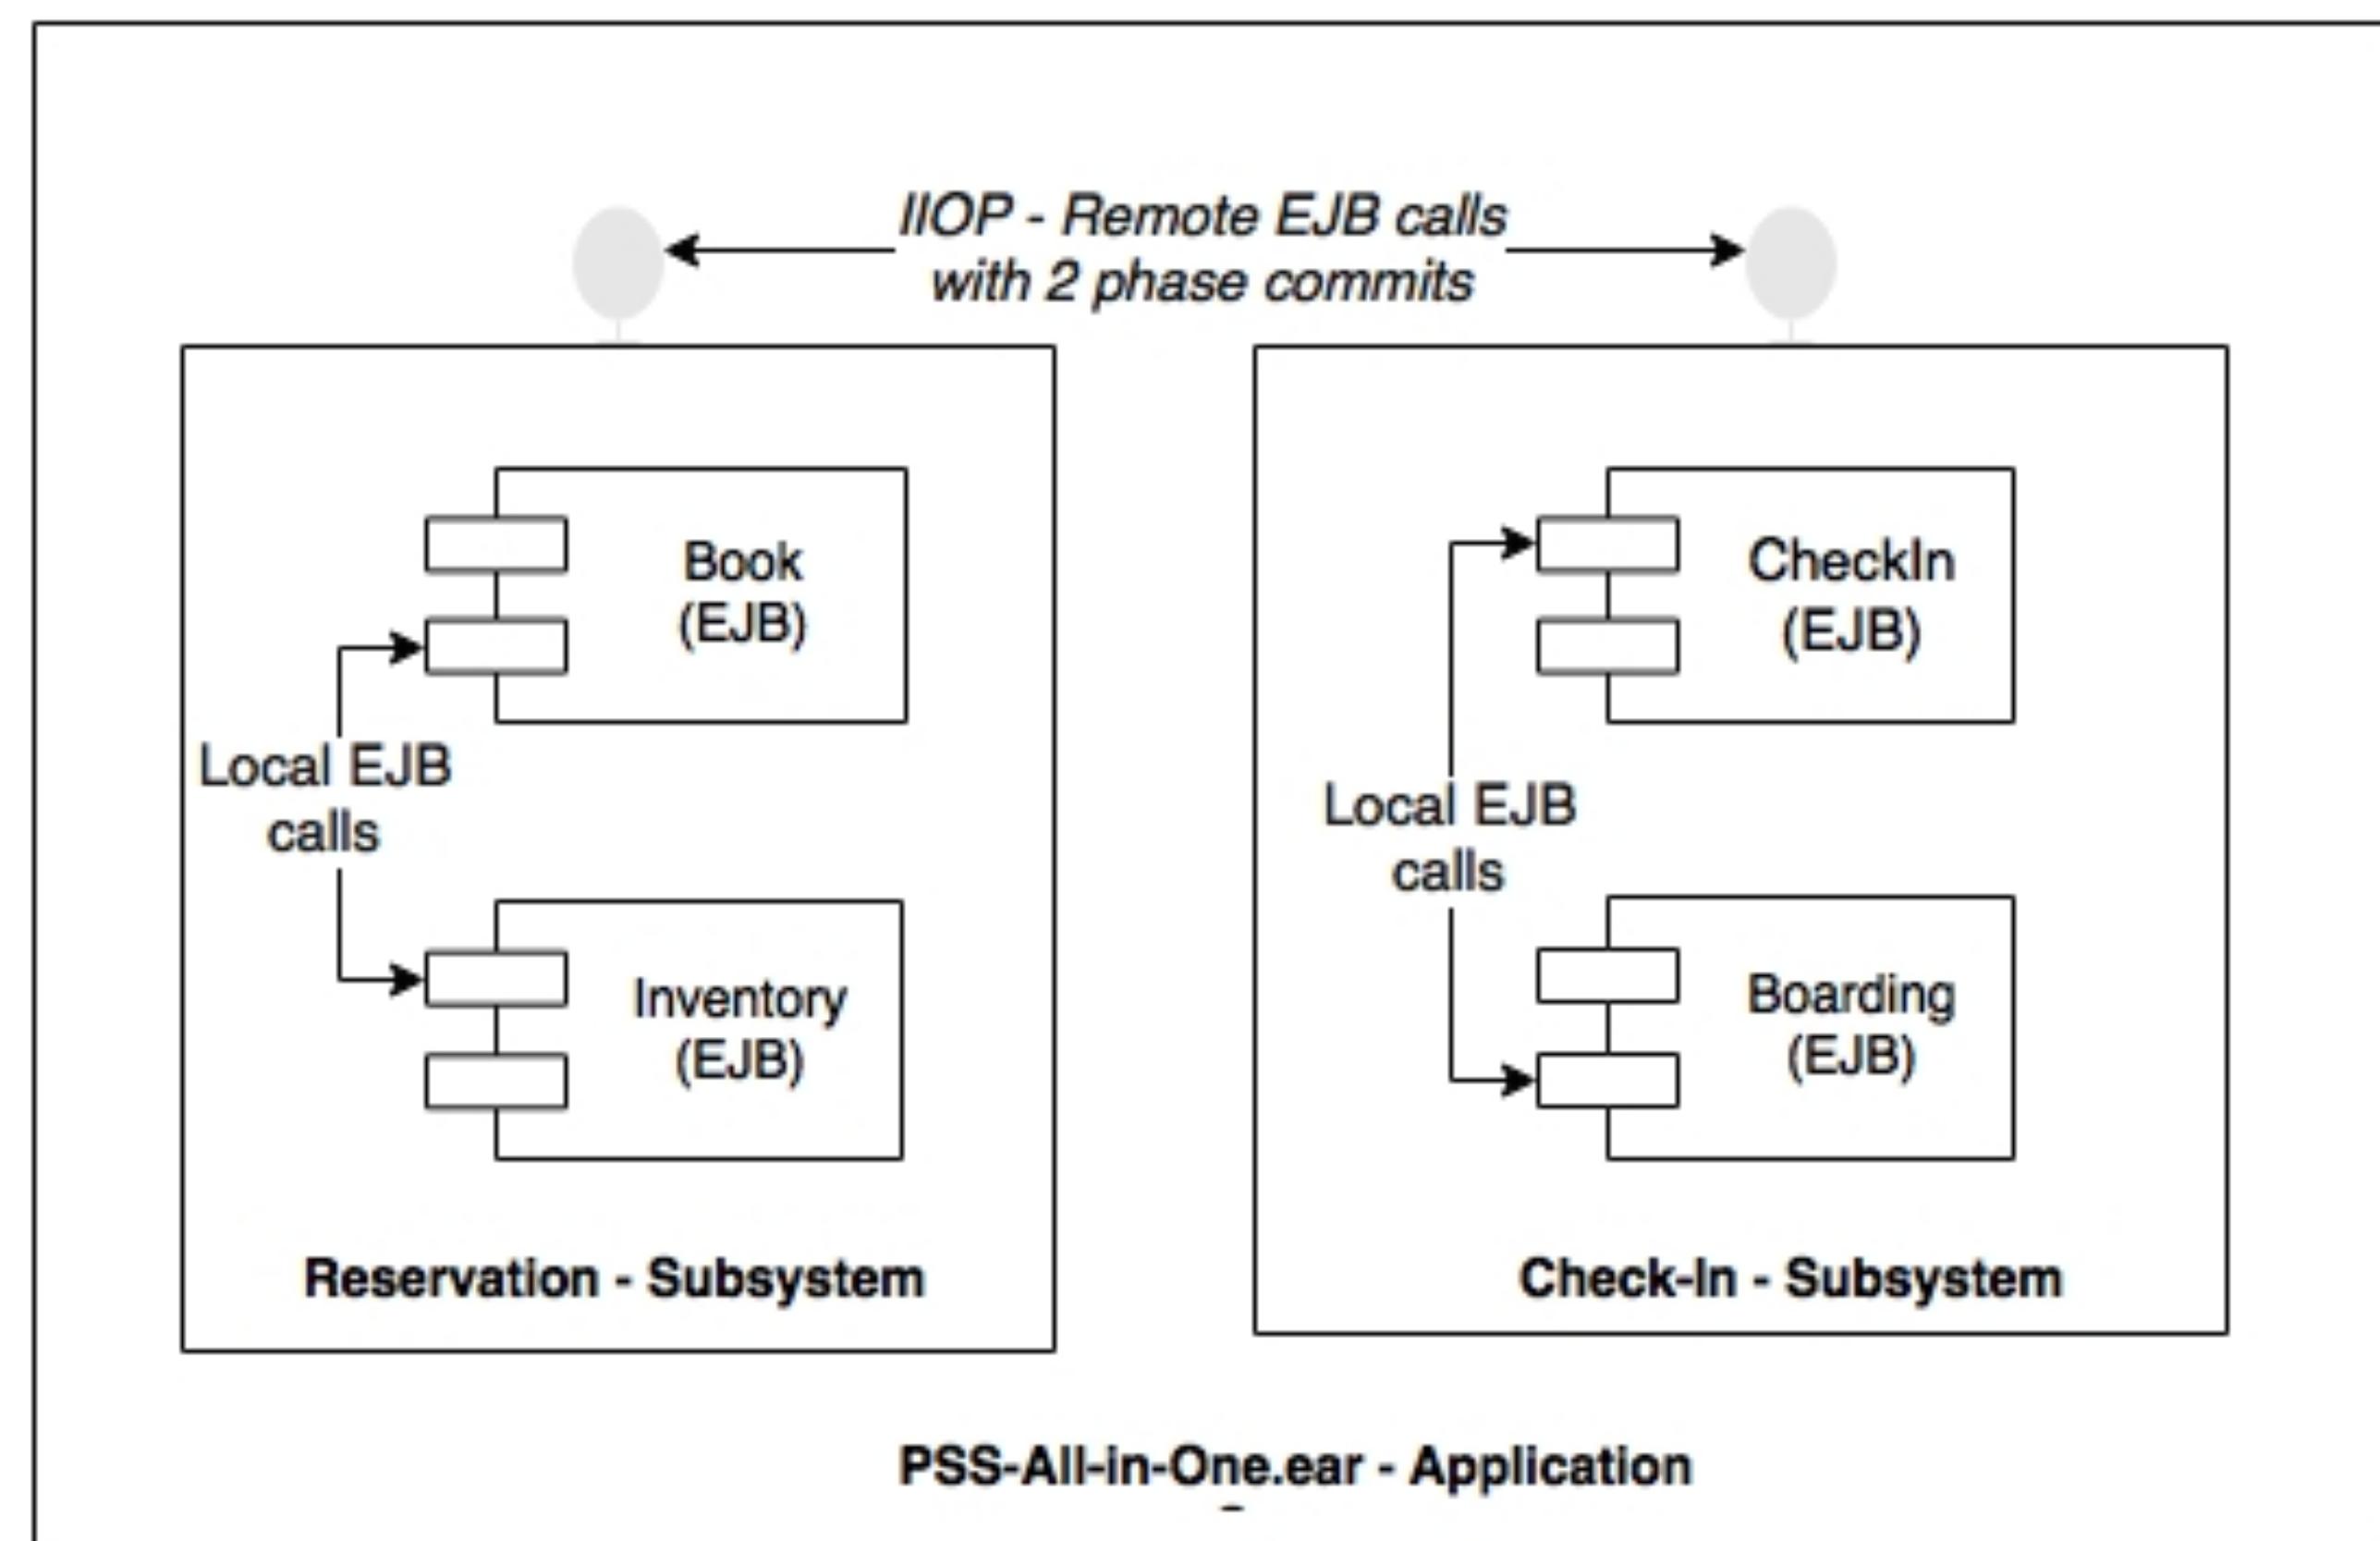
# Architectural view

- This well-architected application was developed using Java and JEE technologies combined with the best-of-the-breed open source technologies available at the time
  - The architecture has well-defined boundaries
  - Different concerns are separated into different layers



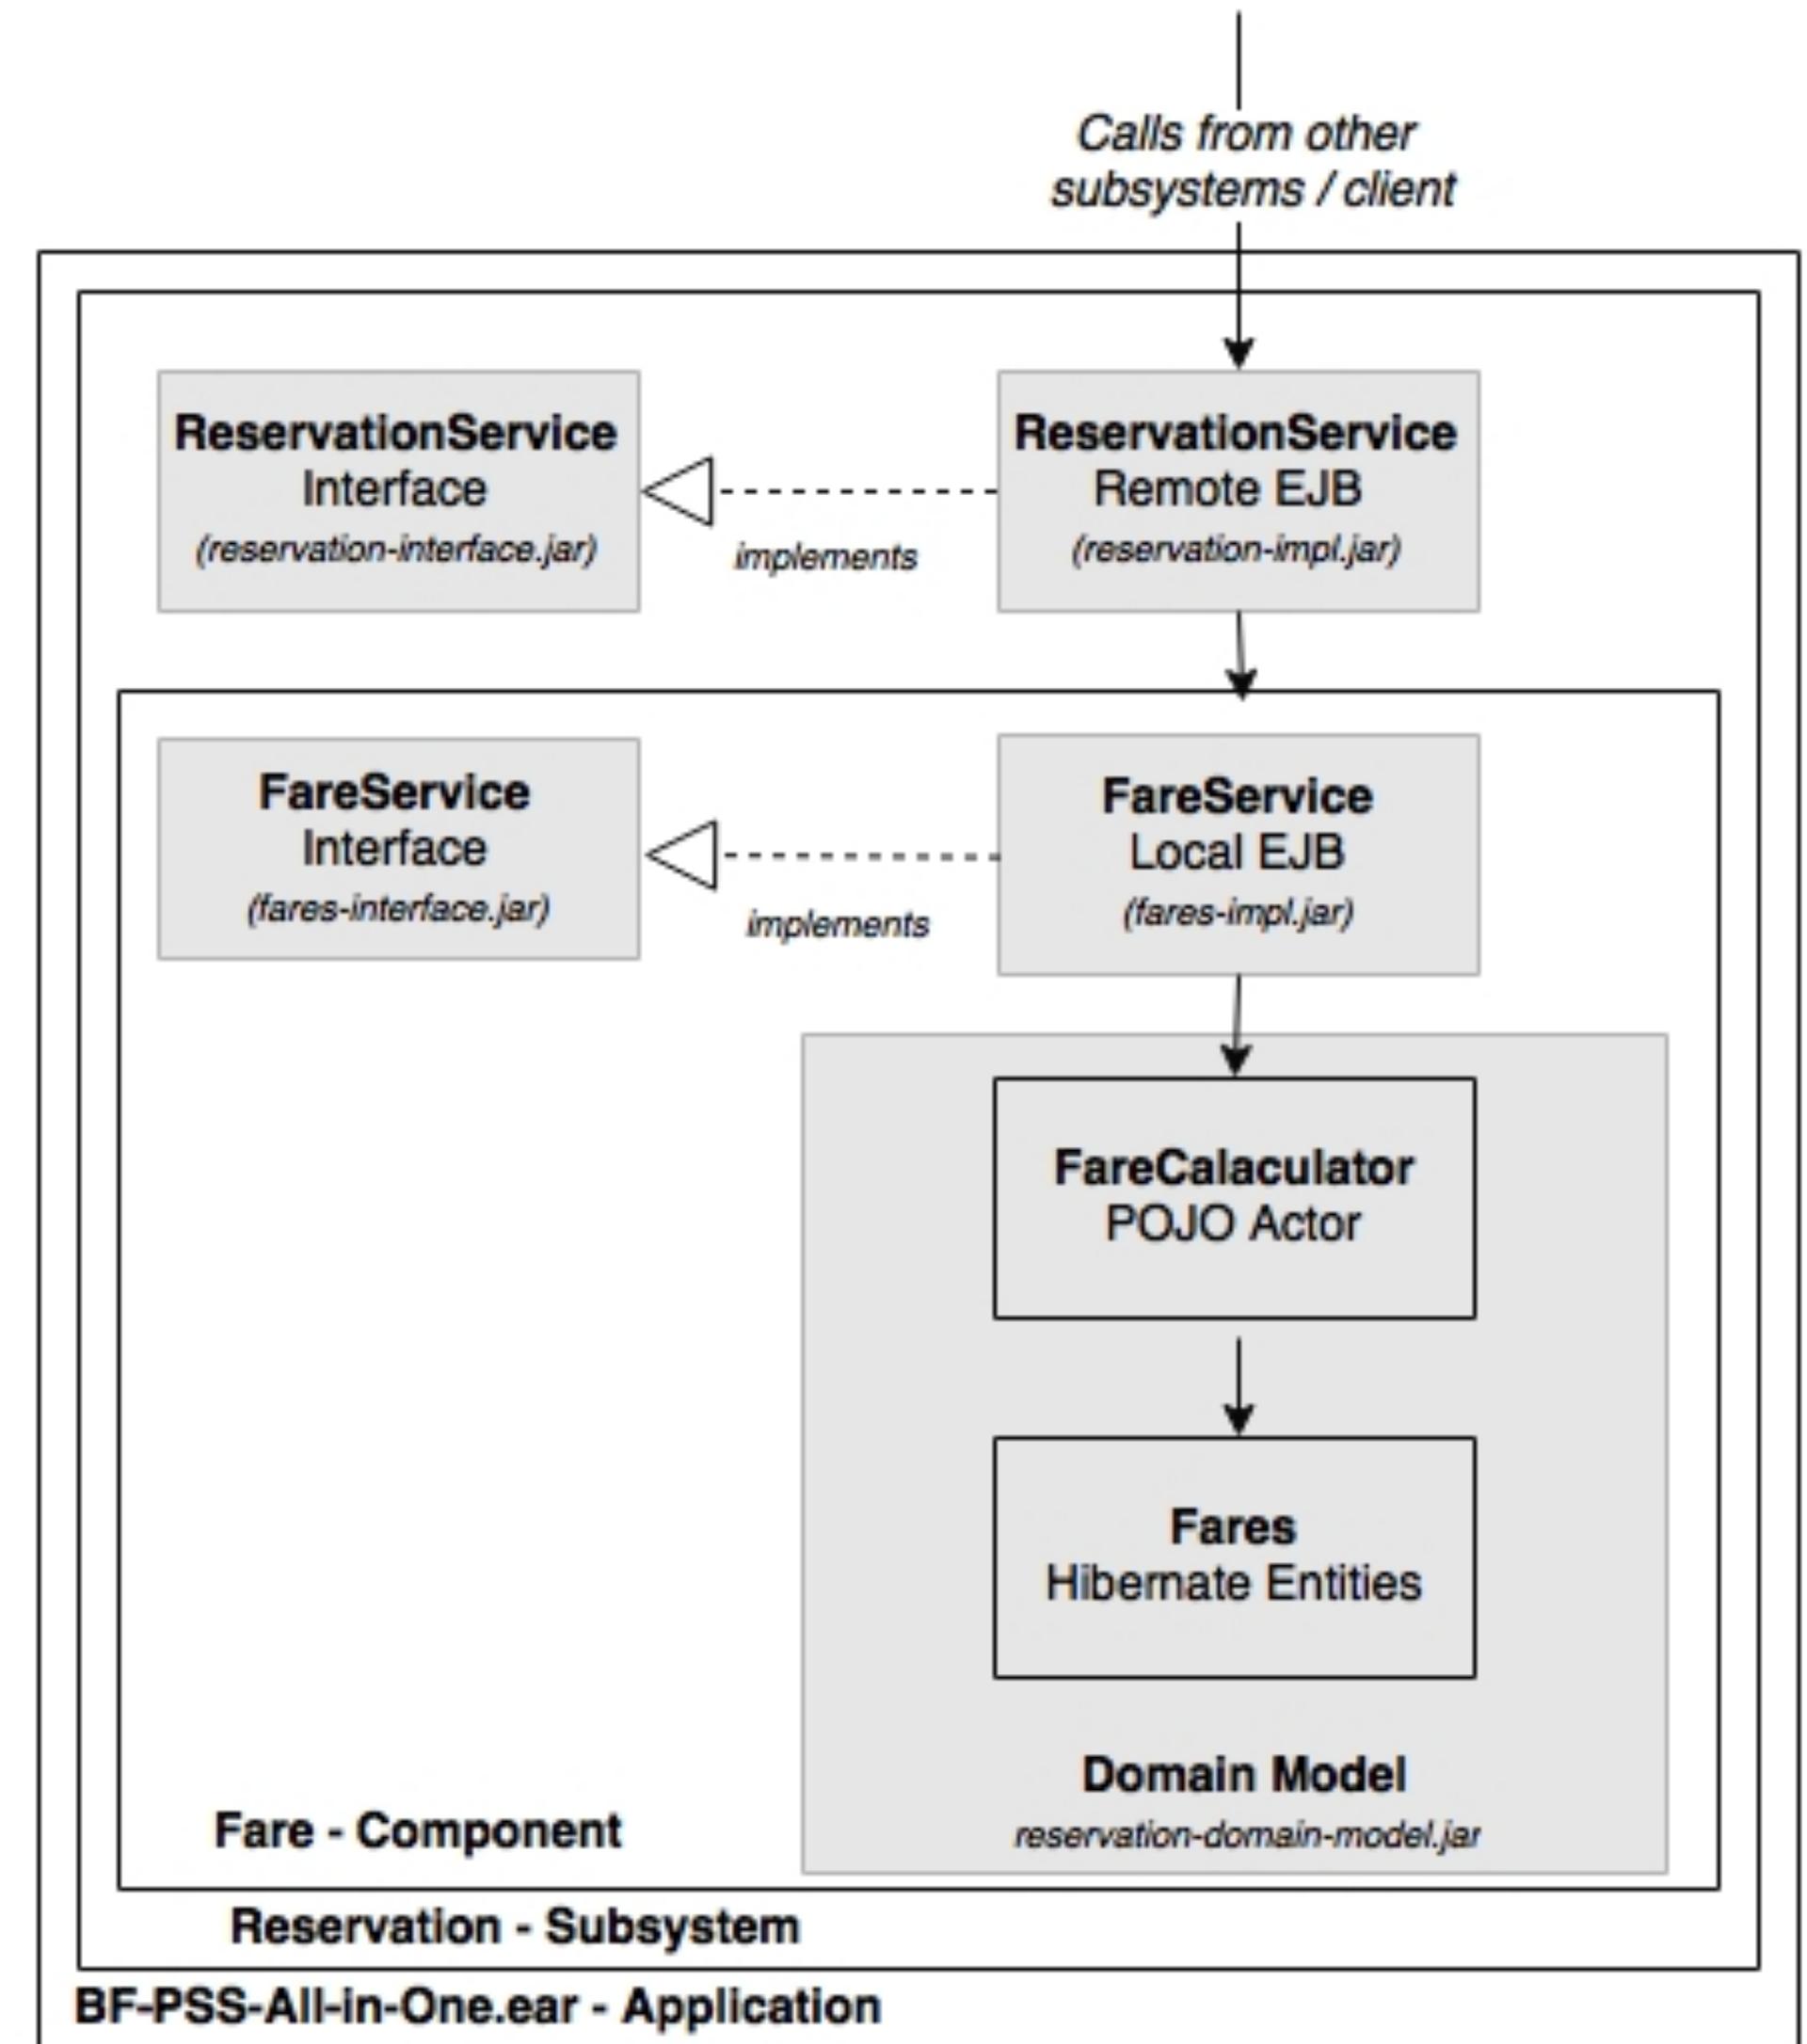
# Design view

- Subsystems interact with each other through remote EJB calls using the IIOP protocol
- The transactional boundaries span across subsystems
- Components within the subsystems communicate with each other through local EJB component interfaces
- In theory, since subsystems use remote EJB endpoints, they could run on different physically separated application servers.
- This was one of the design goals



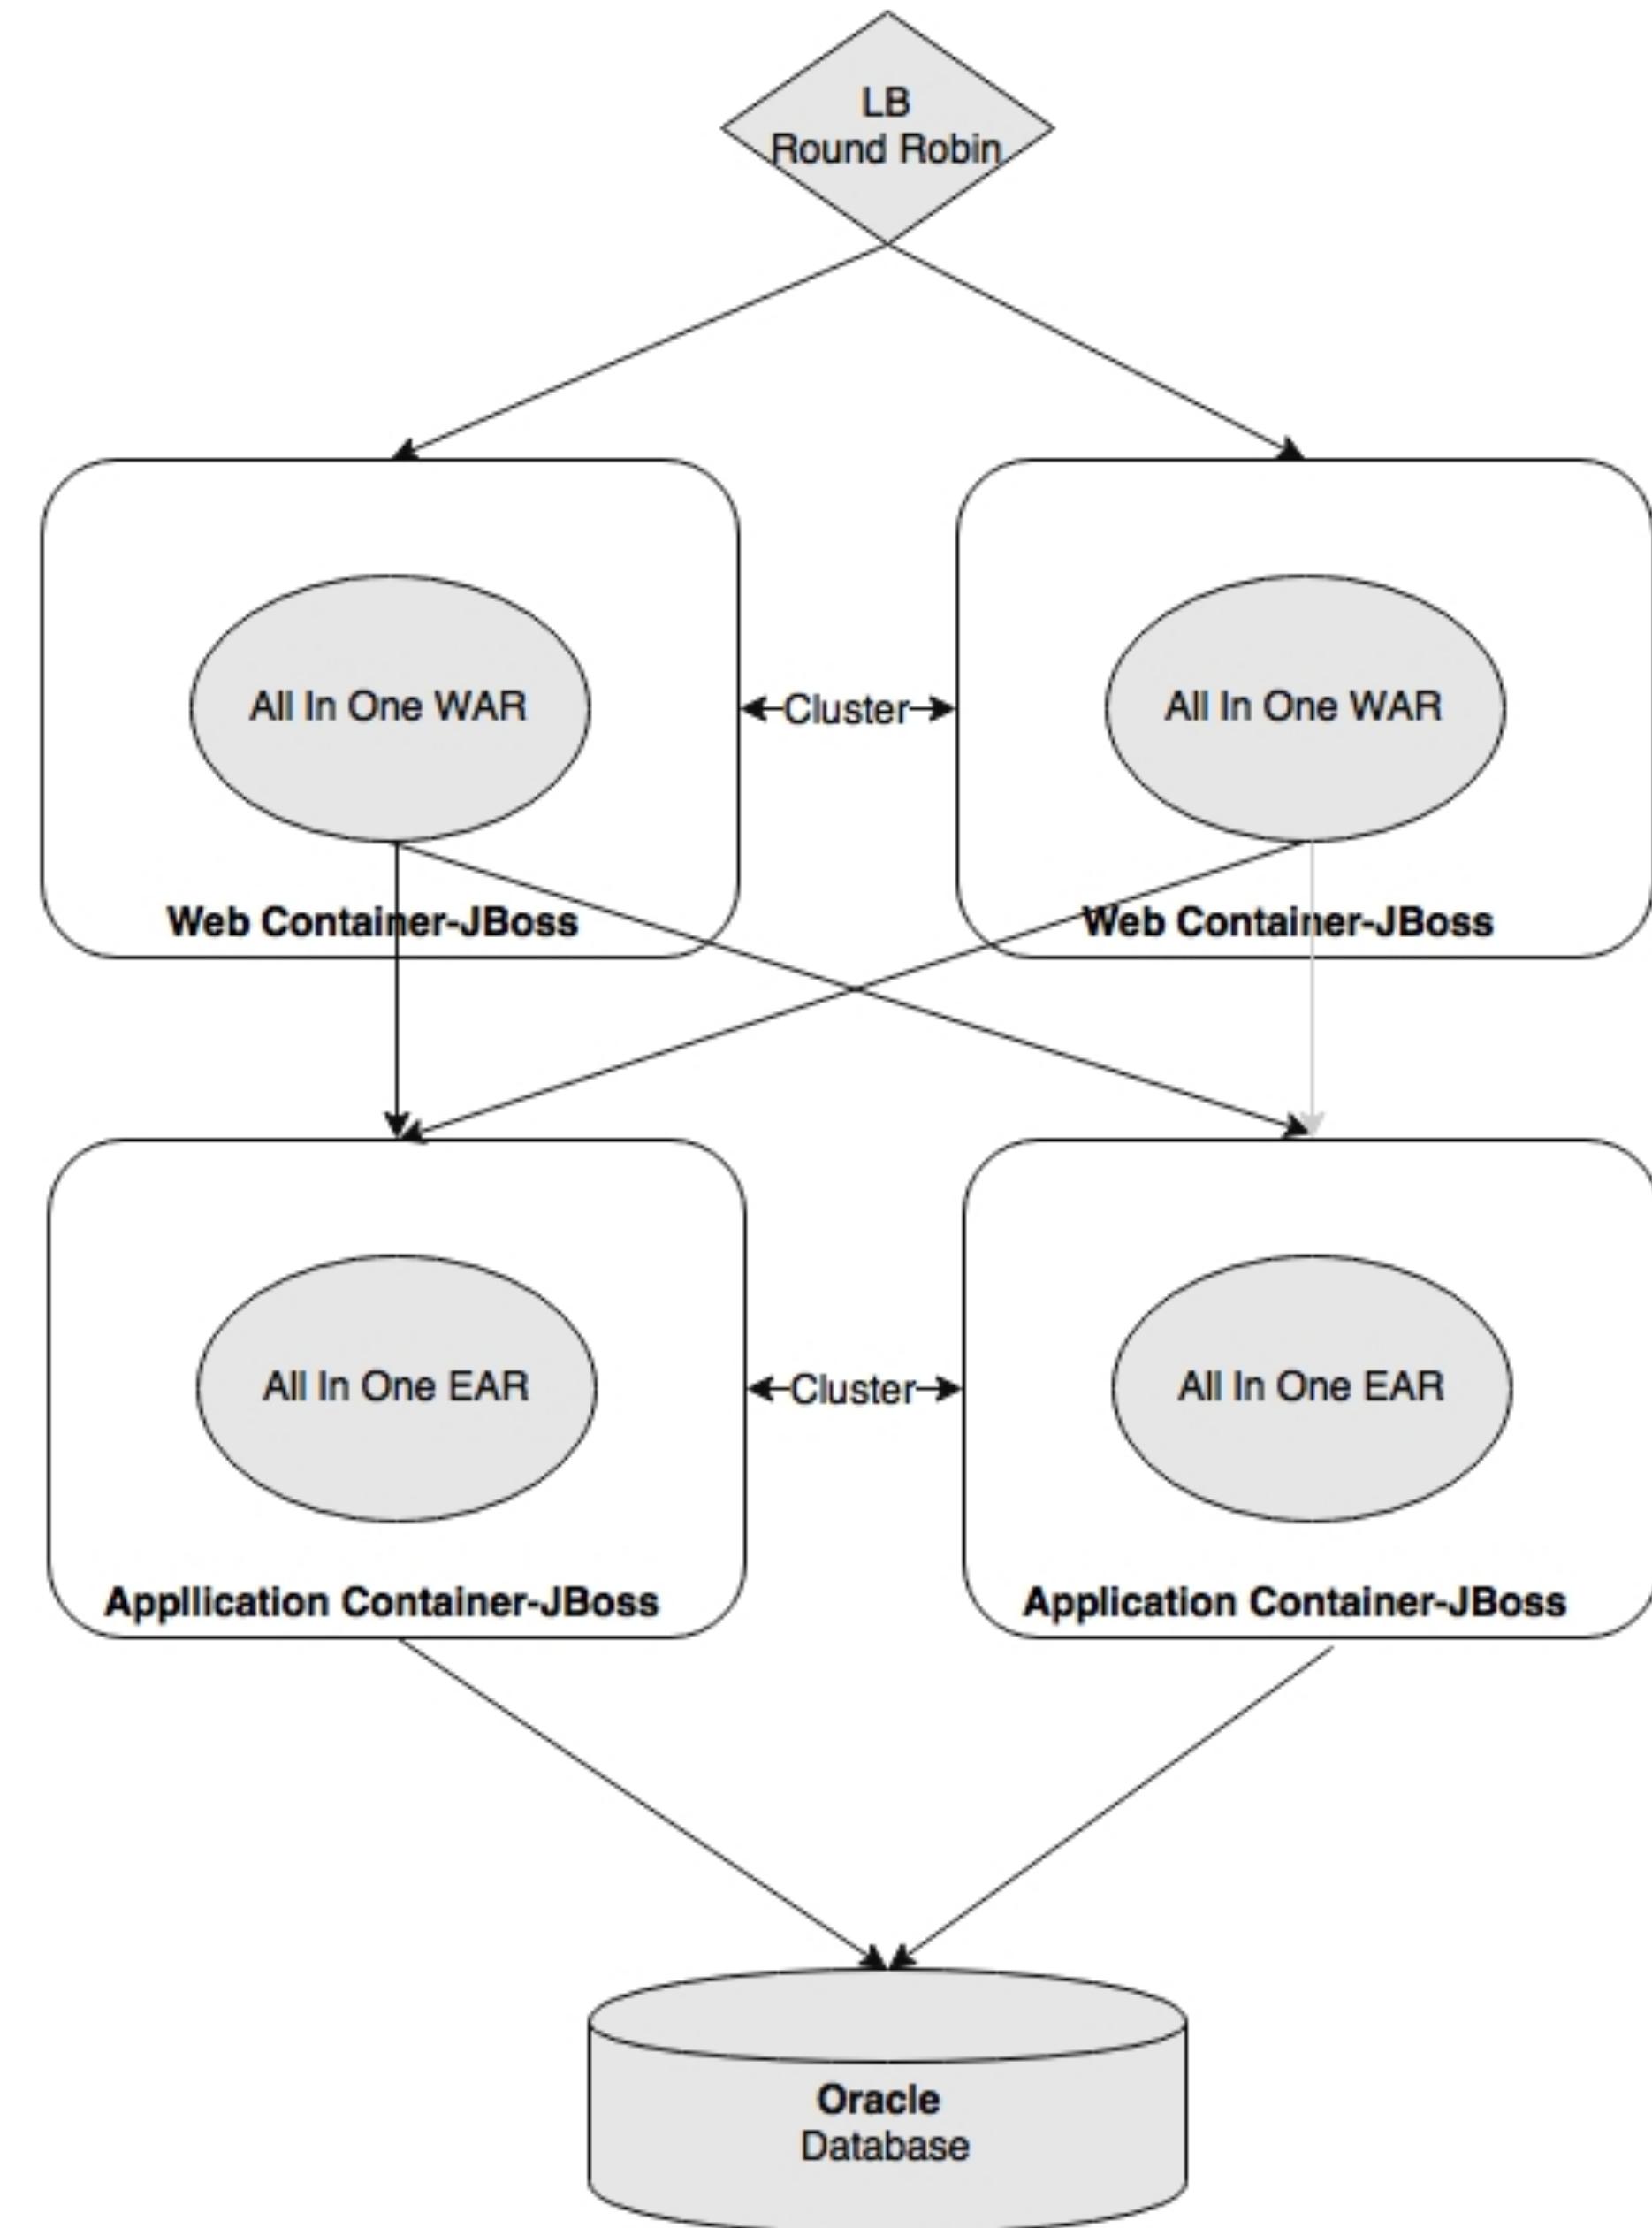
# Implementation view

- The gray-shaded boxes are treated as different Maven projects, and translate into physical artifacts
- Interfaces are packaged as separate JAR files so that clients are abstracted from the implementations
- Local EJBs are used as component interfaces
- Finally, all subsystems are packaged into a single all-in-one EAR, and deployed in the application server



# Deployment view

- The web modules and business modules were deployed into separate application server clusters
- The application was scaled horizontally by adding more and more application servers to the cluster
- Zero downtime deployments were handled by creating a standby cluster, and gracefully diverting the traffic to that cluster
- The standby cluster is destroyed once the primary cluster is patched with the new version and brought back to service



# Death of the monolith

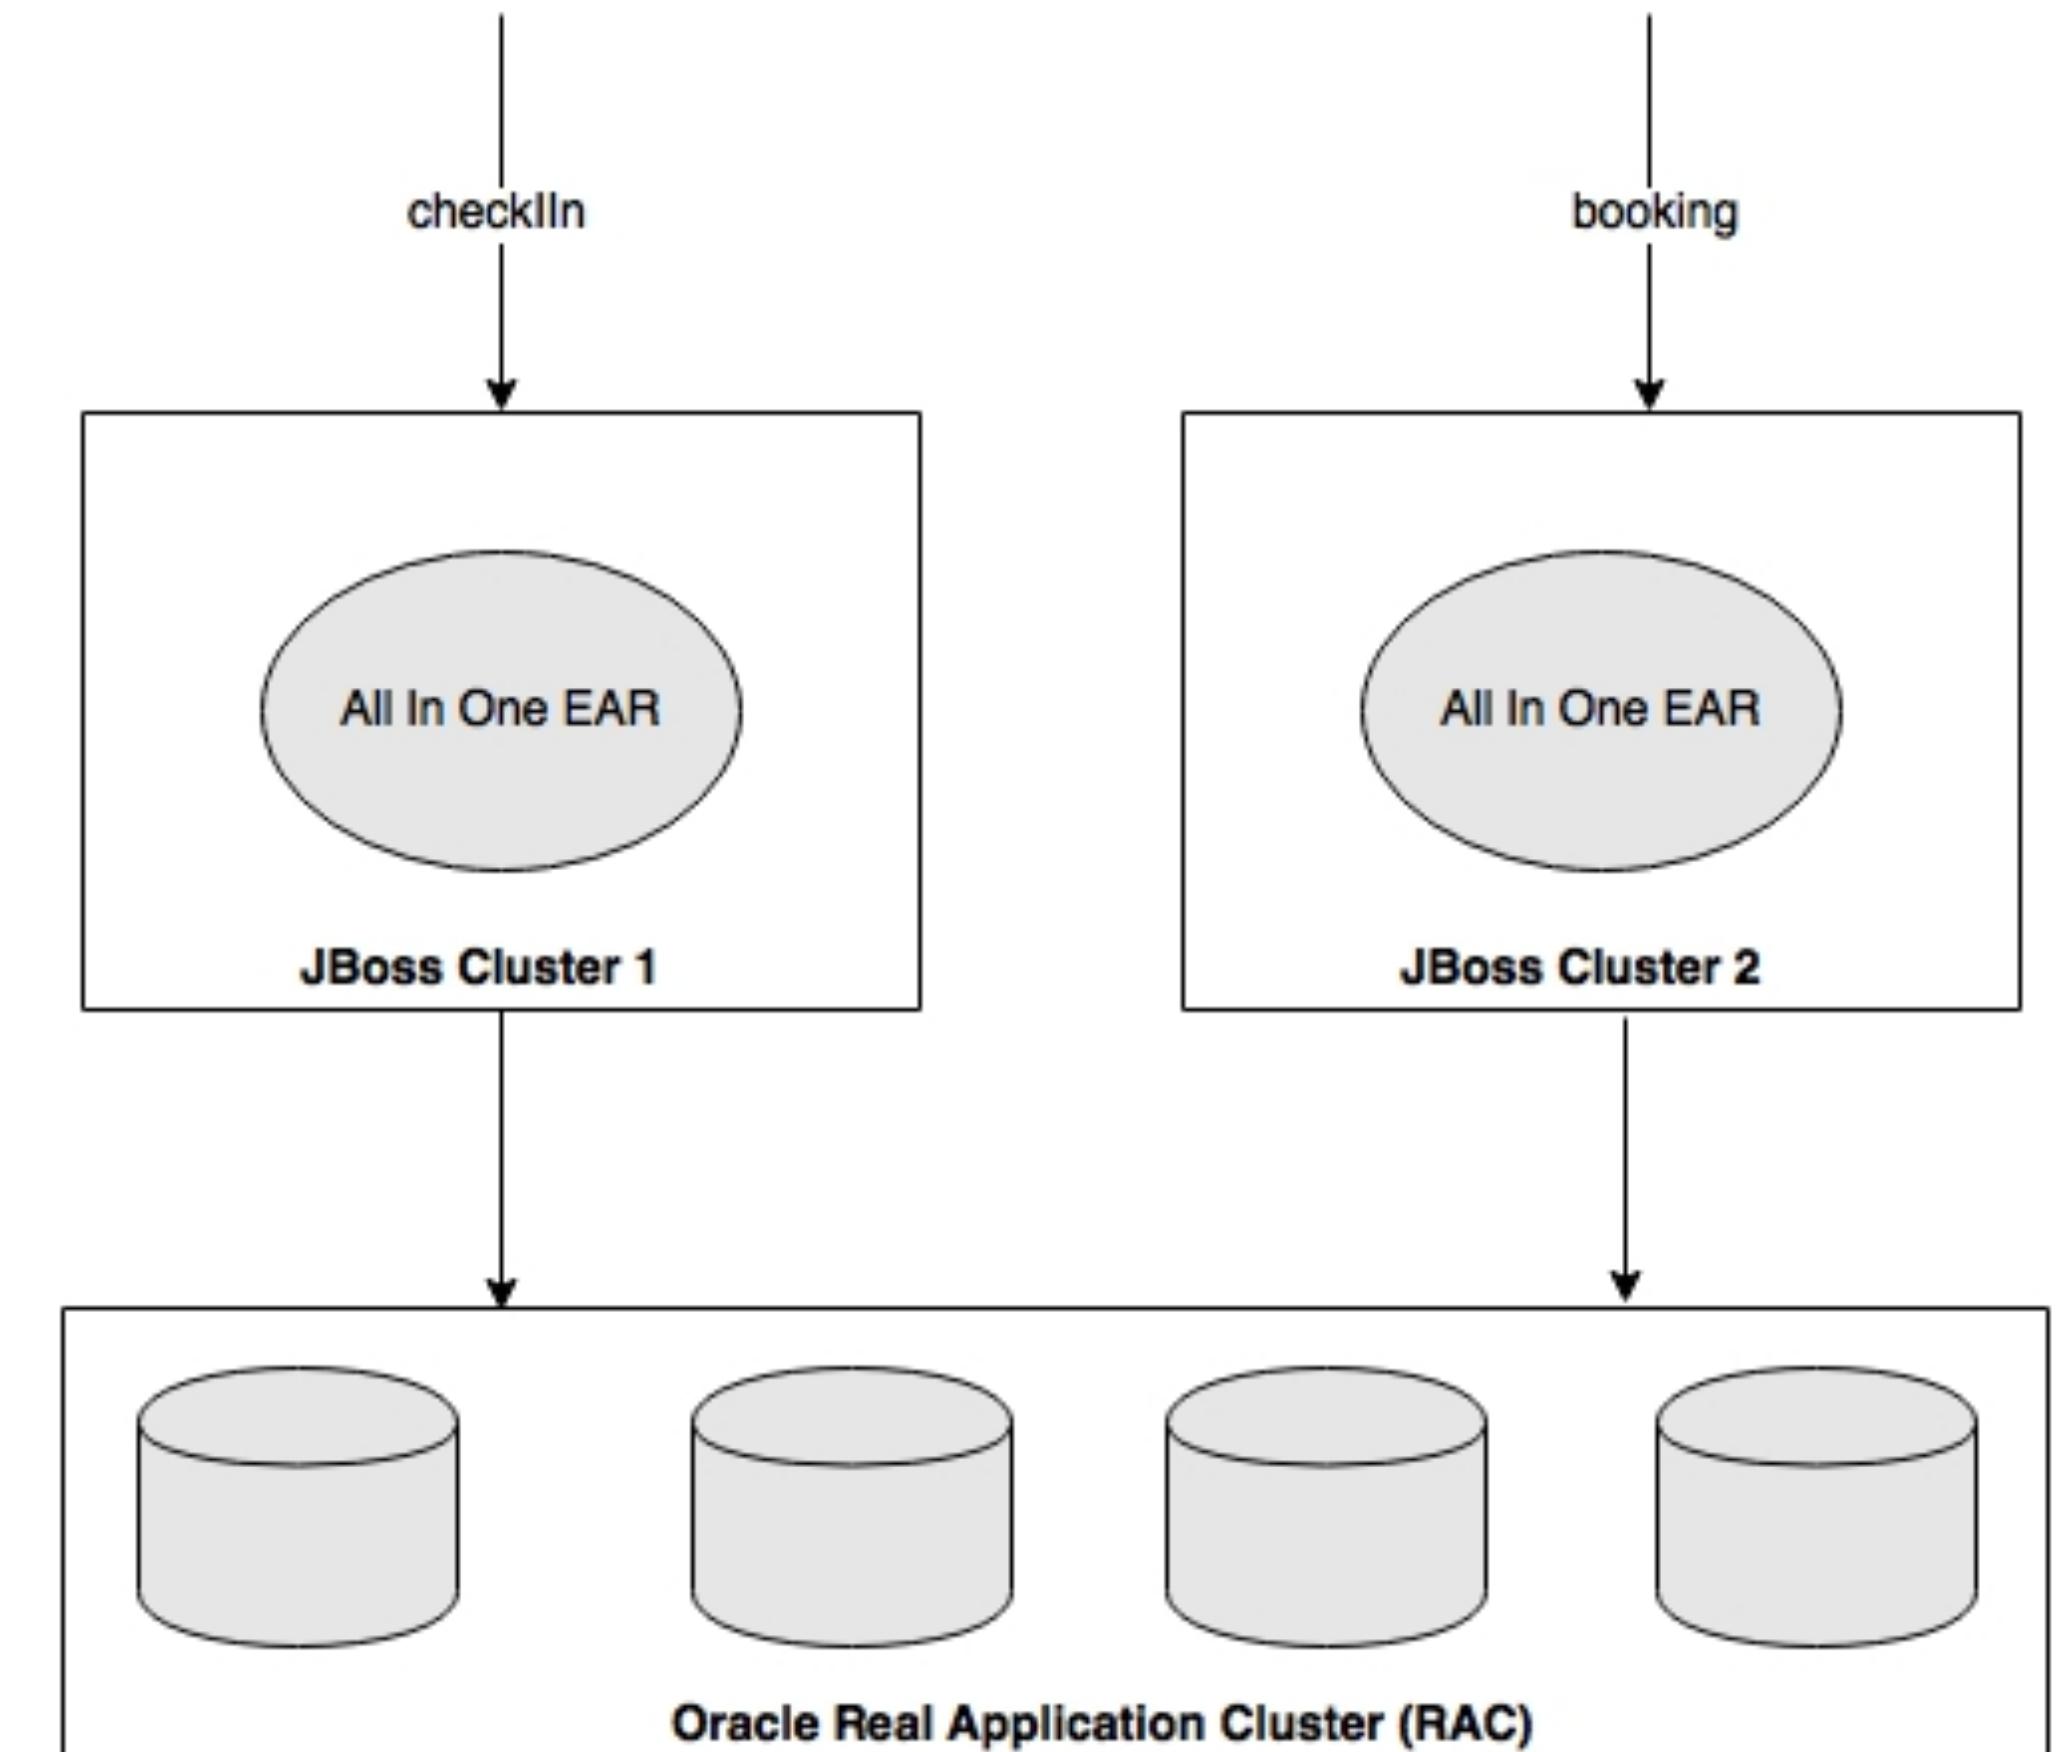
- The PSS application was performing **well**, **successfully** supporting **all business requirements** as well as the **expected service levels**
  - The system had **no issues** in scaling with the organic growth of the business in the initial years
- The business has seen tremendous **growth** over a period of time.
  - The **fleet size increased significantly**, and new destinations got added to the network.
  - As a result of this rapid growth, the number of bookings has gone up, resulting in a steep increase in transaction volumes, up to **200 - to 500 - fold** of what was originally estimated

# Pain points

- The **rapid growth** of the business eventually put the application under **pressure**
- Weaknesses of the system as well as the root causes of many failures
  - Stability - primarily due to stuck threads, which limit the application server's capability to accept more transactions
  - Outages - outage window increased largely because of the increase in server startup time
  - Agility - changes became harder to implement

# Stop gap fix

- Performance issues were partially addressed by applying the Y-axis scale method in the scale cube
- This new scaling model reduced the stability issues, but at a premium of increased complexity and cost of ownership
- The technology debt also increased over a period of time, leading to a state where a complete rewrite was the only option for reducing this technology debt

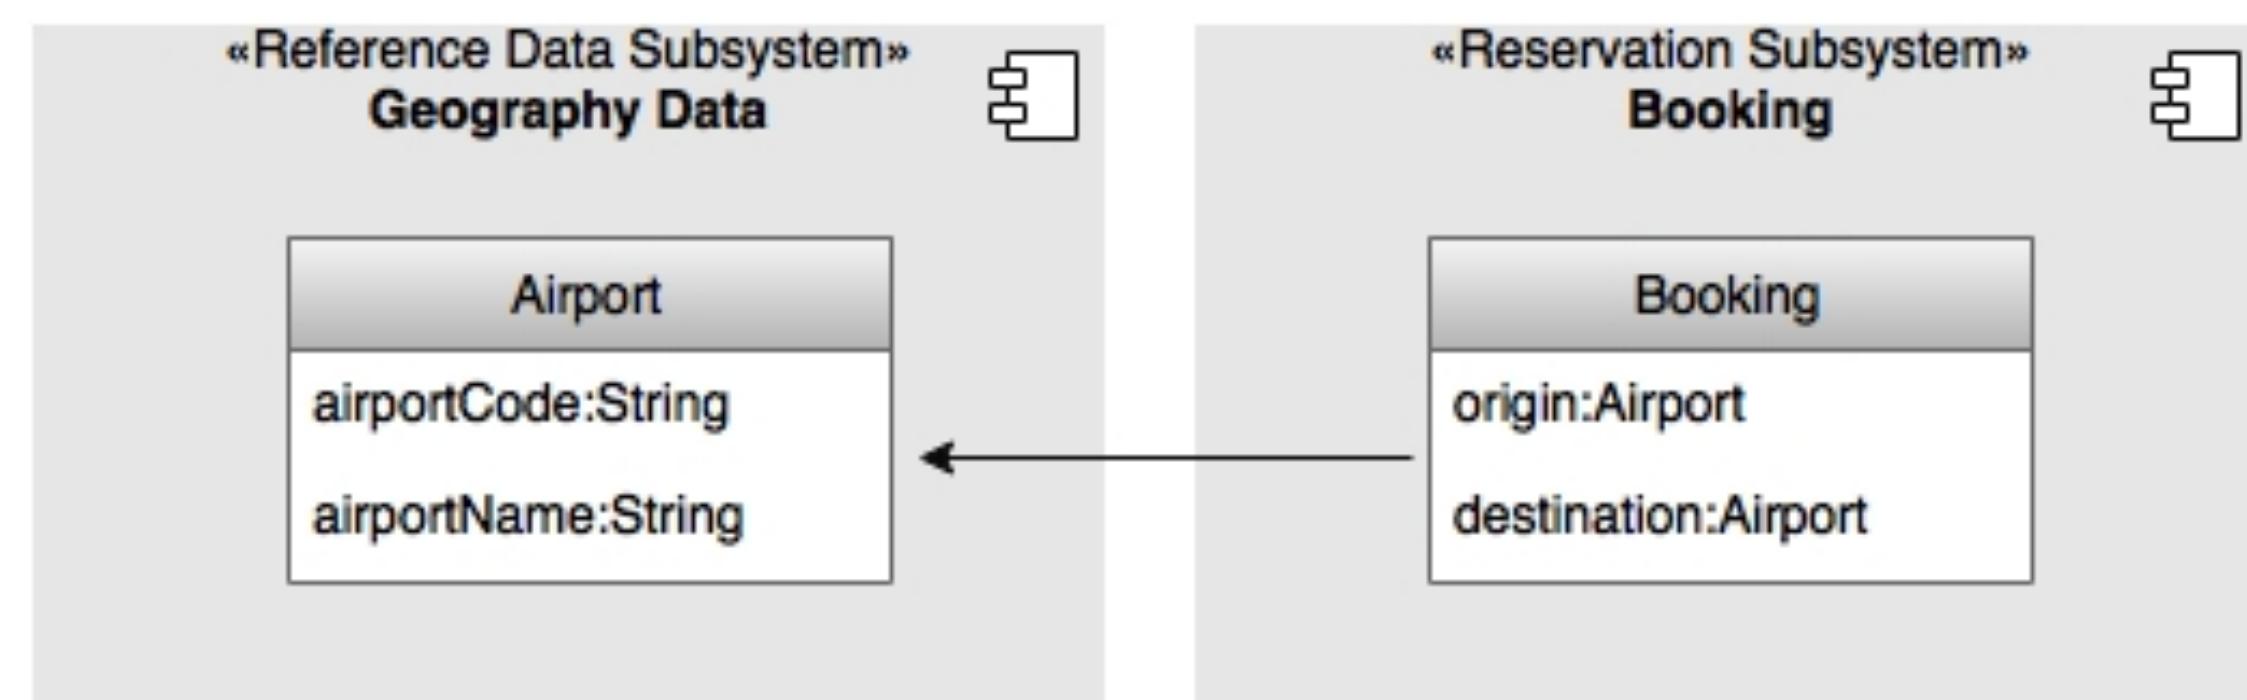


# Retrospection

- Although the application was well-architected, there was a clear segregation between the functional components
- They were loosely coupled, programmed to interfaces, with access through standards-based interfaces, and had a rich domain model
- How come such a well-architected application failed to live up to the expectations? What else could the architects have done?

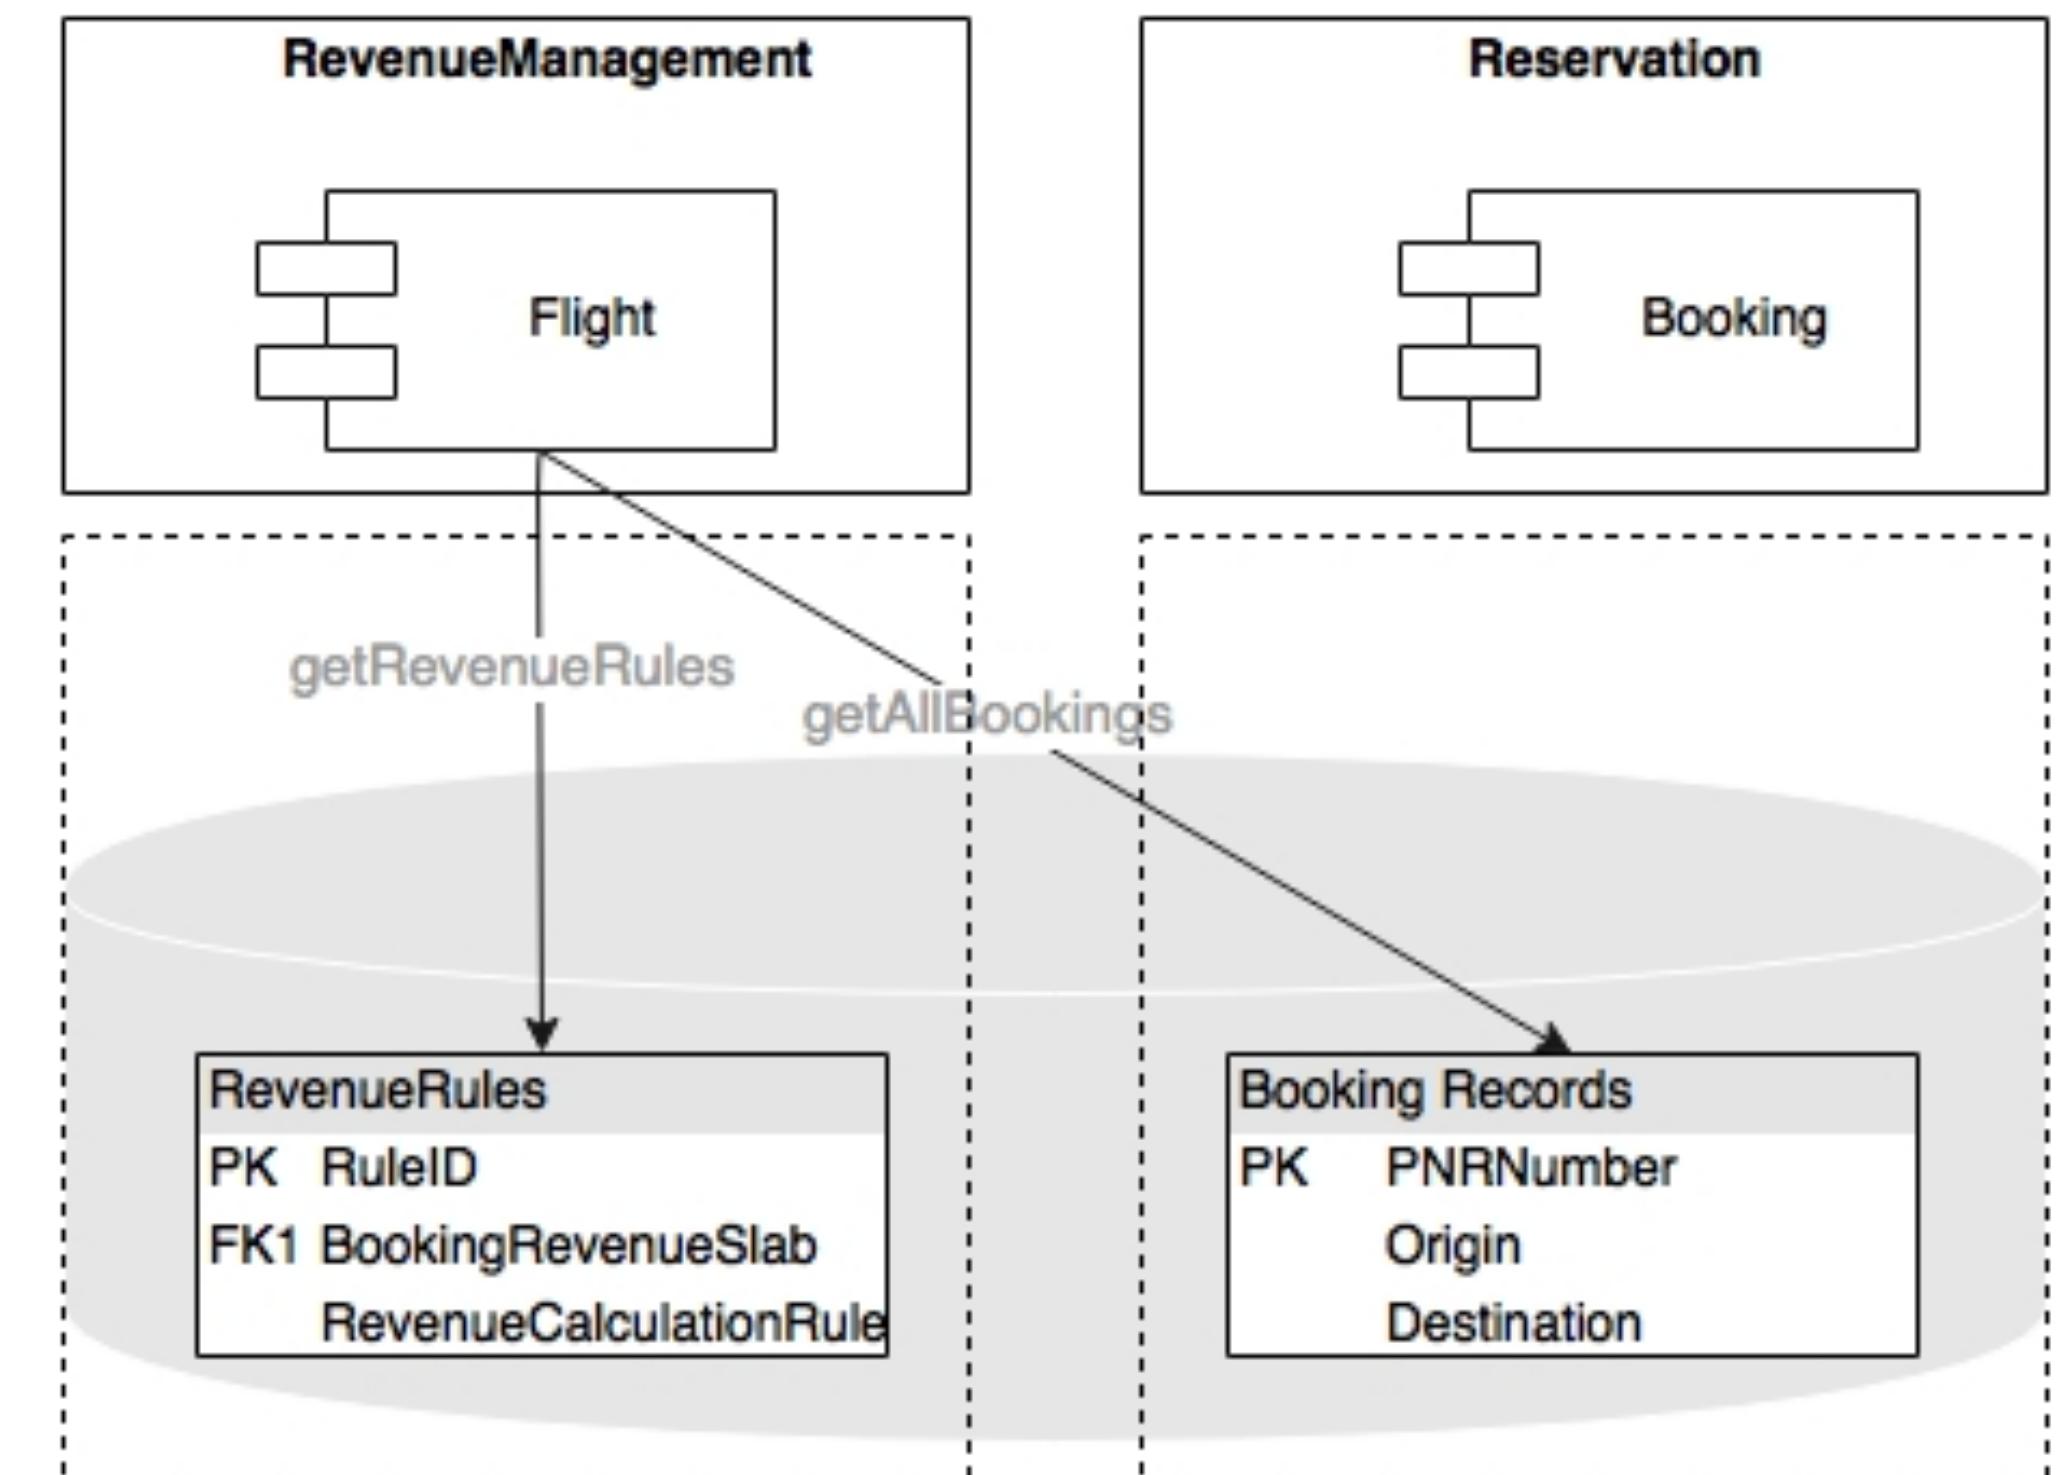
# Retrospection

- Shared data
  - Almost all functional modules require reference data
  - Much of this reference data is neither fully static nor fully dynamic
  - Considering the nature of the query patterns discussed earlier, the approach was to use the reference data as a **shared library**
    - The subsystems were allowed to access the reference data directly using pass-by-reference semantic data instead of going through the EJB interfaces



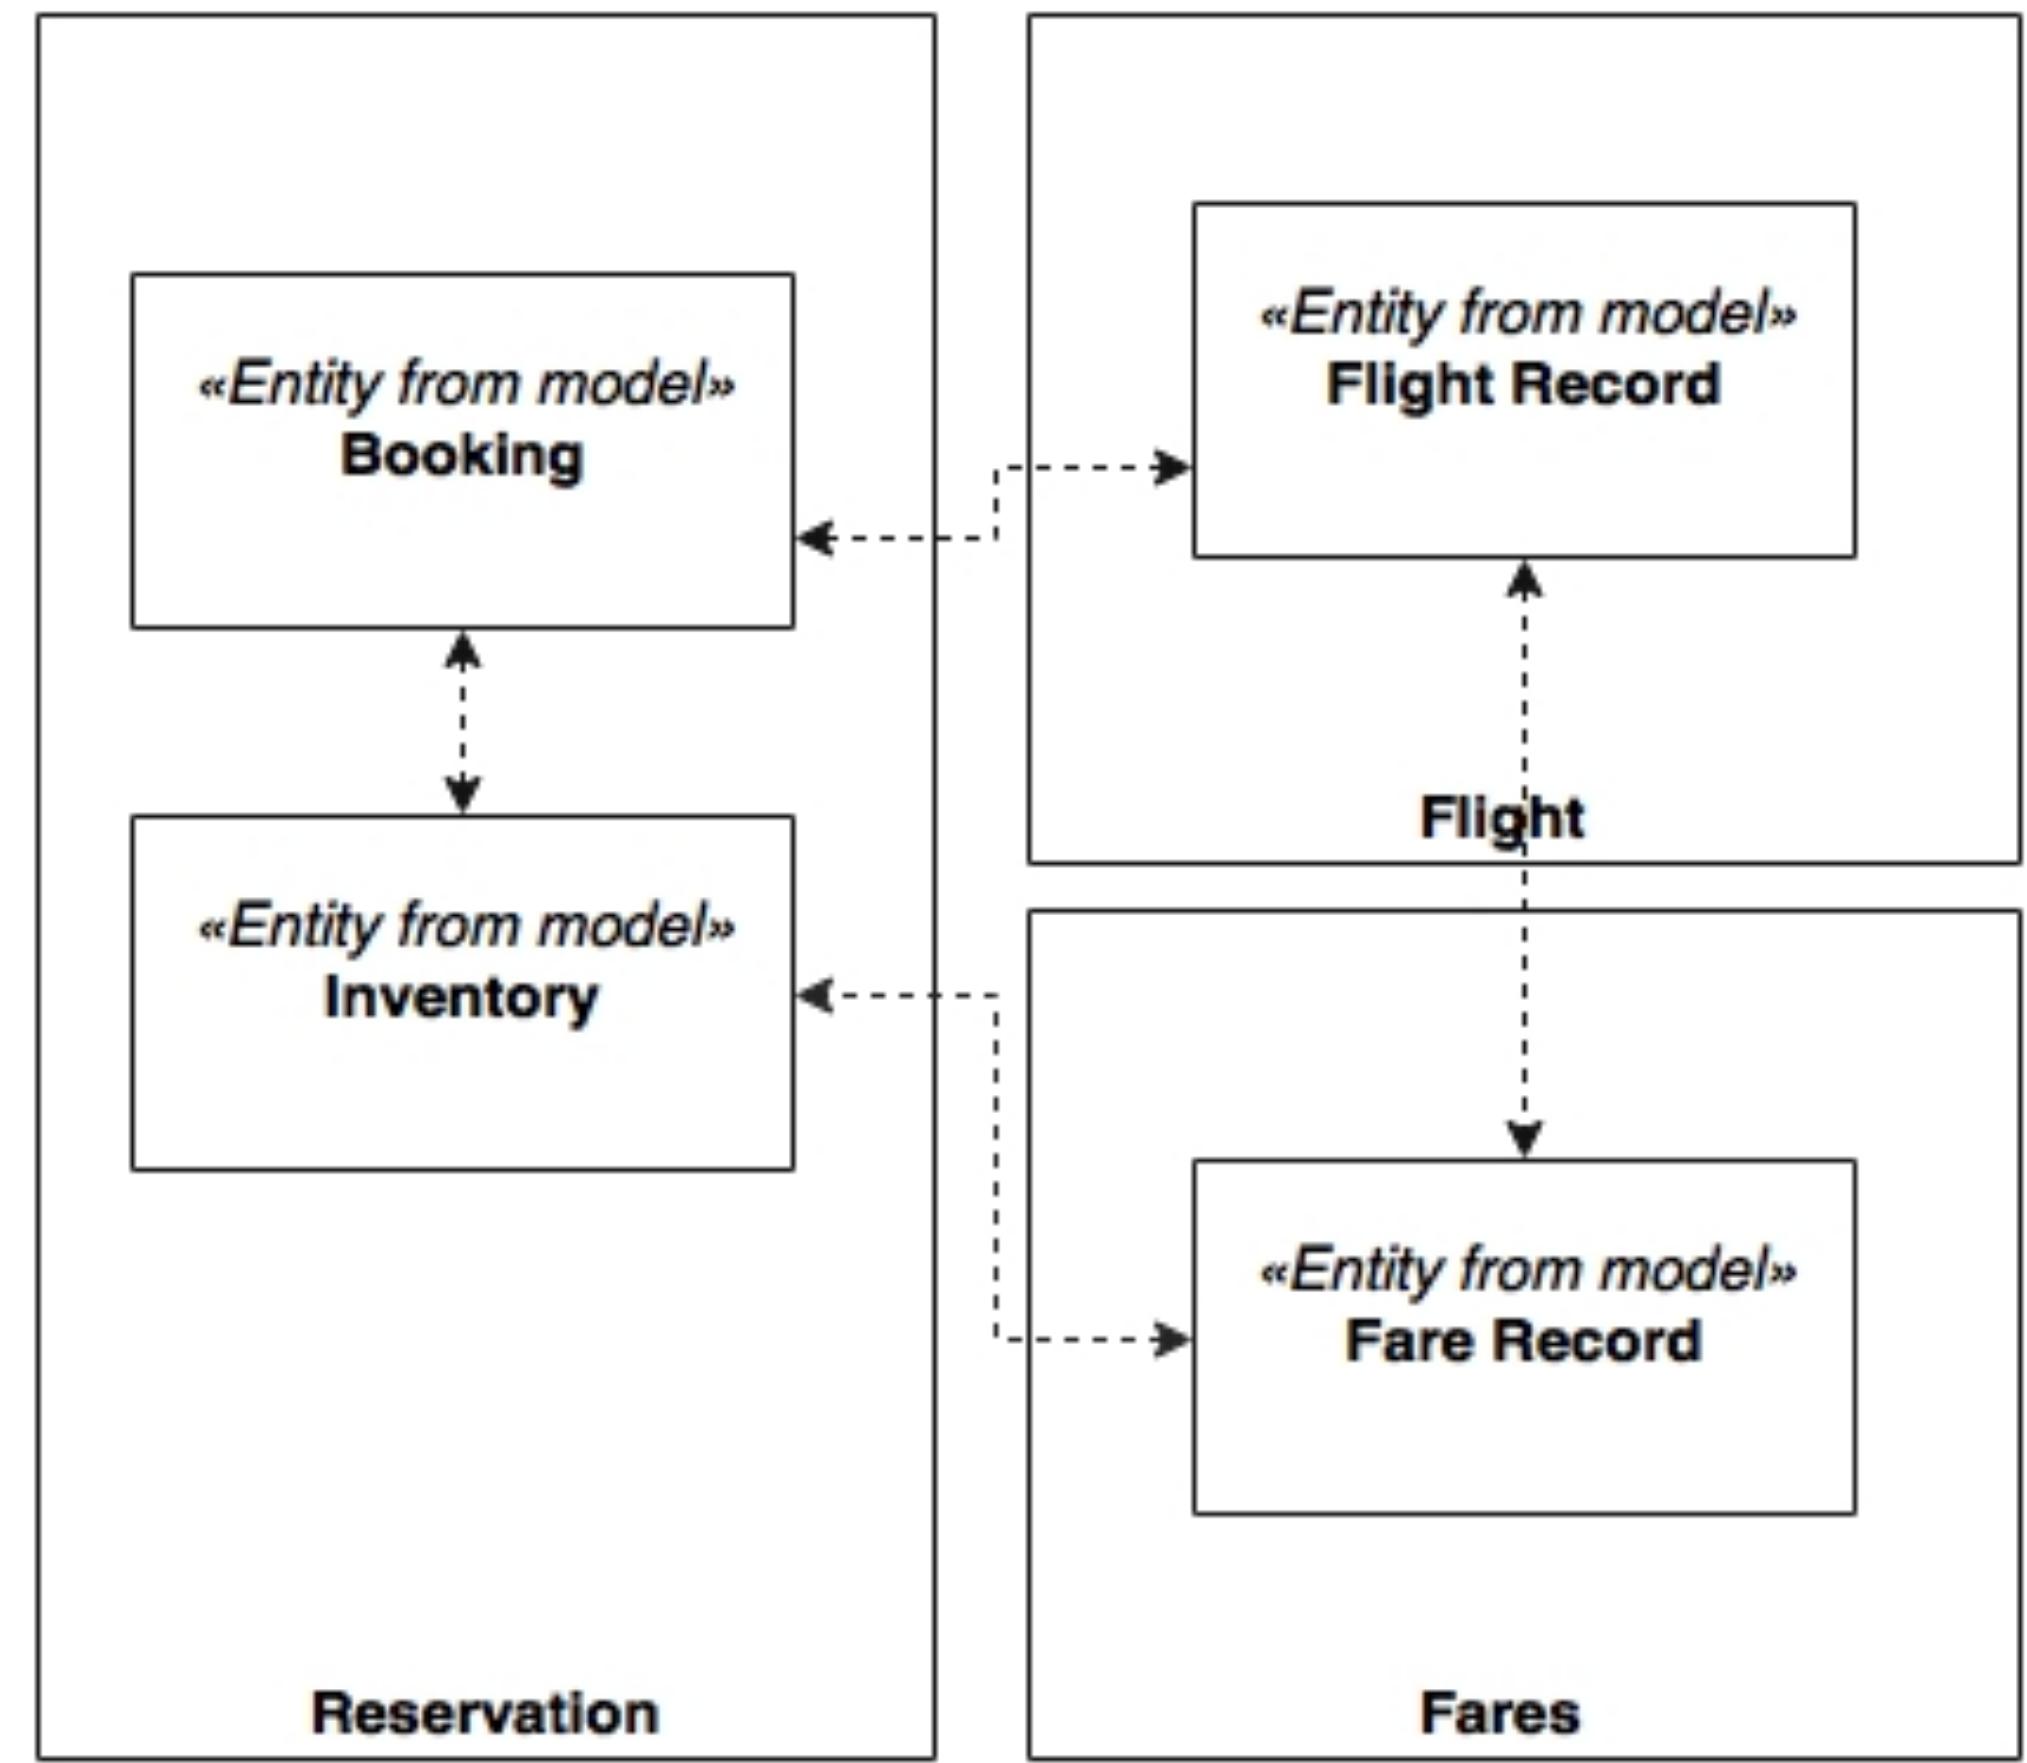
# Retrospection

- Single database
  - The single schema approach opened a **plethora** of issues
  - Native queries
  - Stored procedures



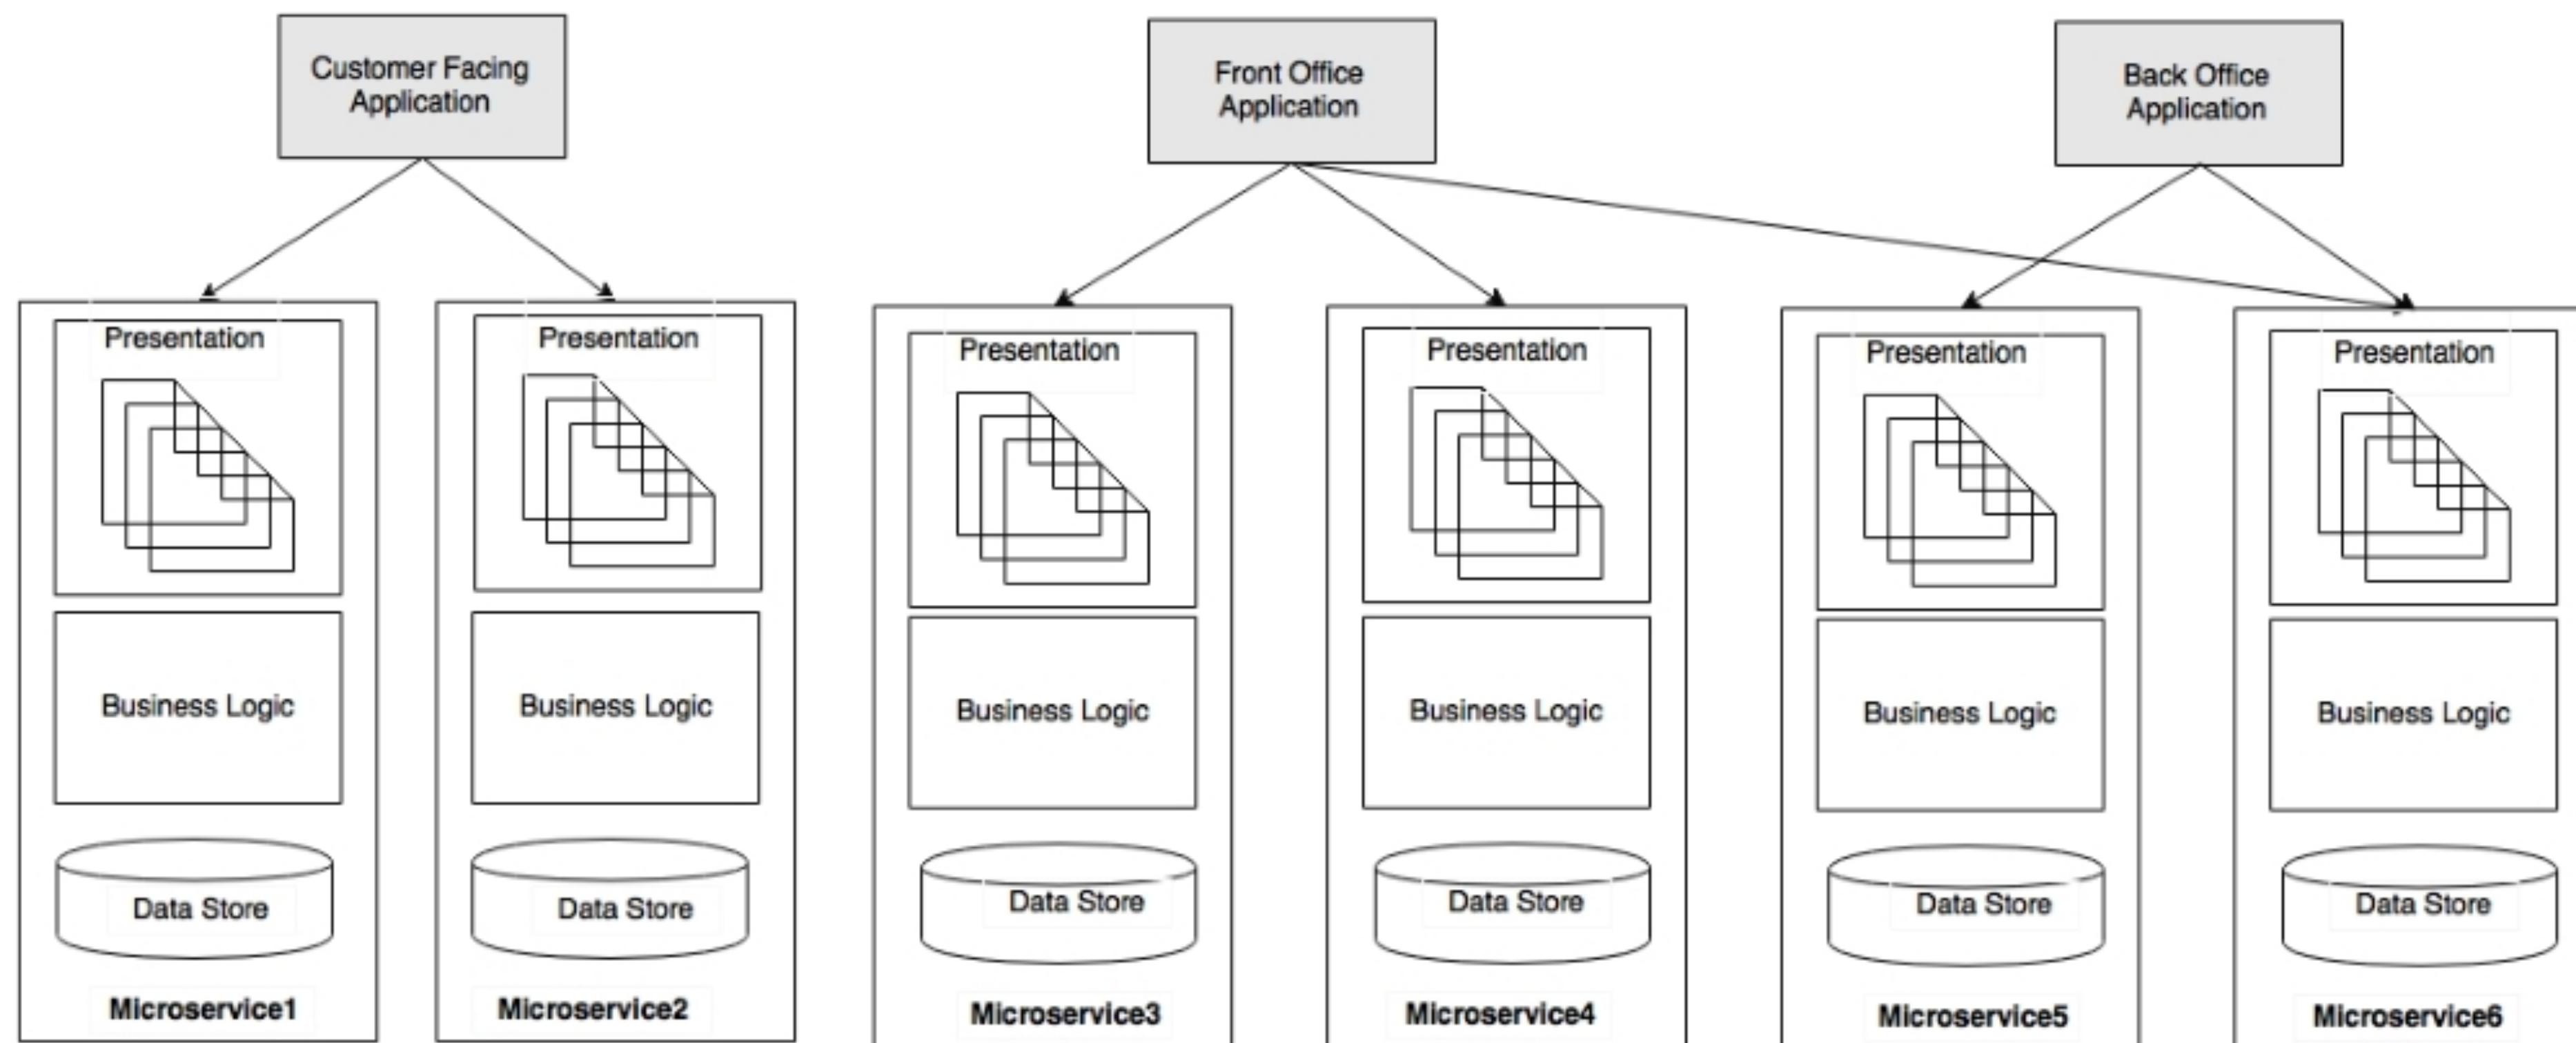
# Retrospection

- Domain boundaries
  - Though the domain boundaries were well established, all the components were packaged as a single EAR file
  - As depicted in the following diagram, hibernate relationships were created across subsystem boundaries



# Microservices to the rescue

- The objective is to move to a microservices-based architecture aligned to the business capabilities. Each microservice will hold the data store, the business logic, and the presentation layer



# The business case

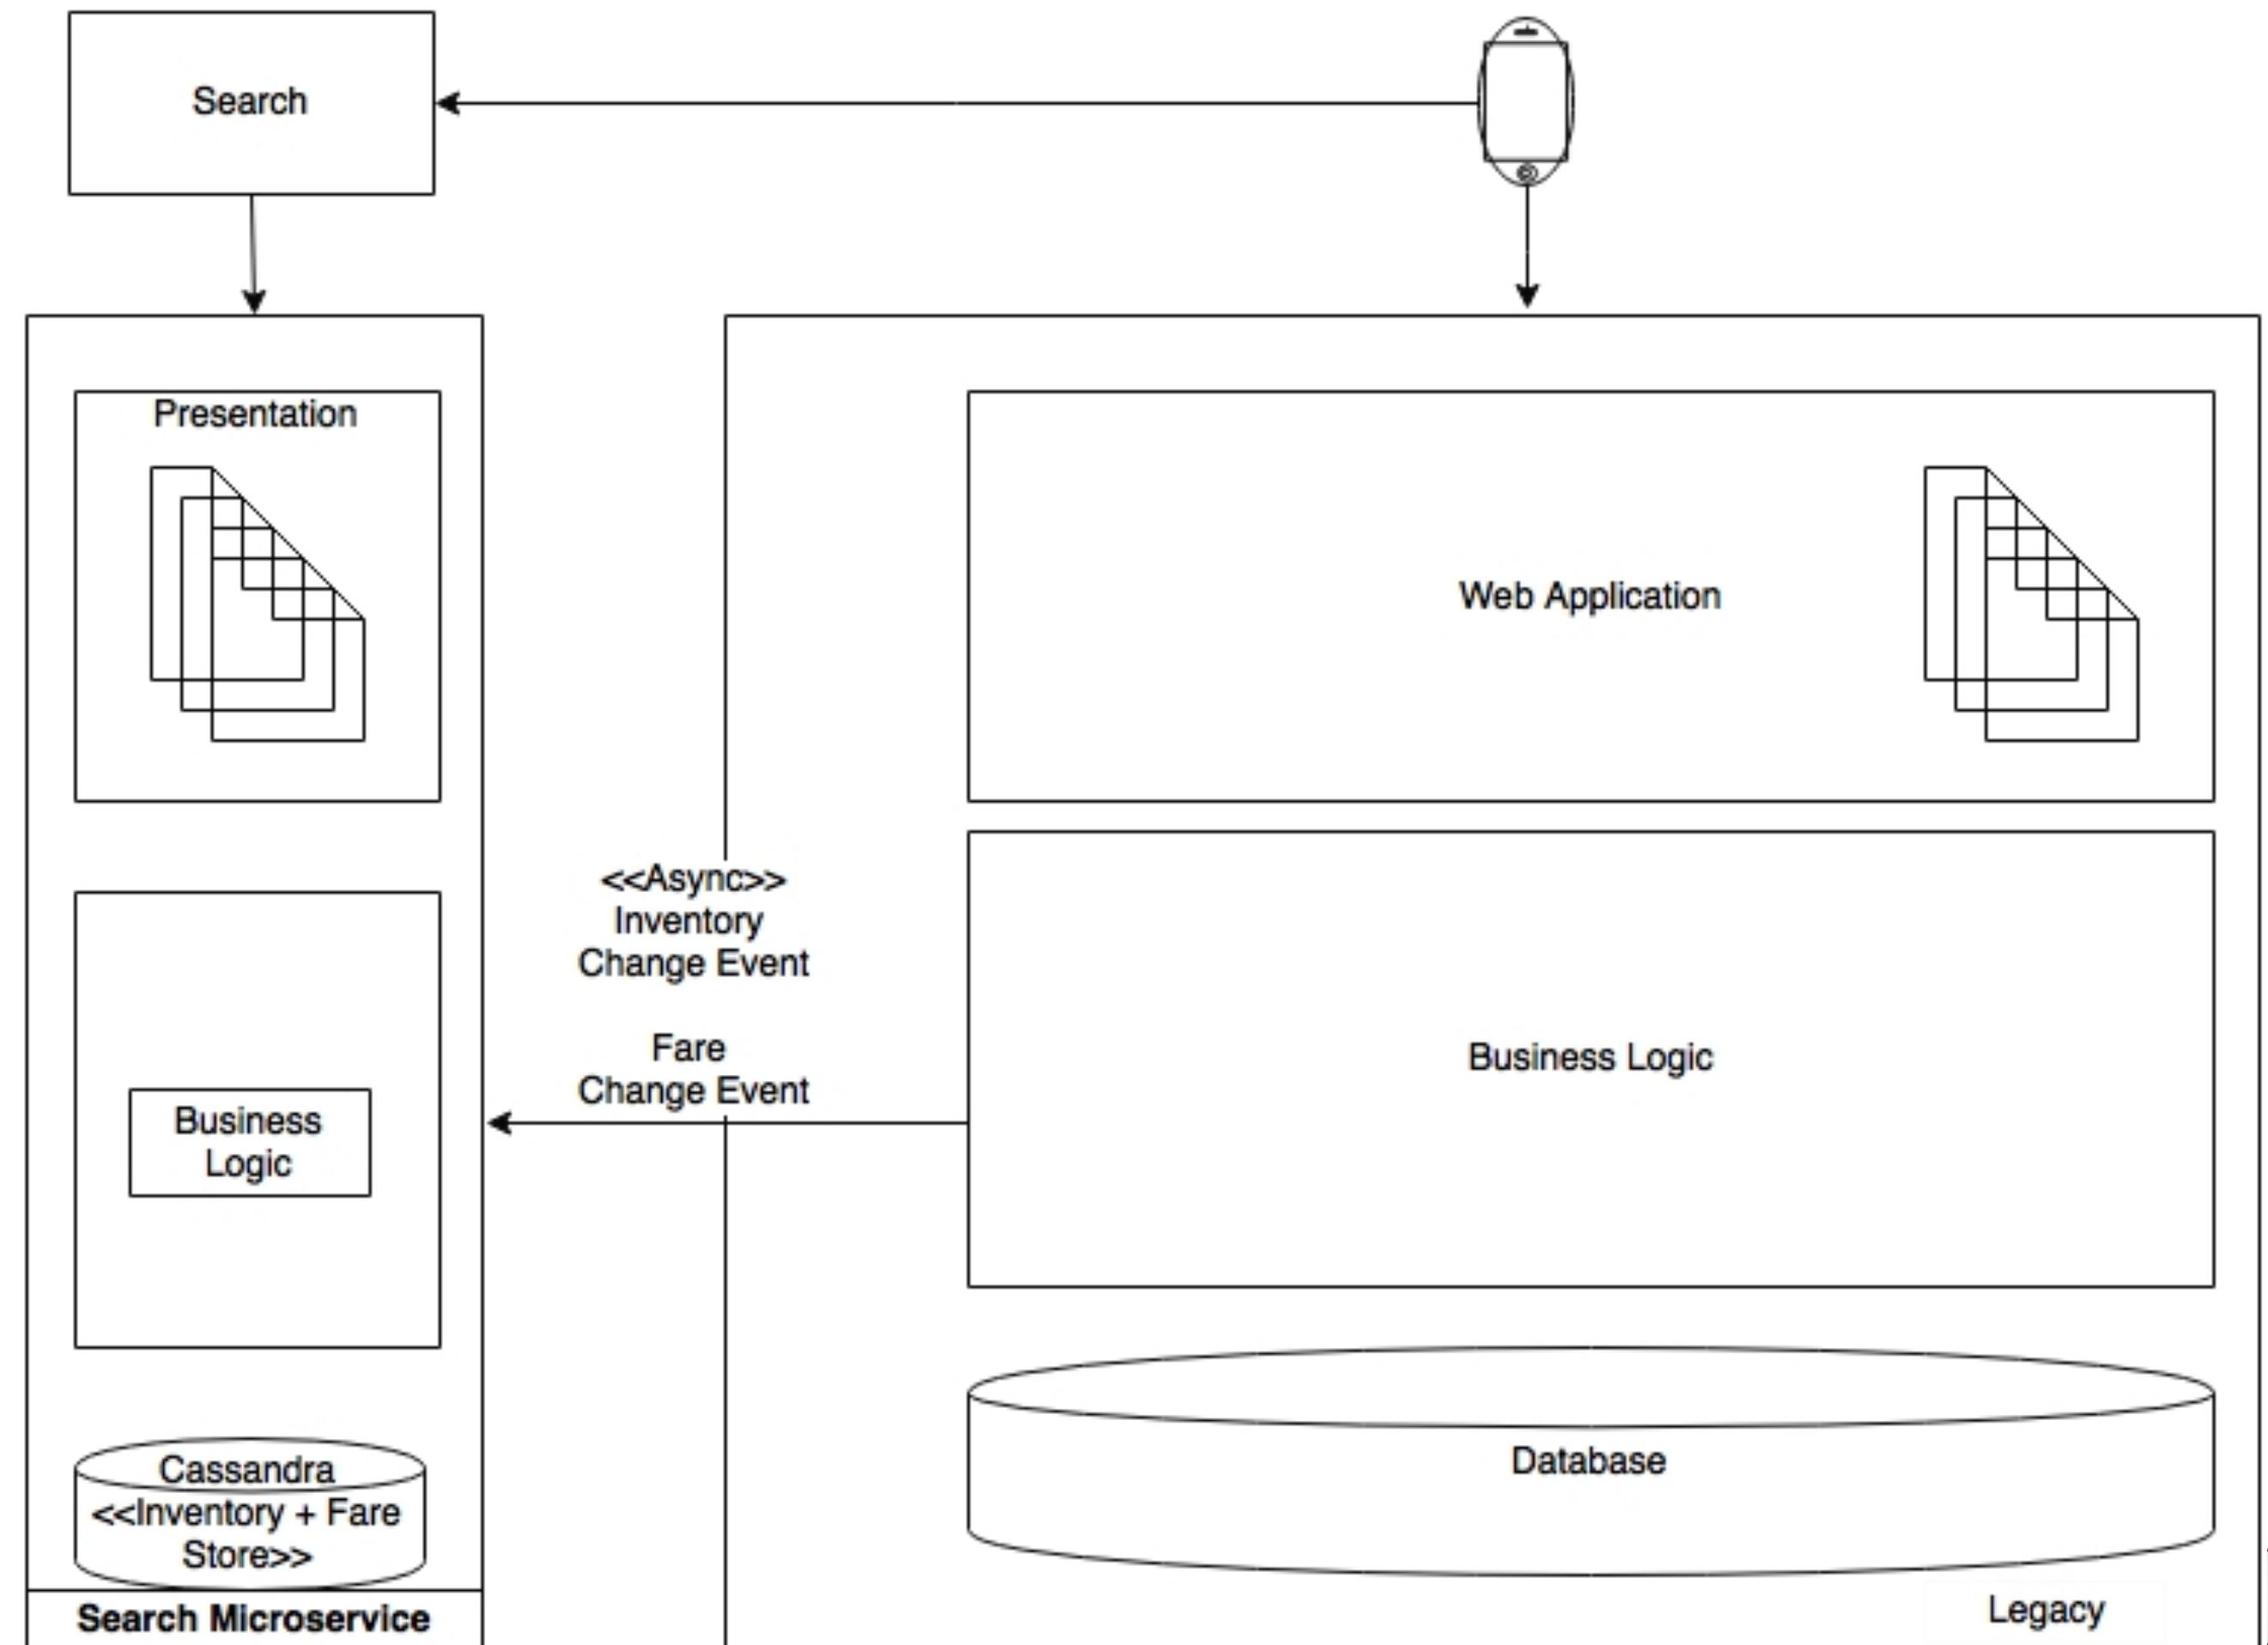
- Microservices offers a full list of benefits (see lecture 2)
  - Service dependencies
  - Physical boundaries (also technology)
  - Selective scaling
  - Technology obsolescence

# Plan the evolution

- It is not simple to break an application that has millions of lines of code, especially if the code has complex dependencies
- How do we break it?
- More importantly, where do we start, and how do we approach this problem?

# Evolutionary approach

- At every step, a microservice will be created outside of the monolithic application, and traffic will be diverted to the new service
- A number of key questions need to be answered from the transition point of view
  - Identification of microservices' boundaries
  - Prioritizing microservices for migration
  - Handling data synchronization during the transition phase
  - Handling user interface integration, working with old and new user interfaces
  - Handling of reference data in the new system
  - Testing strategy to ensure the business capabilities are intact and correctly reproduced
  - Identification of any prerequisites for microservice development such as microservices capabilities, frameworks, processes, and so on

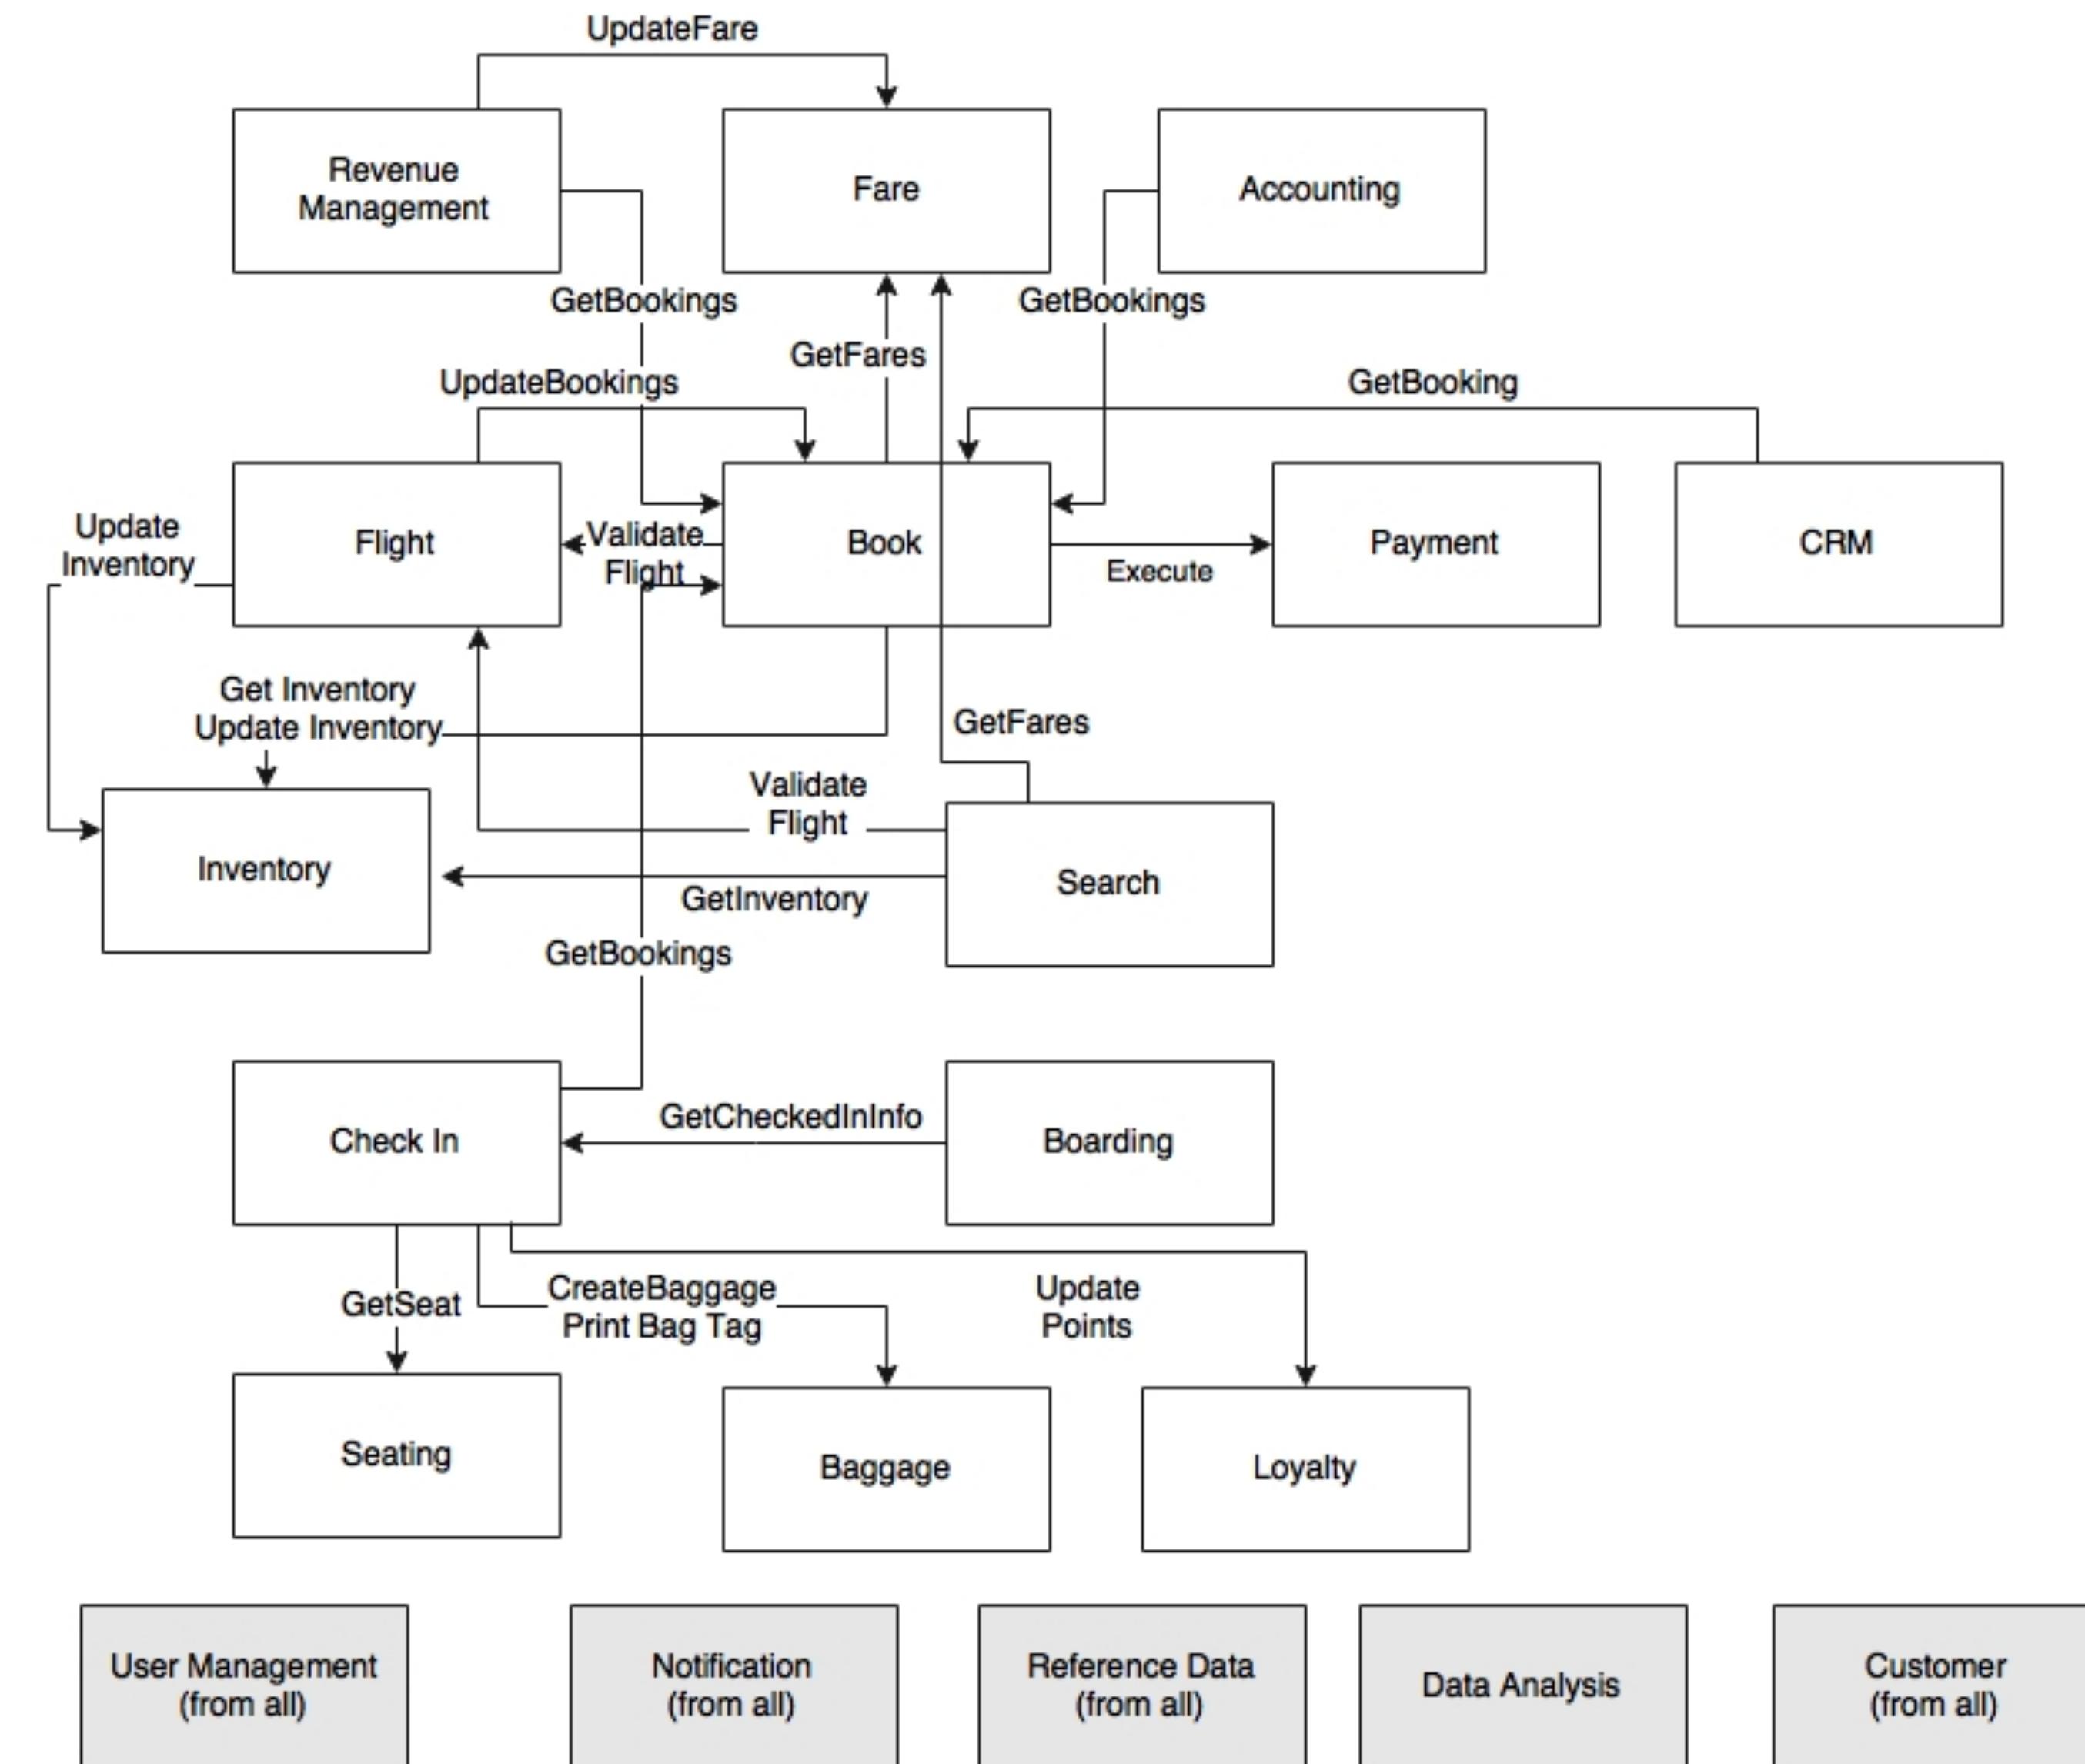


# Identification of microservices boundaries

- Like in SOA, a service decomposition is the best way to identify services
  - Decomposition stops at a business capability or bounded context
- A top-down approach is typically used for domain decomposition.
- The bottom-up approach is also useful in the case of breaking an existing system, as it can utilize a lot of practical knowledge, functions, and behaviors of the existing monolithic application

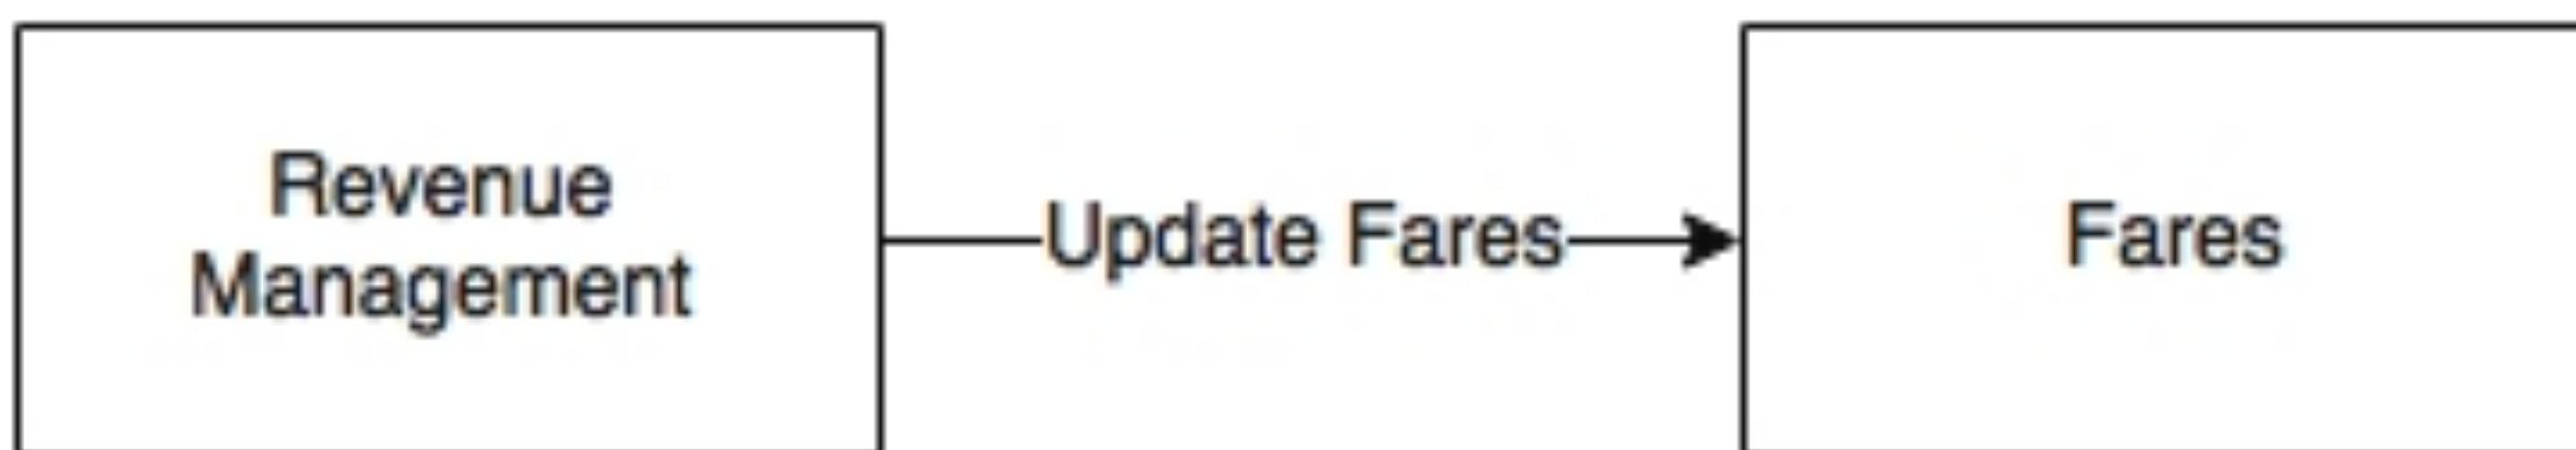
# Analyze dependencies

- Analyze the **dependencies** between the initial set of candidate microservices that we created
  - **Dependency graph**
- The bottom layer shows cross-cutting capabilities that are used across multiple modules



# Events as opposed to query

- Dependencies could be query-based or event-based

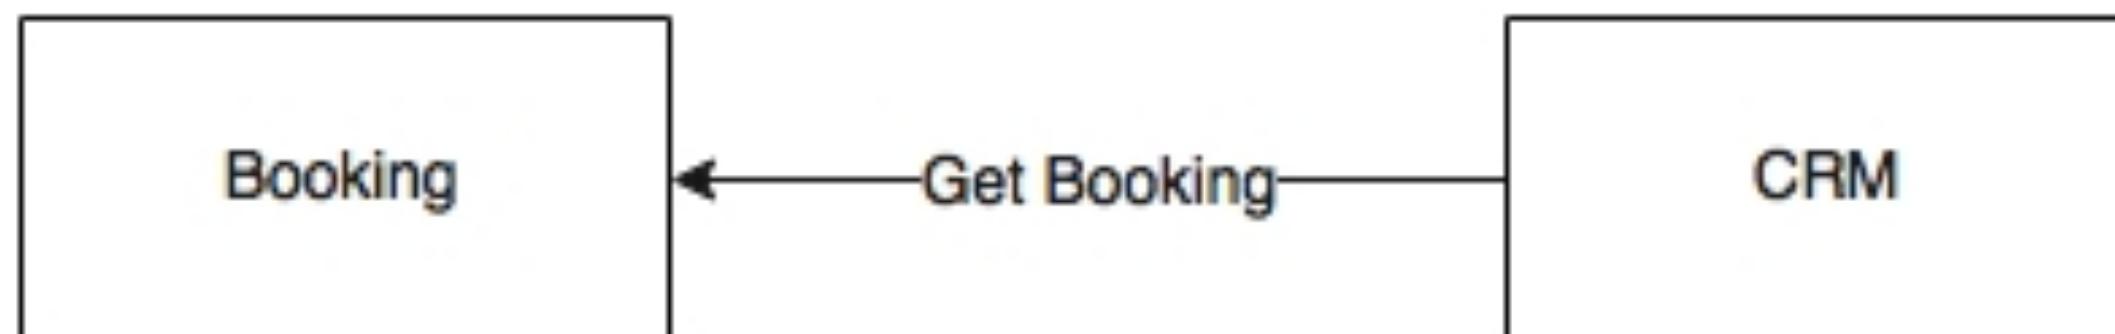


- Revenue Management is a module used for calculating optimal fare values, based on the booking demand forecast. In case of a fare change between an origin and a destination, Update Fare on the Fare module is called by Revenue Management to update the respective fares in the Fare module.

# Events as opposed to query

- An alternate way of thinking is that the Fare module is subscribed to Revenue Management for any changes in fares, and Revenue Management publishes whenever there is a fare change.
- This reactive programming approach gives an added flexibility by which the Fares and the Revenue Management modules could stay independent, and connect them through a reliable messaging system

# Events as opposed to query

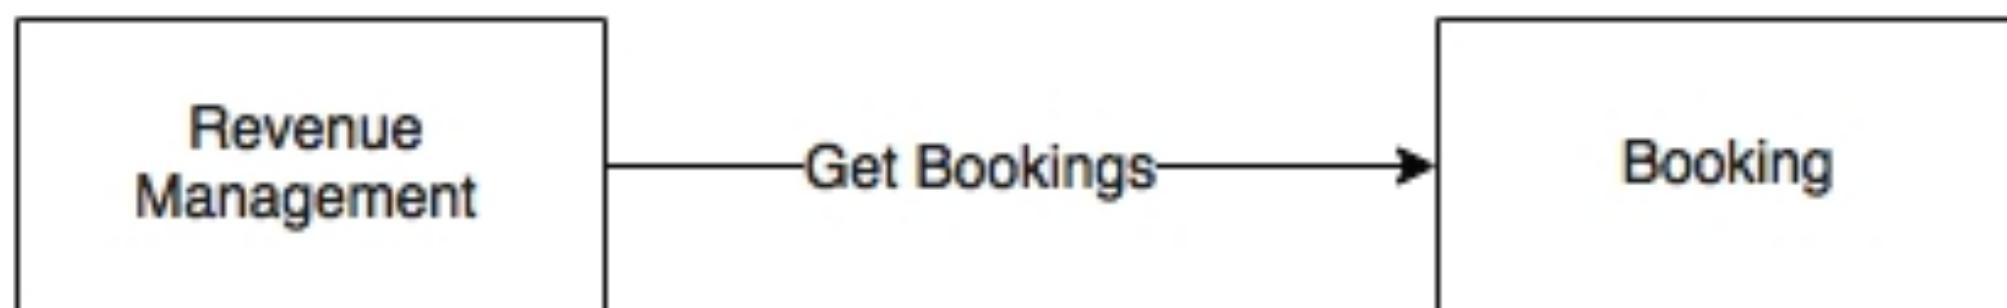


- The CRM module is used to manage passenger complaints. When CRM receives a complaint, it retrieves the corresponding passenger's Booking data
  - In reality, the number of complaints are negligibly small when compared to the number of bookings
- If we blindly apply the previous pattern where CRM subscribes to all bookings, we will find that it is not cost effective

# Homework 5.1

- Examine another scenario between the Check-in and Booking modules
- Instead of Check-in calling the Get Bookings service on Booking, can Check-in listen to booking events?
- This is possible, but the challenge here is that a booking can happen 360 days in advance, whereas Check-in generally starts only 24 hours before the fight departure.

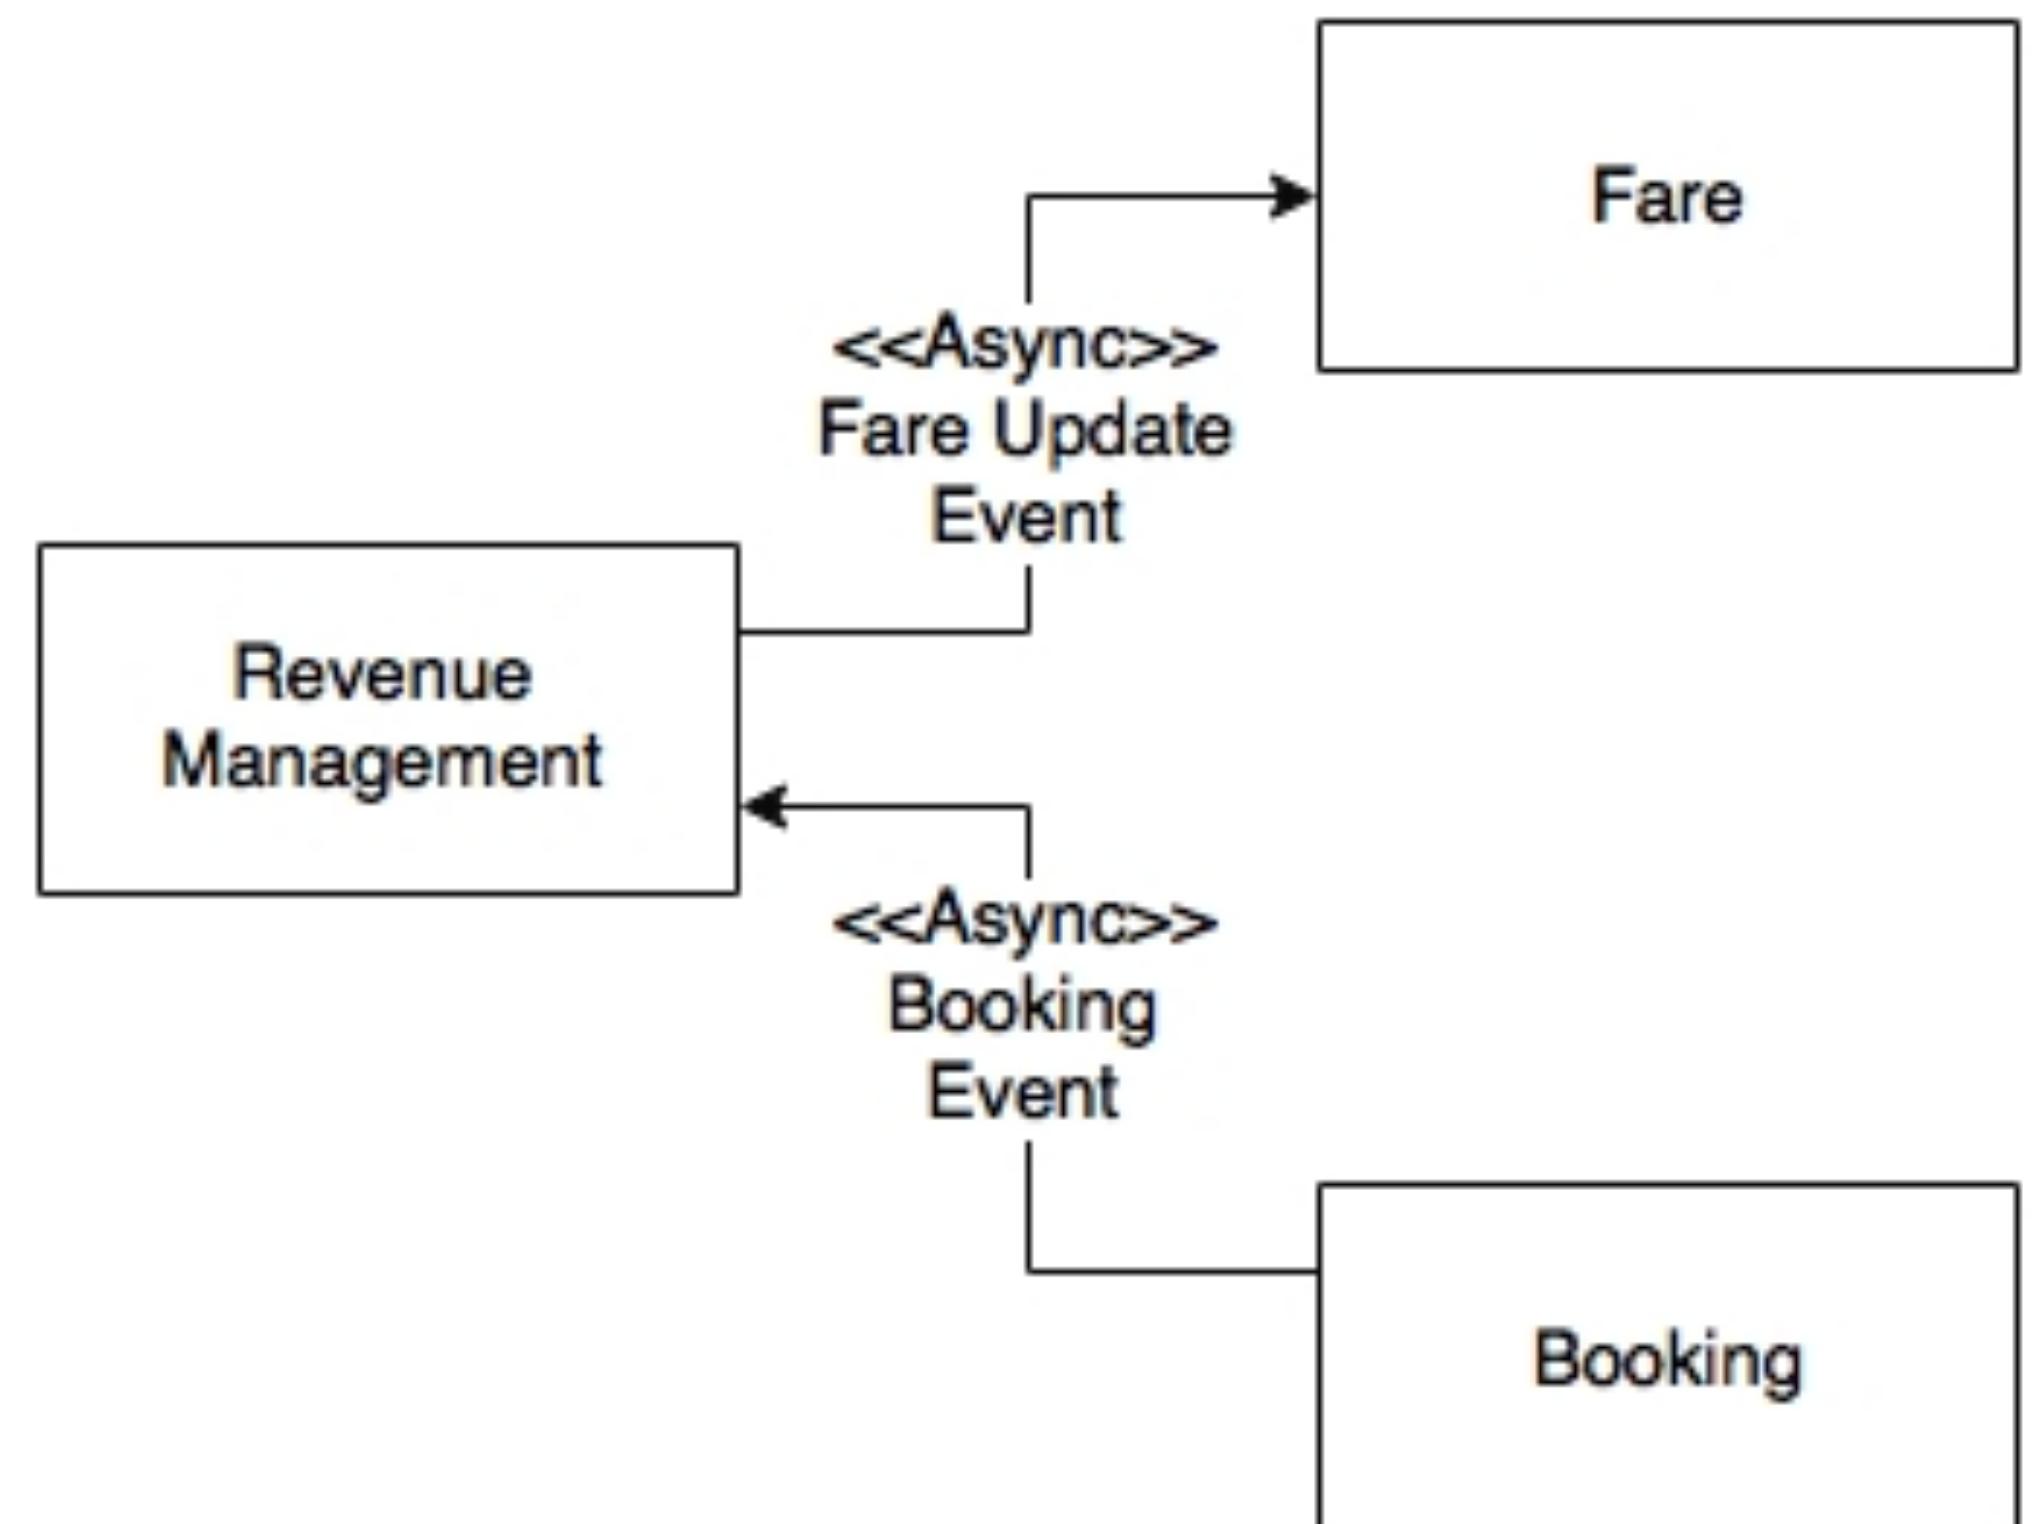
# Events as opposed to synchronous updates



- Apart from the query model, a dependency could be an update transaction as well
- Revenue Management has a schedule job that calls Get Booking on Booking to get all incremental bookings (new and changed) since the last synchronization

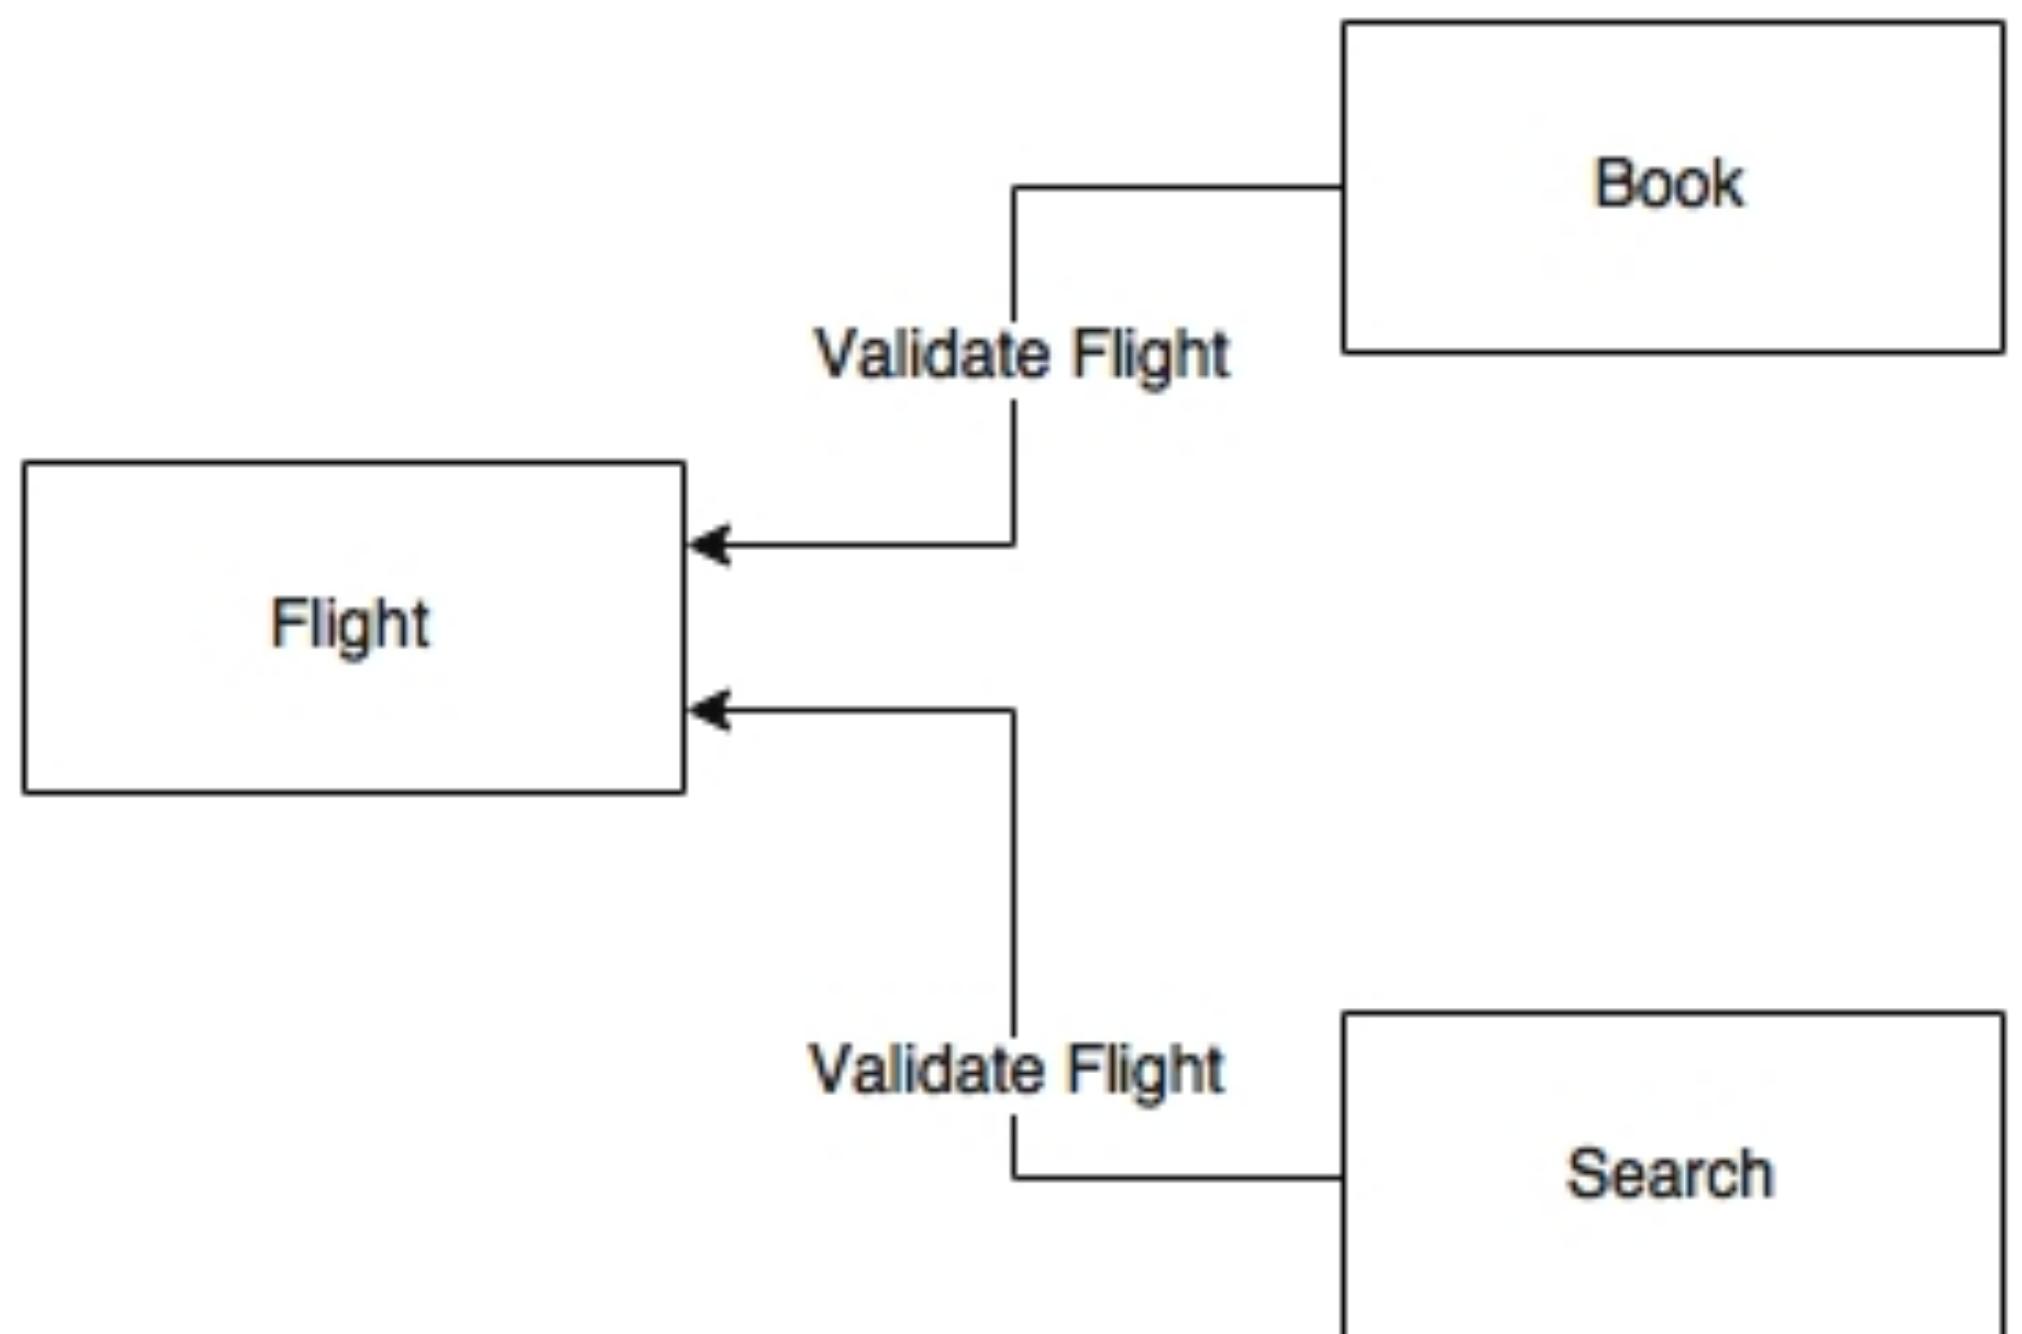
# Events as opposed to synchronous updates

- The same pattern could be applied in many other scenarios such as from Booking to Accounting, from Flight to Inventory, and also from Flight to Booking



# Challenge requirements

- The Validate Flight call is to validate the input flight data coming from different channels
- An alternate way of solving this is to adjust the inventory of the flight based on these given conditions
- As far as Search and Booking are concerned, both just look up the inventory instead of validating flights for every request

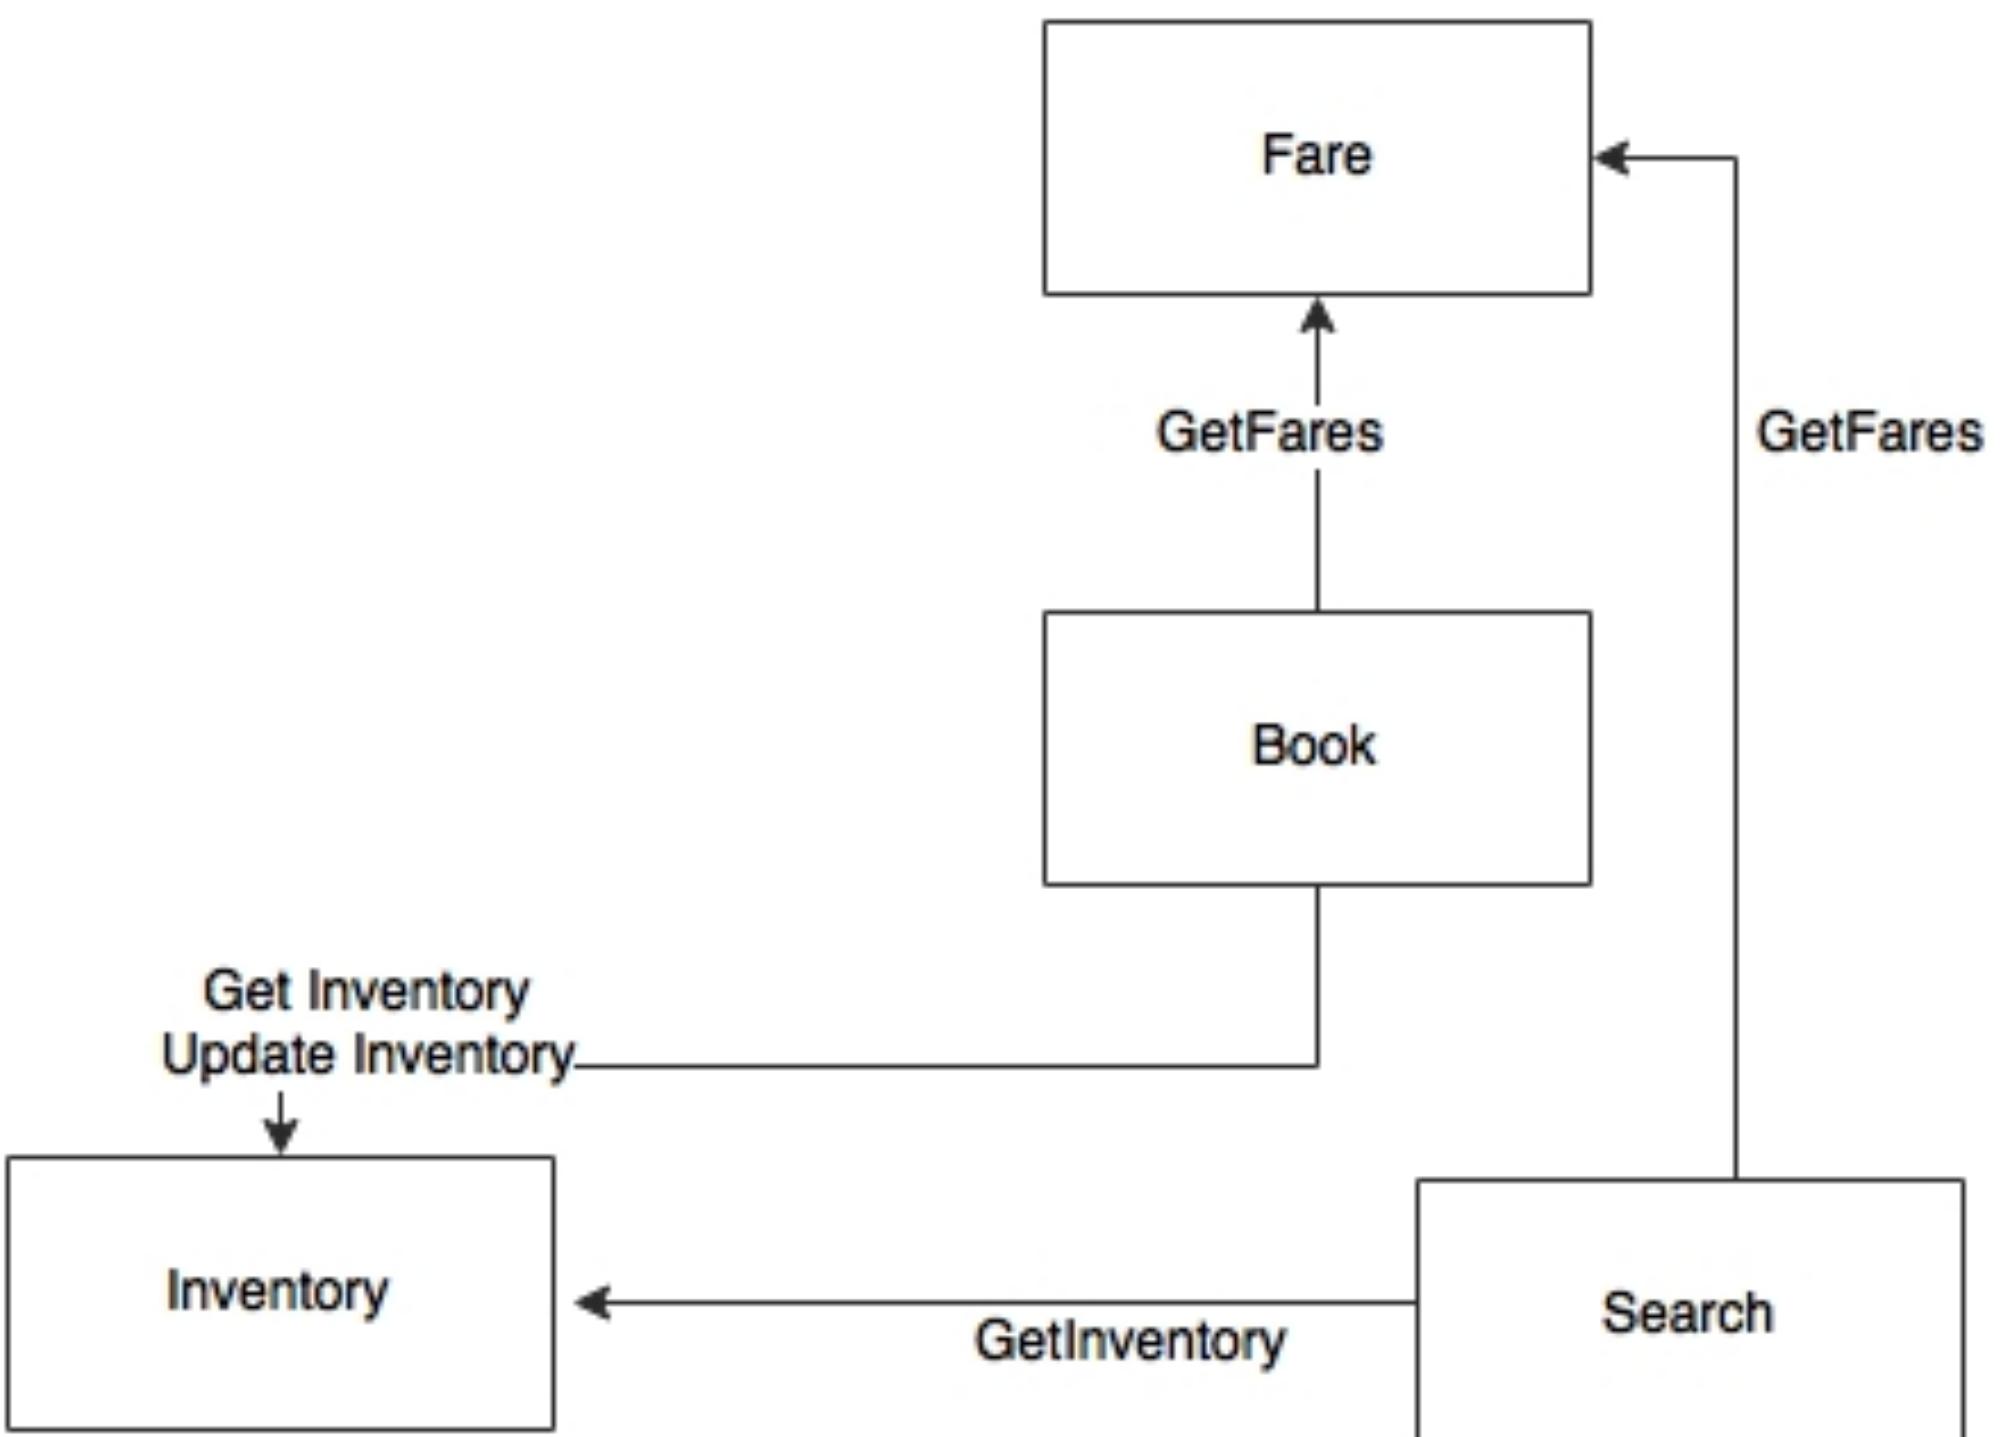


# Challenge service boundaries

- The Seating function runs a few algorithms based on the current state of the seat allocation in the airplane
- However, other than Check-in, no other module is interested in the Seating function
- From a business capability perspective, Seating is just a function of Check-in, not a business capability by itself
- Therefore, it is better to embed this logic inside Check-in itself

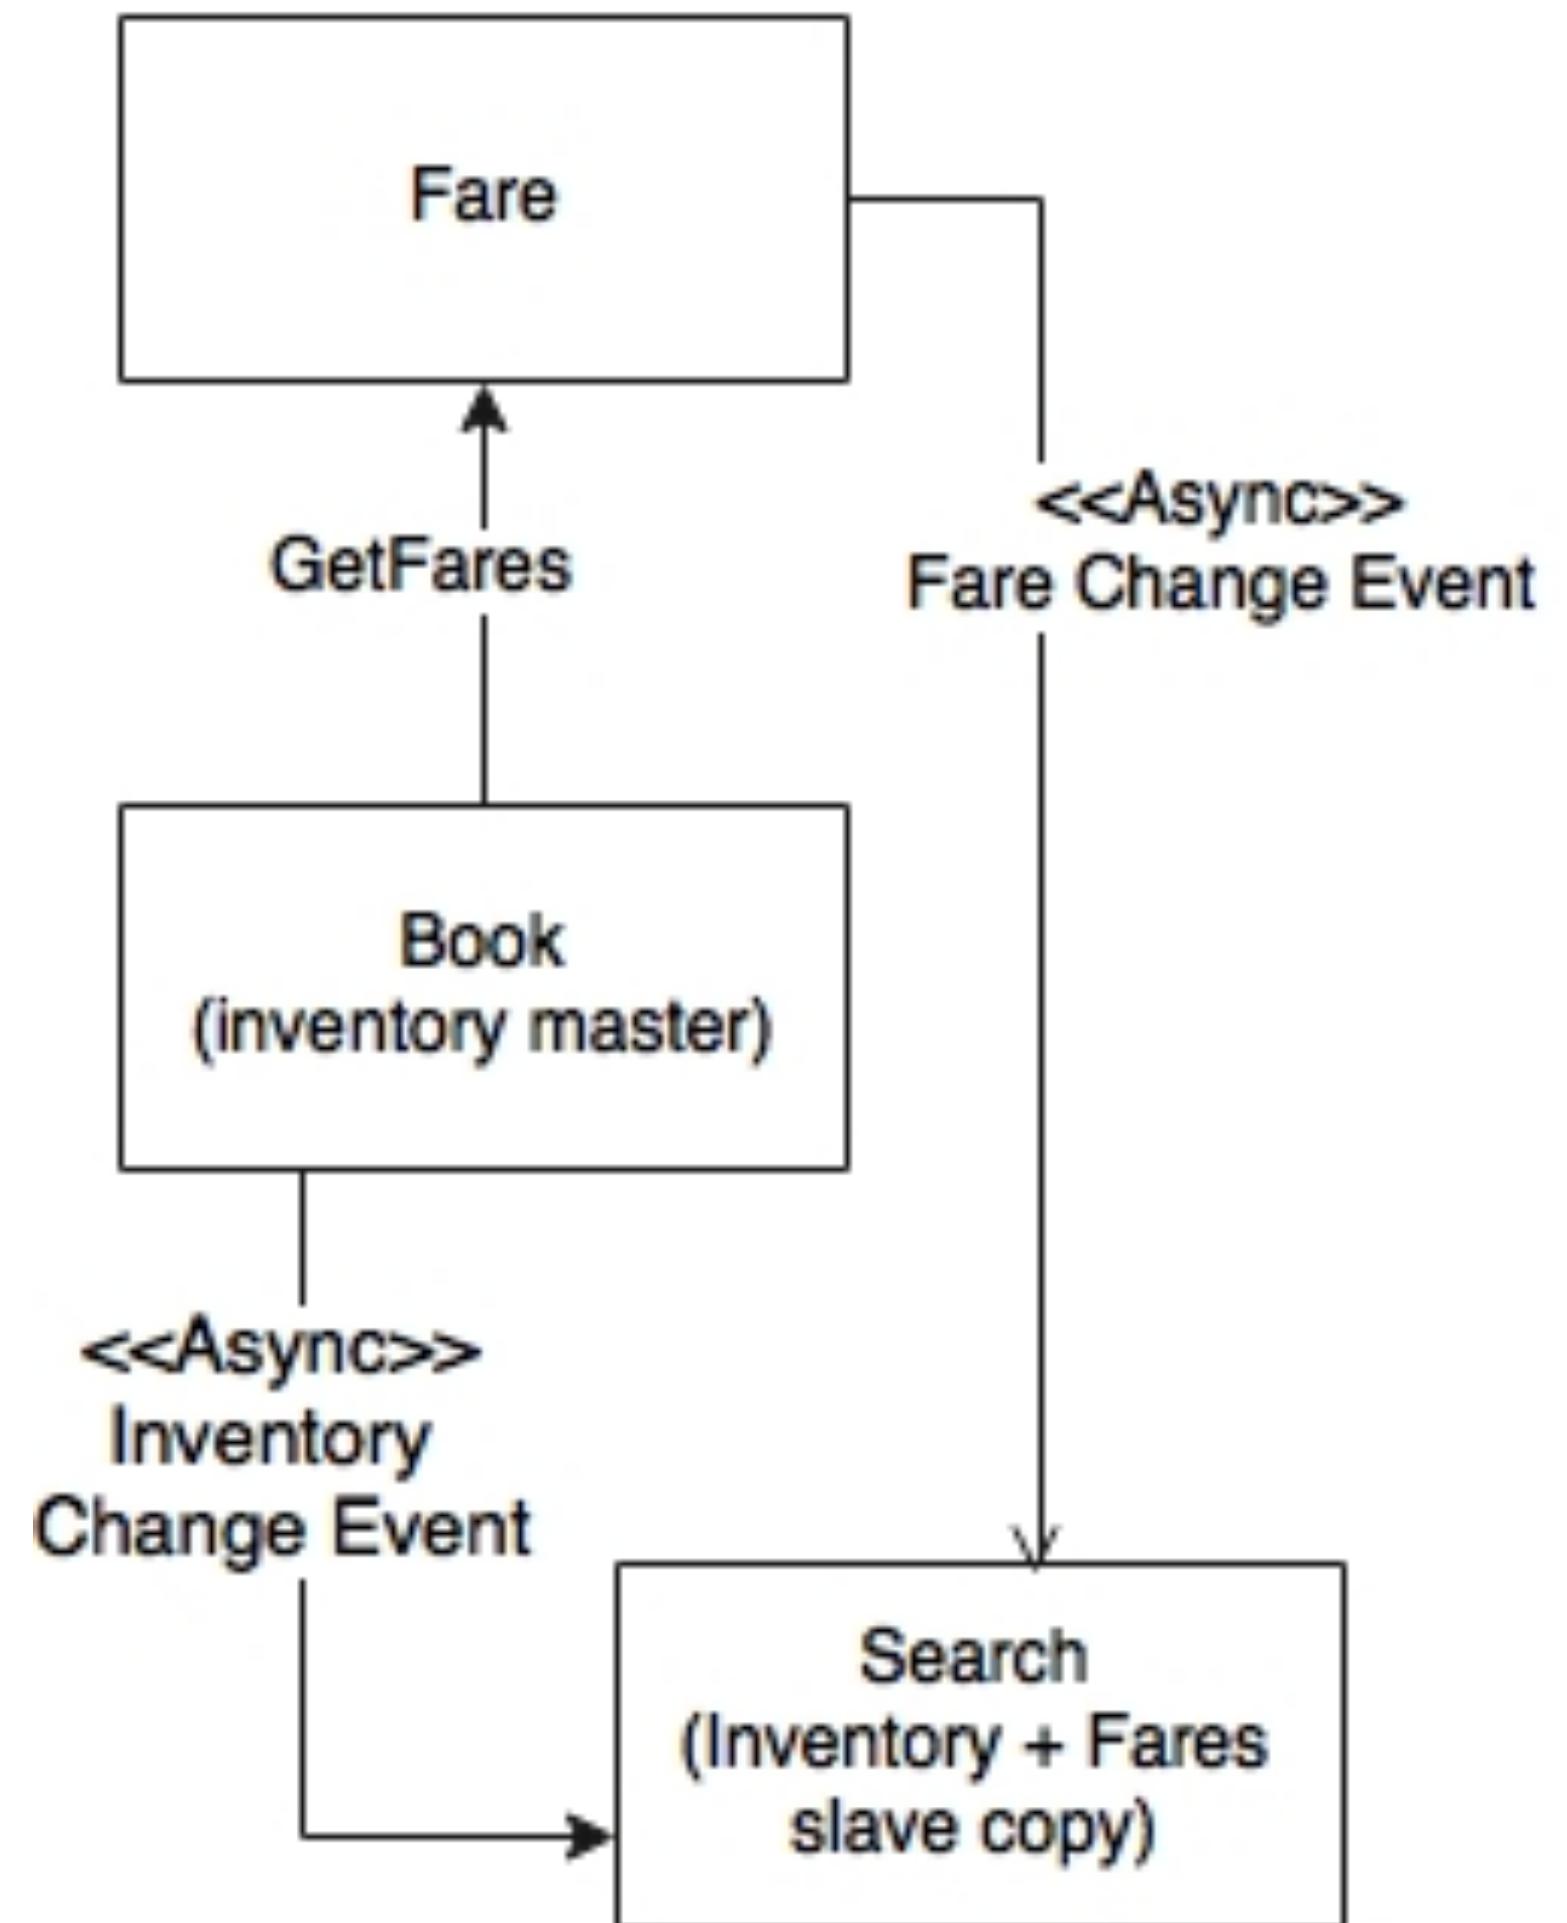
# The Book, Search, and Inventory functions

- Assume, for the time being, that Search, Inventory, and Booking are moved to a single microservice named Reservation
- Search transactions are 10 times more frequent than the booking transactions
  - We need different scalability models for search and booking



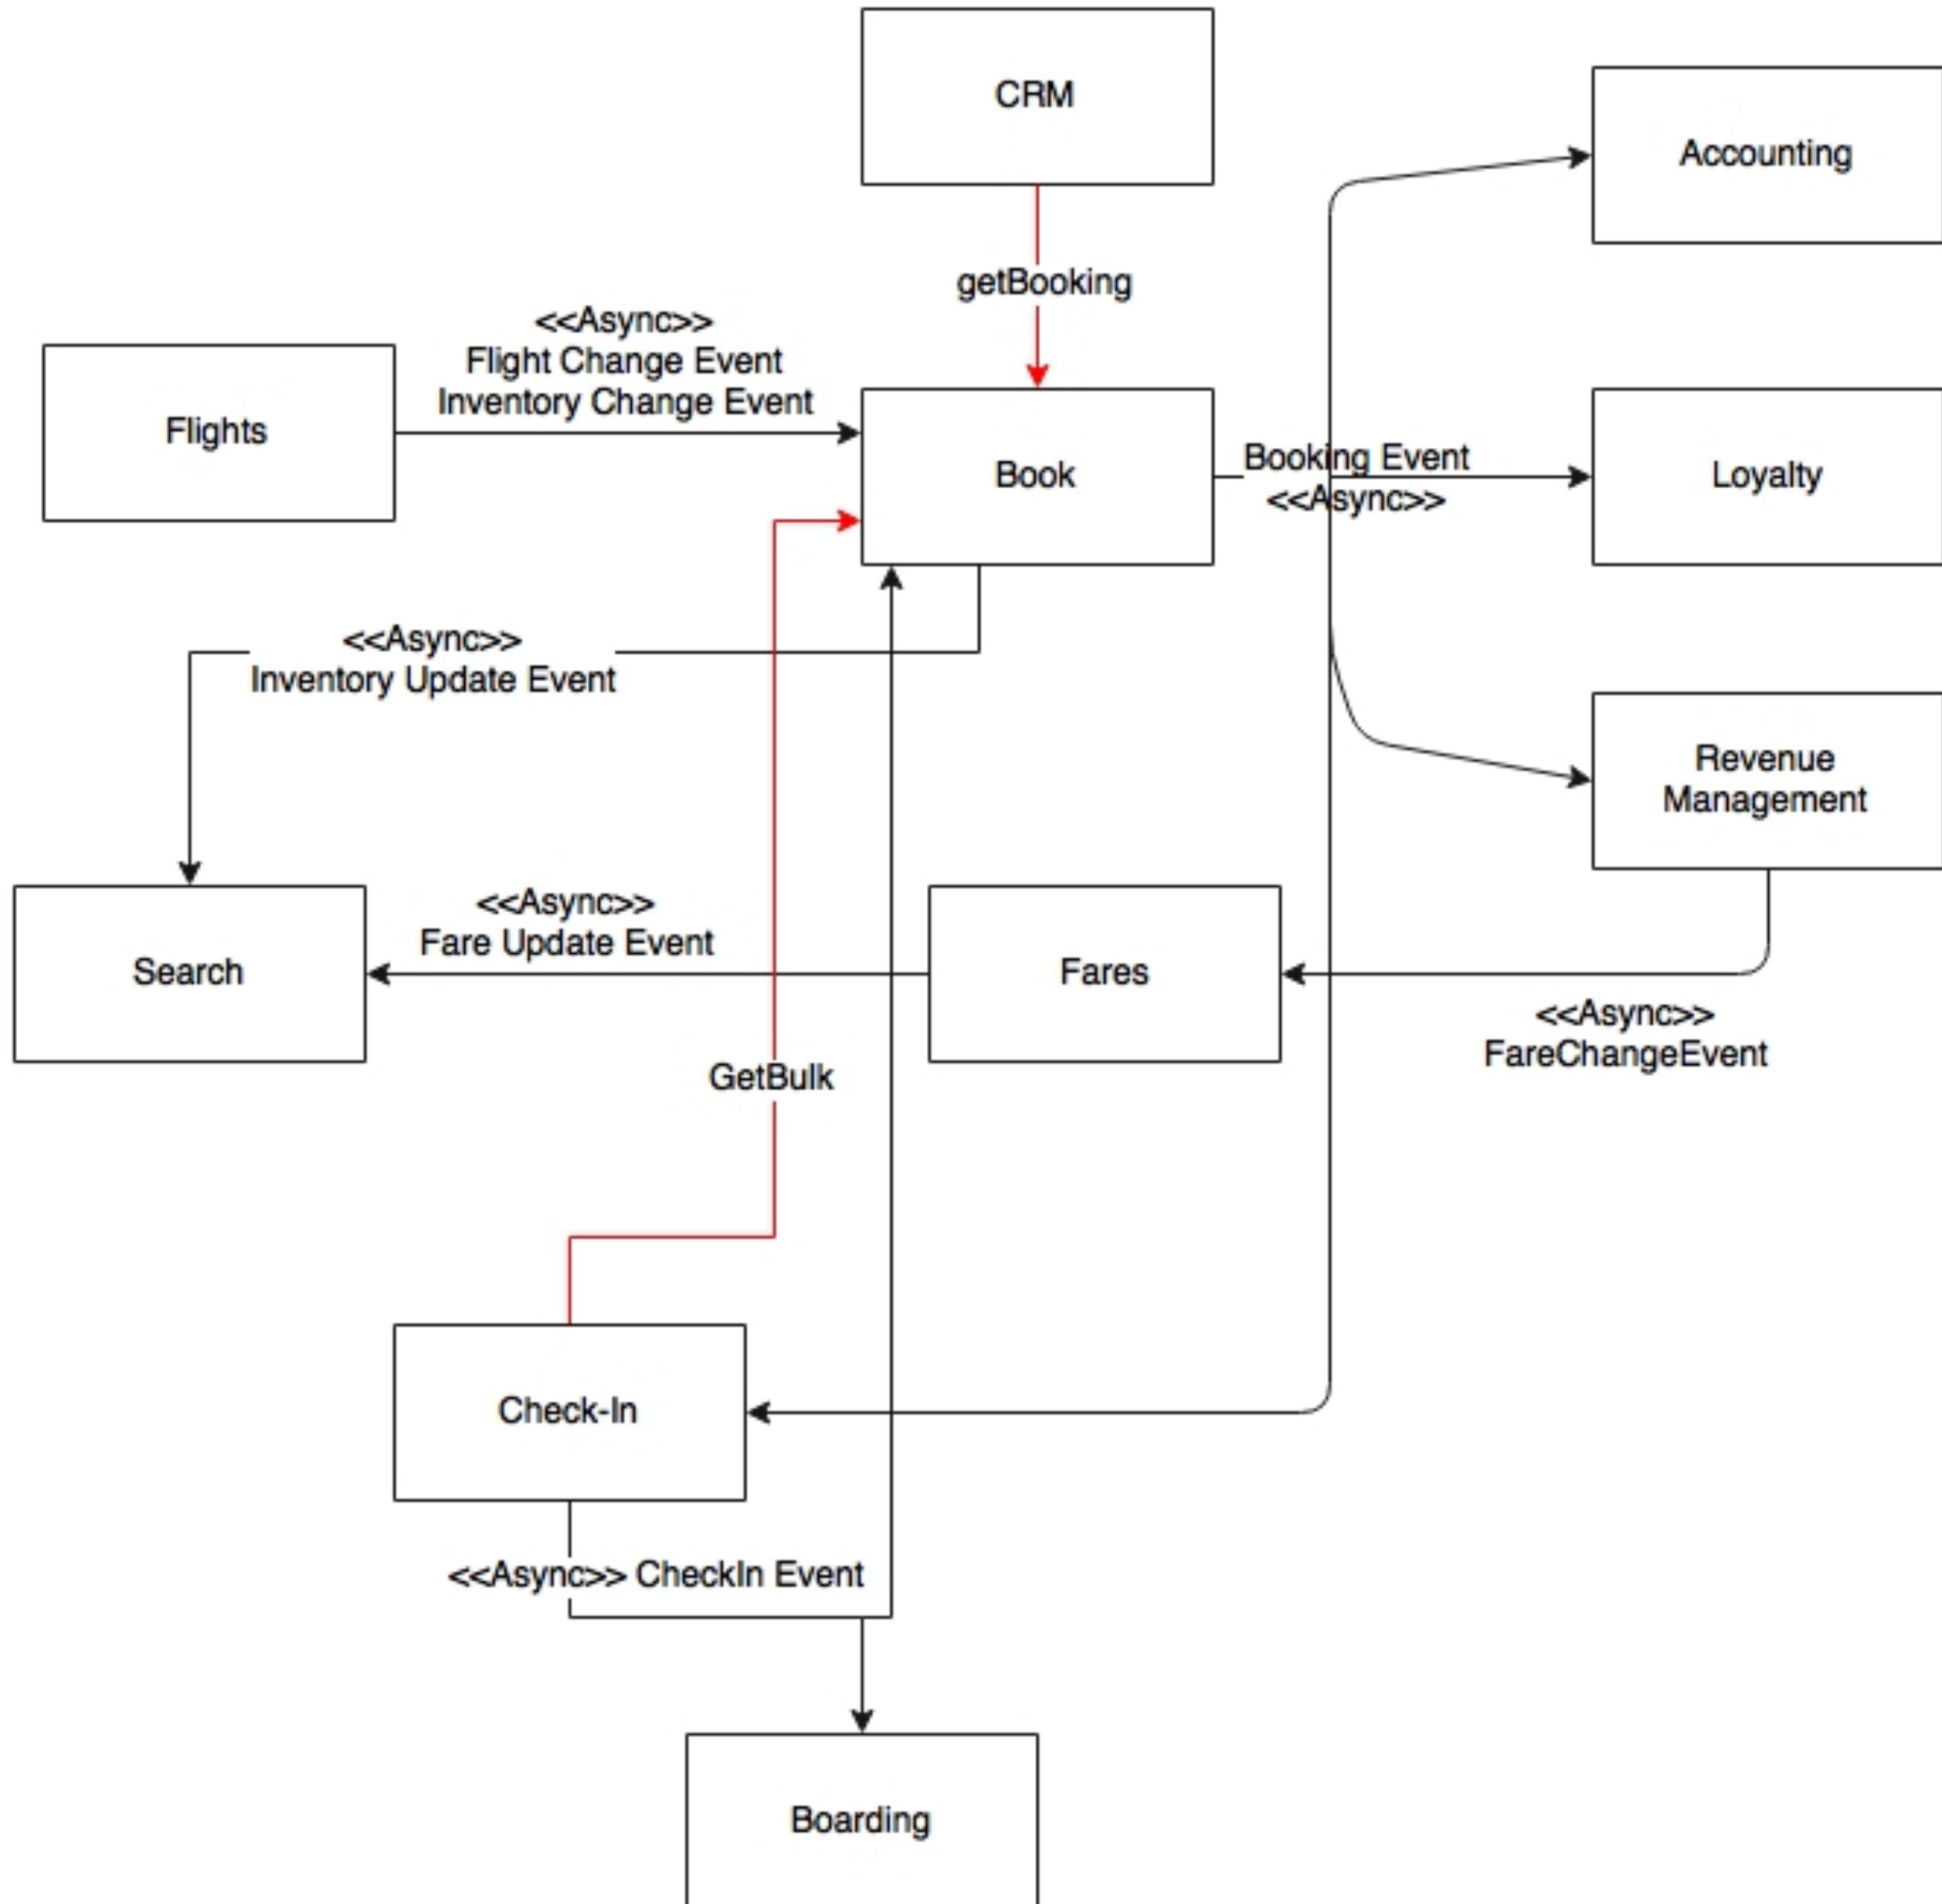
# The Book, Search, and Inventory functions

- Let us assume that we remove Search
- Only Inventory and Booking remain under Reservation
- Now Search has to hit back to Reservation to perform inventory searches



# Final dependency graph

- There are still a few synchronized calls, which, for the time being, we will keep as they are
- By applying all these changes, the final dependency diagram will look like the following one



# Prioritizing microservices for migration

- As the next step, we will analyze the **priorities**, and identify the order of migration
  - **Dependency**: services with **less dependency** or **no dependency** at all are **easy** to migrate
  - **Transaction volume**: will have more value from an IT support and maintenance perspective
  - **Resource utilization**: helps the remaining modules to function better
  - **Complexity**: less complex modules are easy to migrate

# Prioritizing microservices for migration

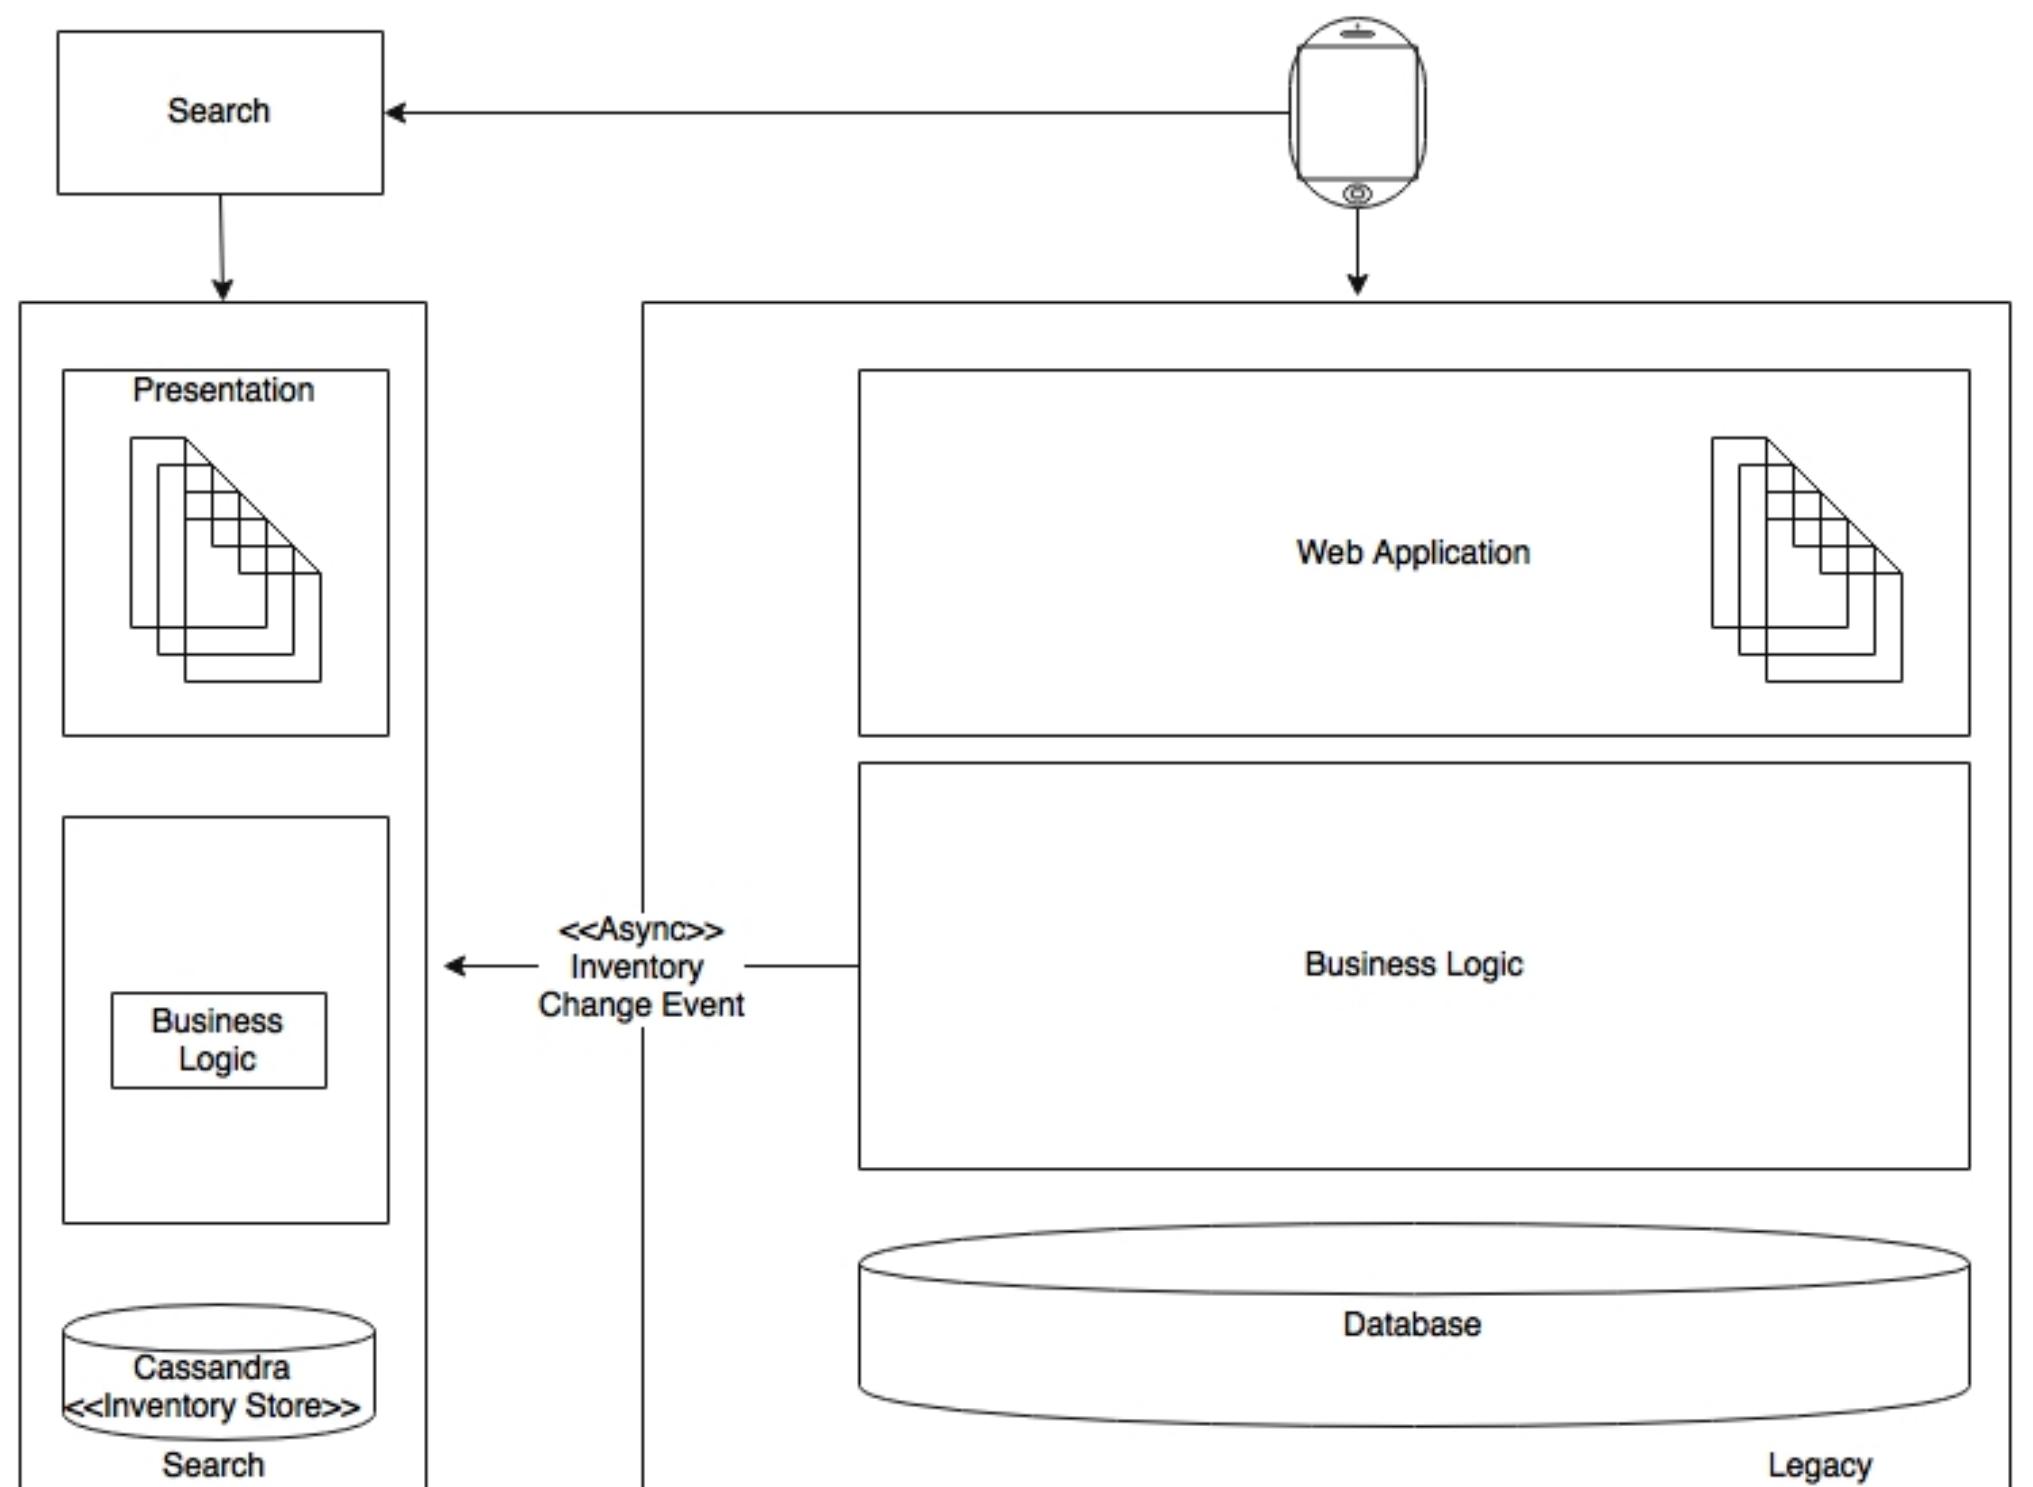
- **Business criticality:** highly critical modules deliver higher business value
- **Velocity of changes:** indicates the number of change requests targeting a function in a short time frame
- **Innovation:** innovations in legacy systems are harder to achieve as compared to applying innovations in the microservices world

# Data synchronization during migration

- During the transition phase, the legacy system and the new microservices will run in parallel
  - synchronize the data between the two systems at the database level by using any data synchronization tool
    - built on the same data store technologies
    - we allow a **backdoor entry**, hence exposing the microservices' internal data store outside

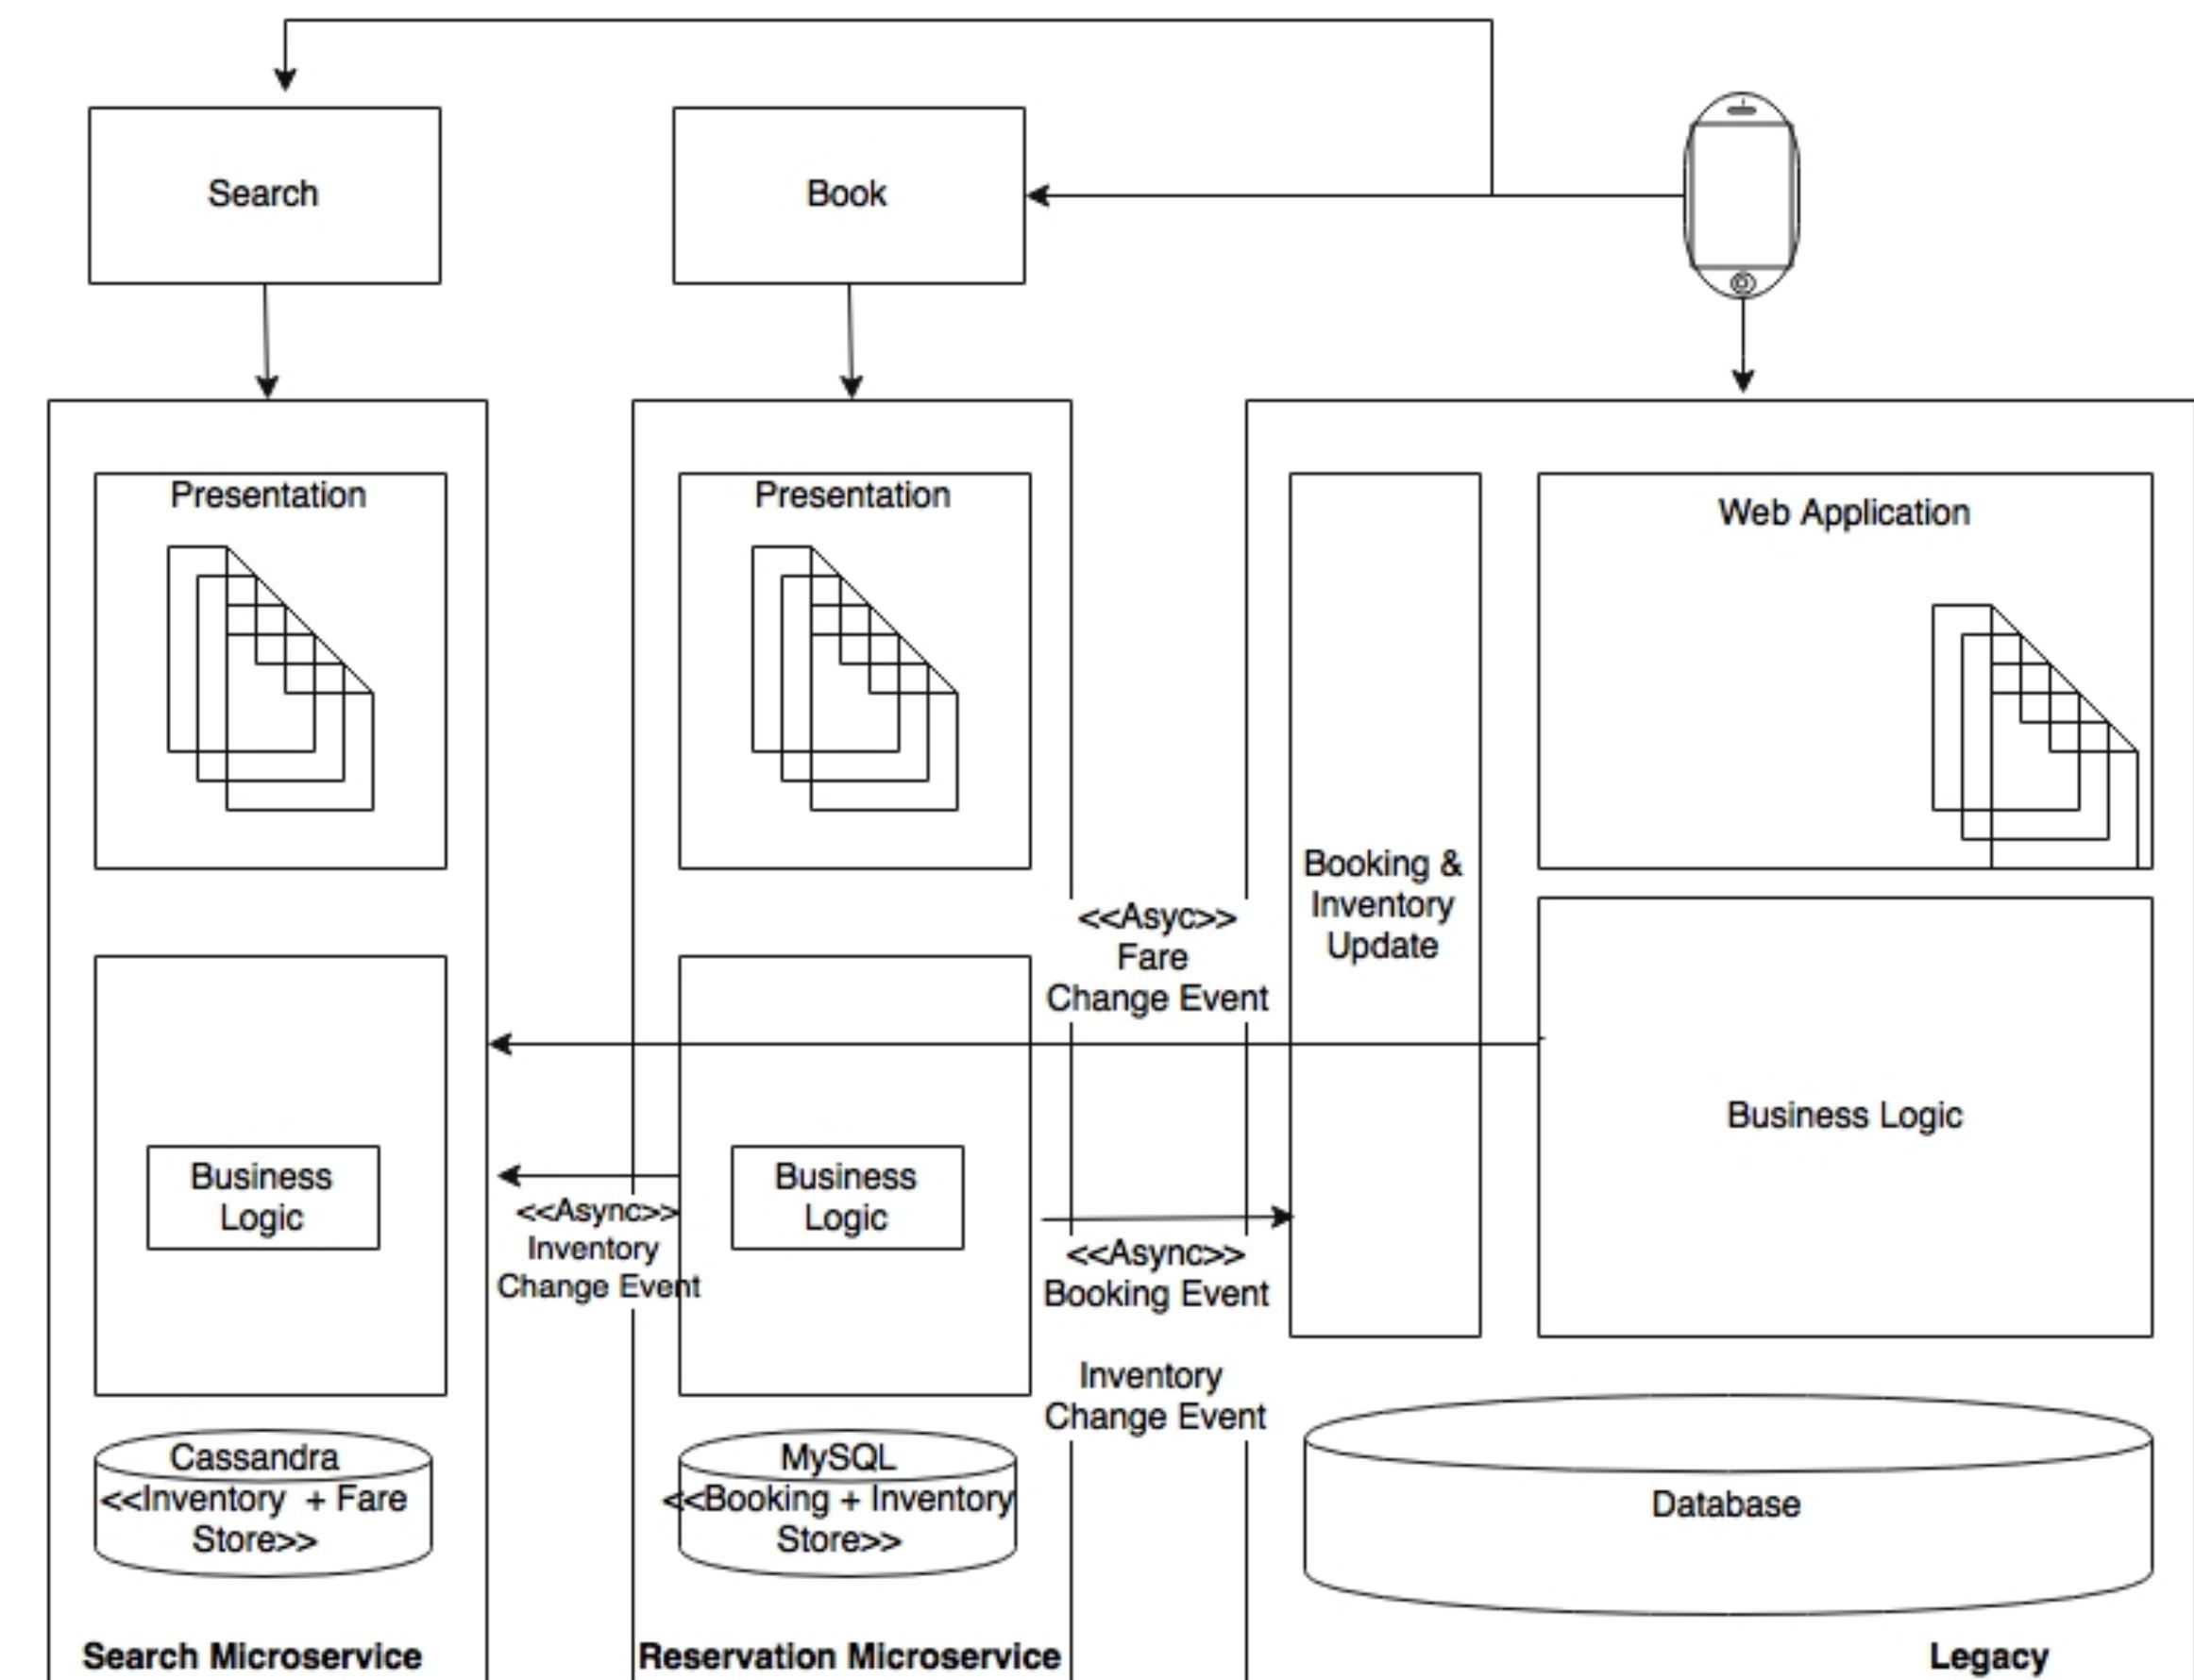
# Data synchronization during migration

- Let us assume that we use a NoSQL database for keeping inventory and fares under the Search service
- In this particular case, all we need is the **legacy system** to supply data to the new service using **asynchronous events**
- The Search service then **accepts** these events, and **stores them locally** into the local NoSQL store



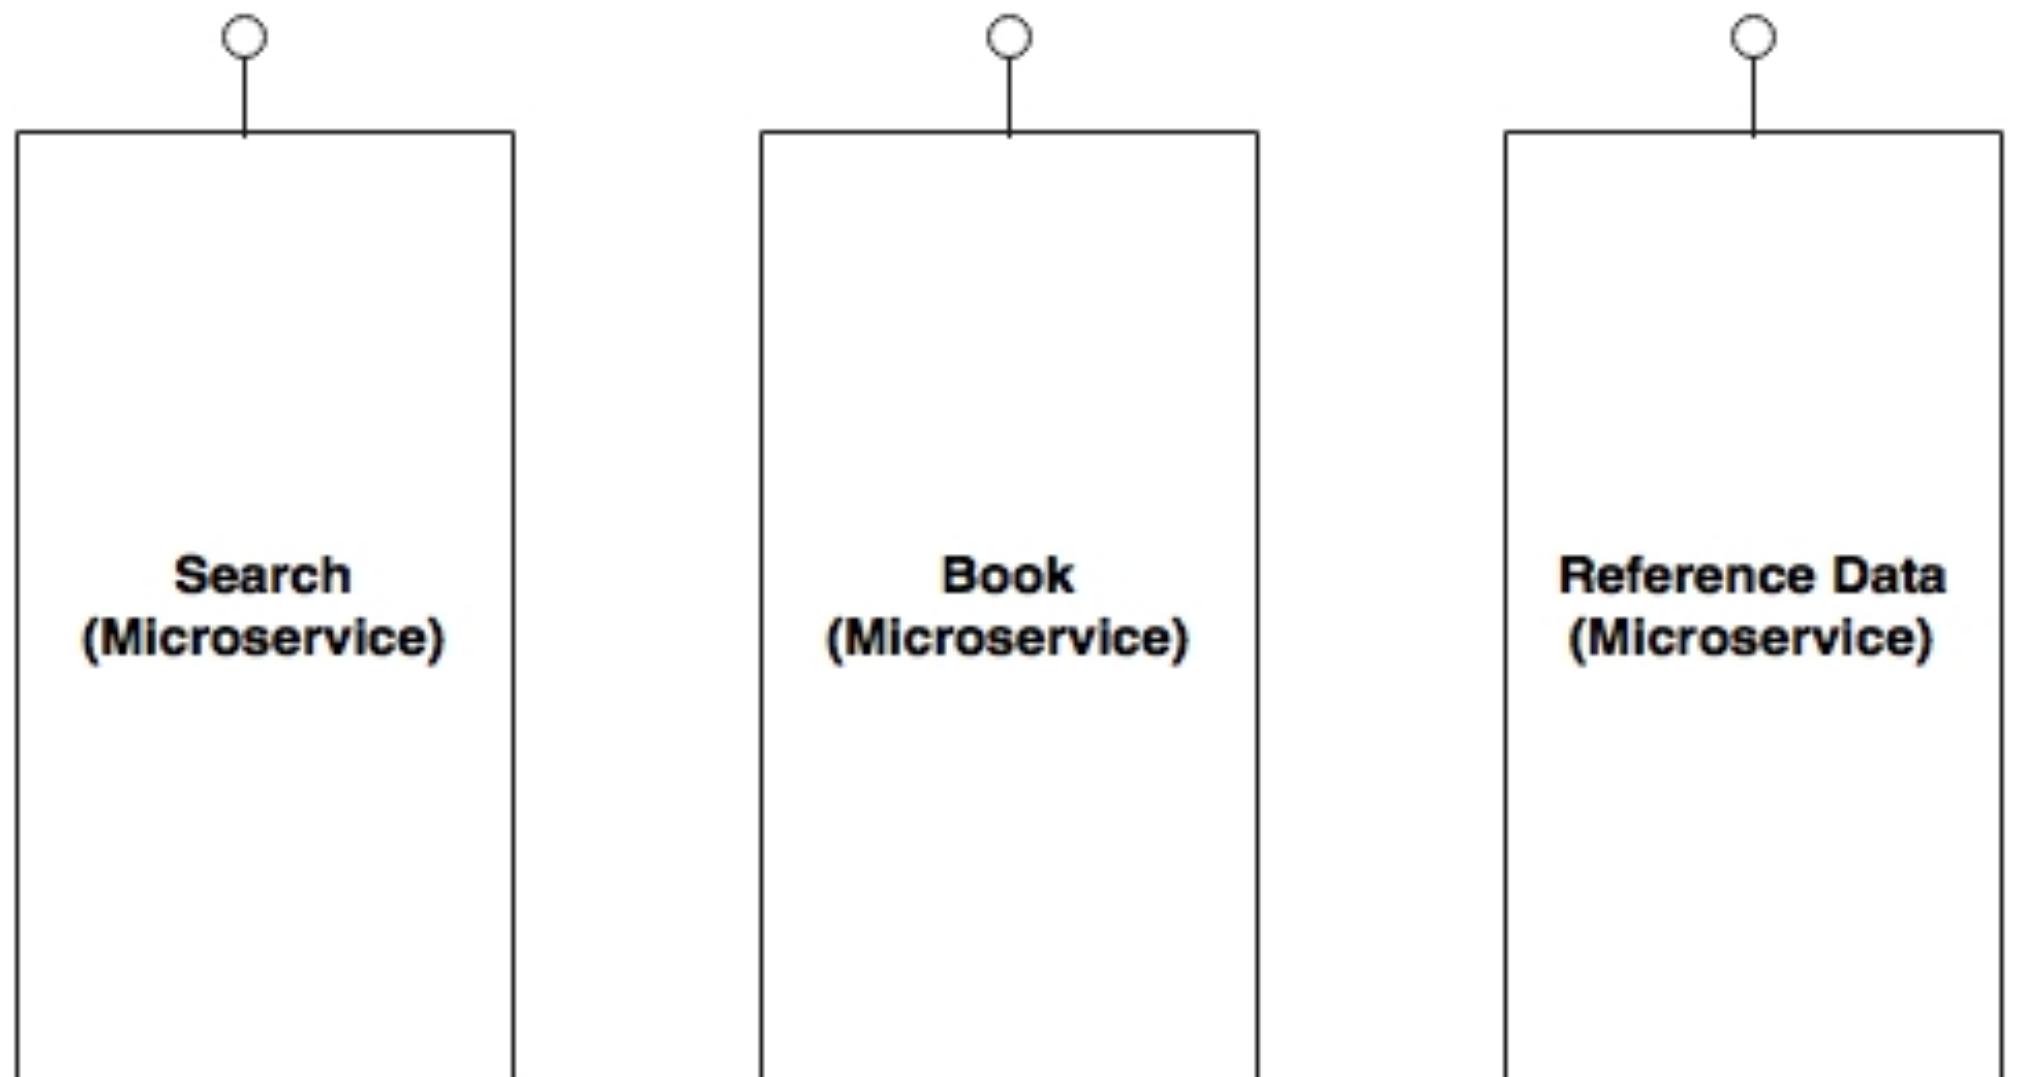
# Data synchronization during migration

- The new Booking microservice **sends** the **inventory change events** to the Search service
- In addition to this, the **legacy application** also has to send the **fare change events** to Search
- Booking will then **store** the new Booking service in its MySQL data store



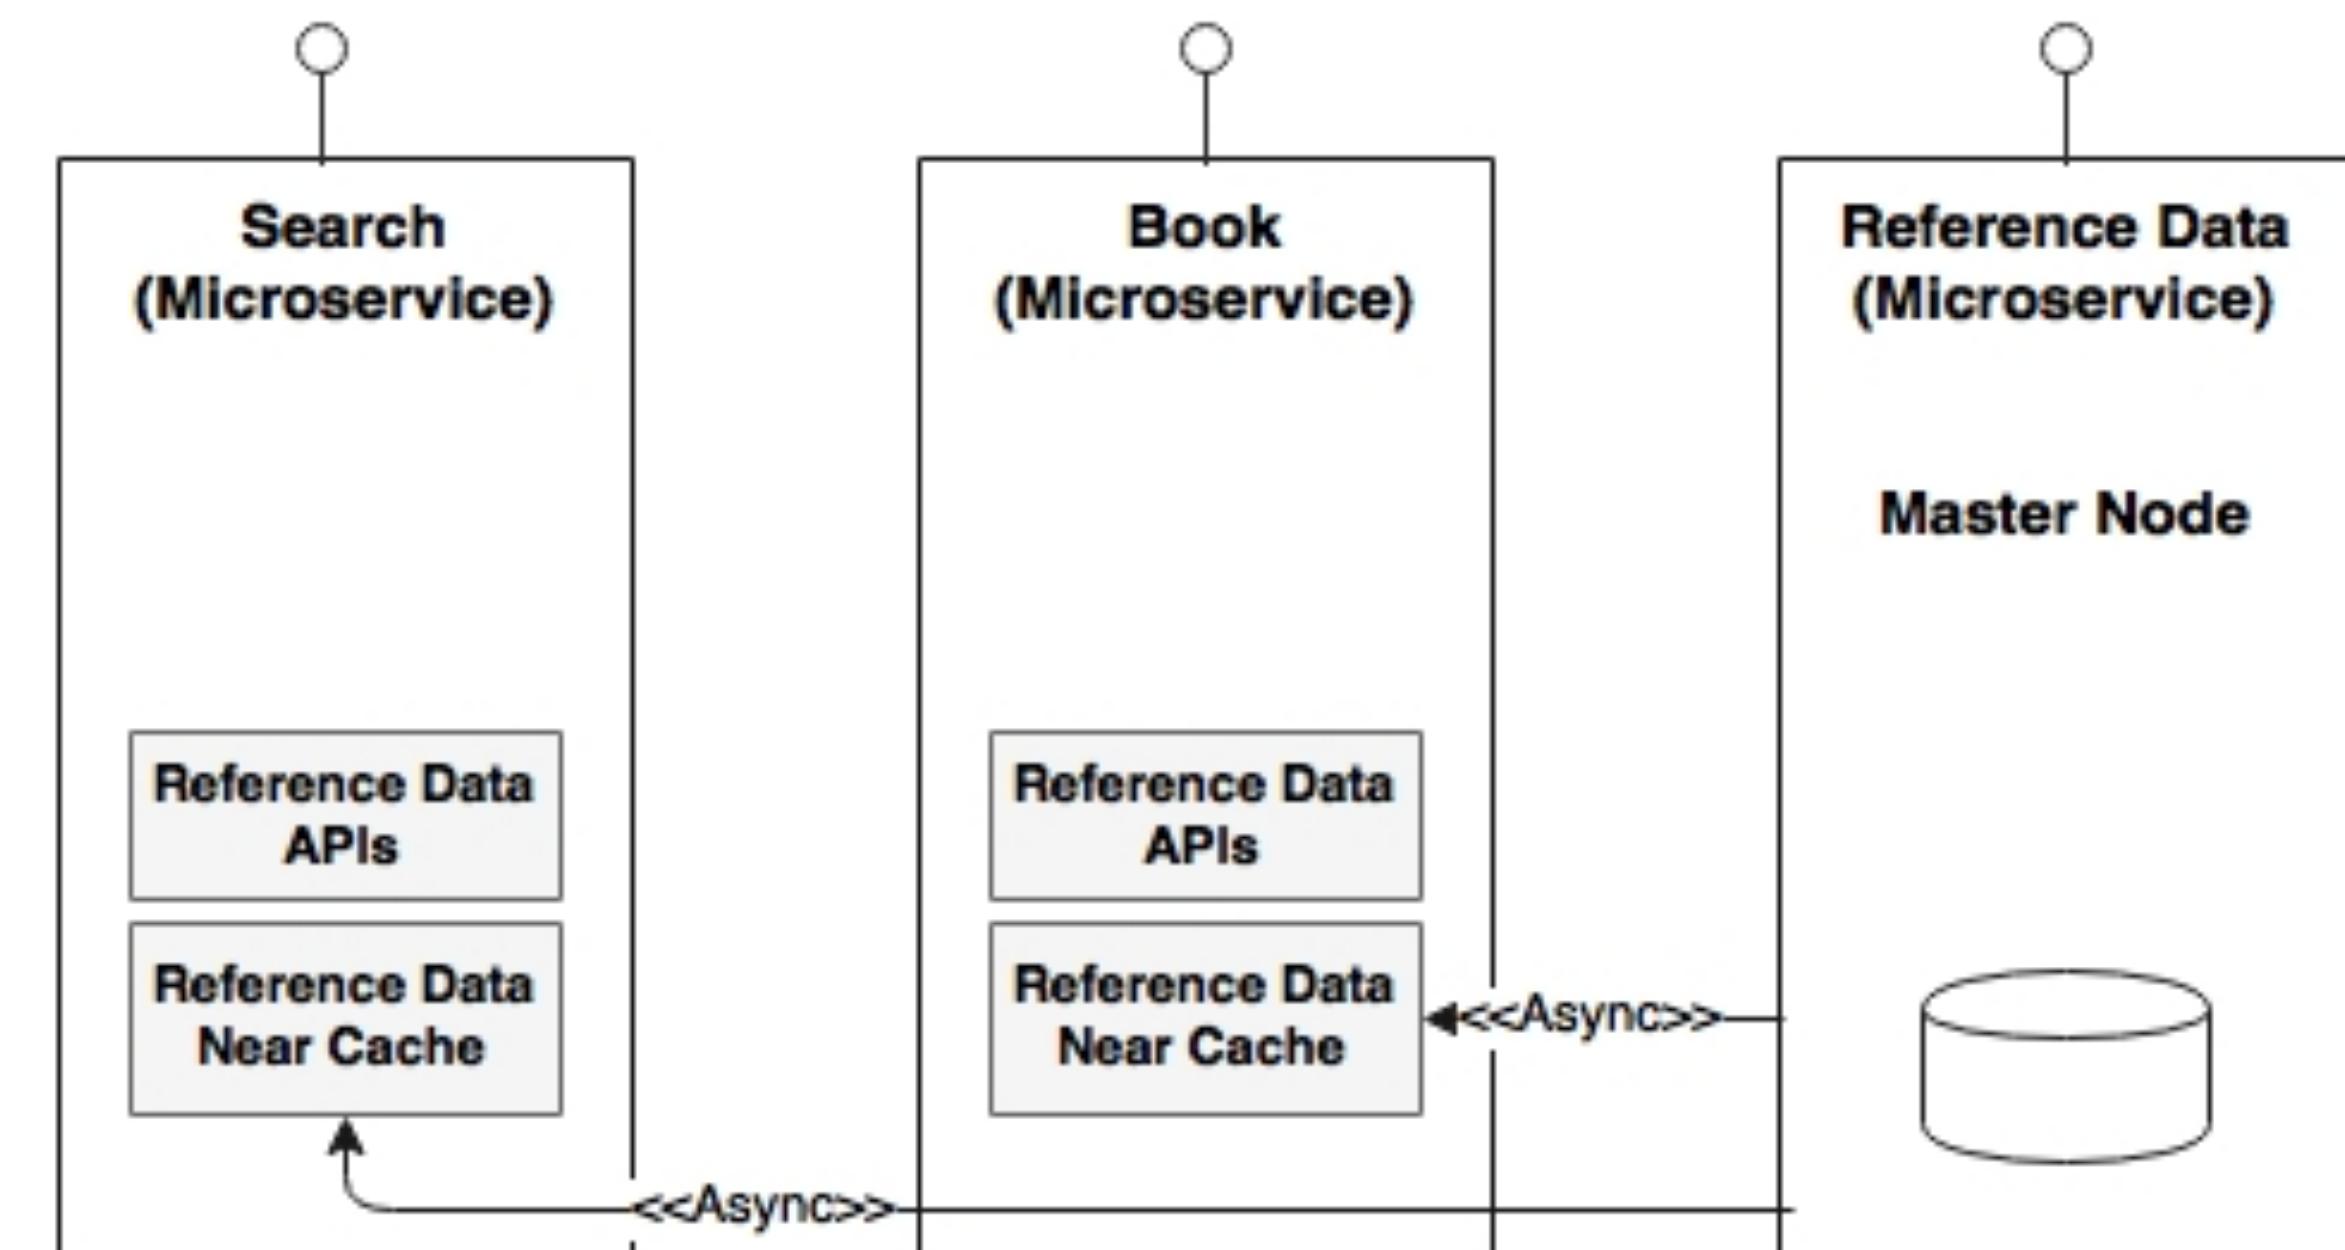
# Managing reference data

- One of the biggest challenges
- A simple approach is to **build the reference data as another microservice itself**
- In this case, whoever needs reference data **should access it through the microservice endpoints**
- This is a well-structured approach, but could lead to **performance issues** as encountered in the original legacy system



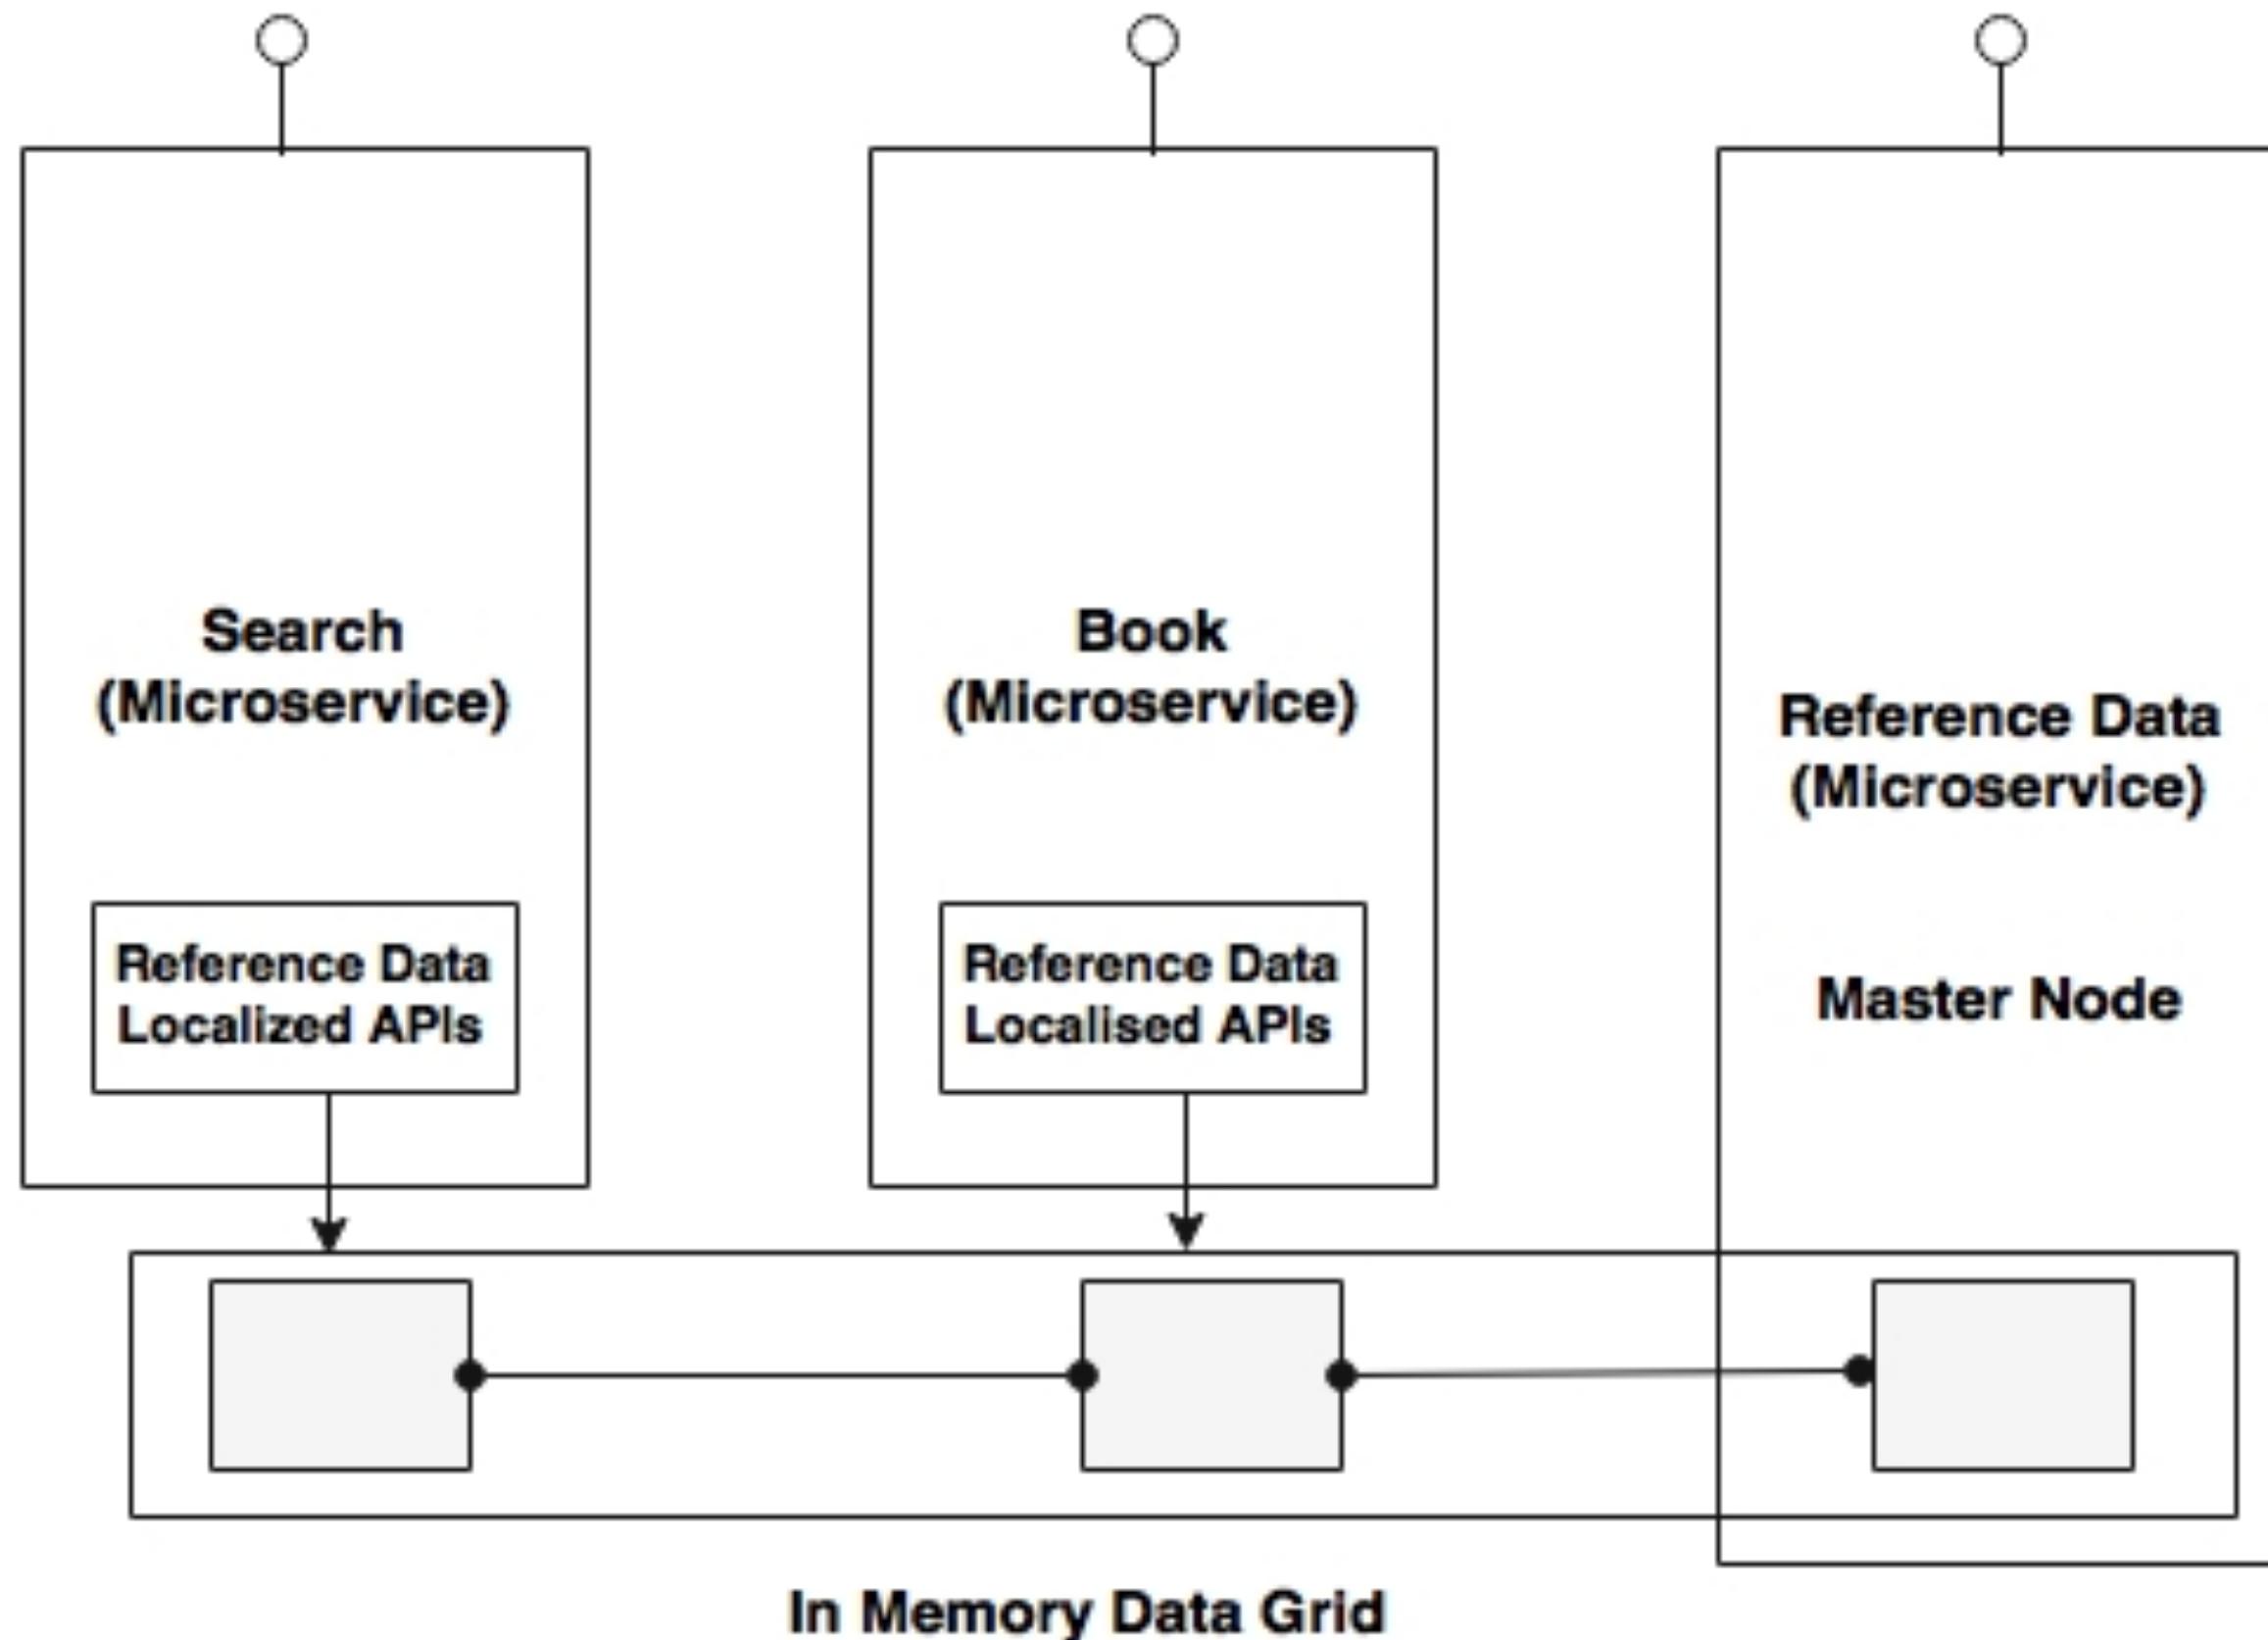
# Managing reference data

- An alternate approach is to have reference data as a microservice service for all the admin and CRUD functions
  - A near cache will then be created under each service to incrementally cache data from the master services
  - A thin reference data access proxy library will be embedded in each of these services
- The reference data access proxy abstracts whether the data is coming from cache or from a remote service.



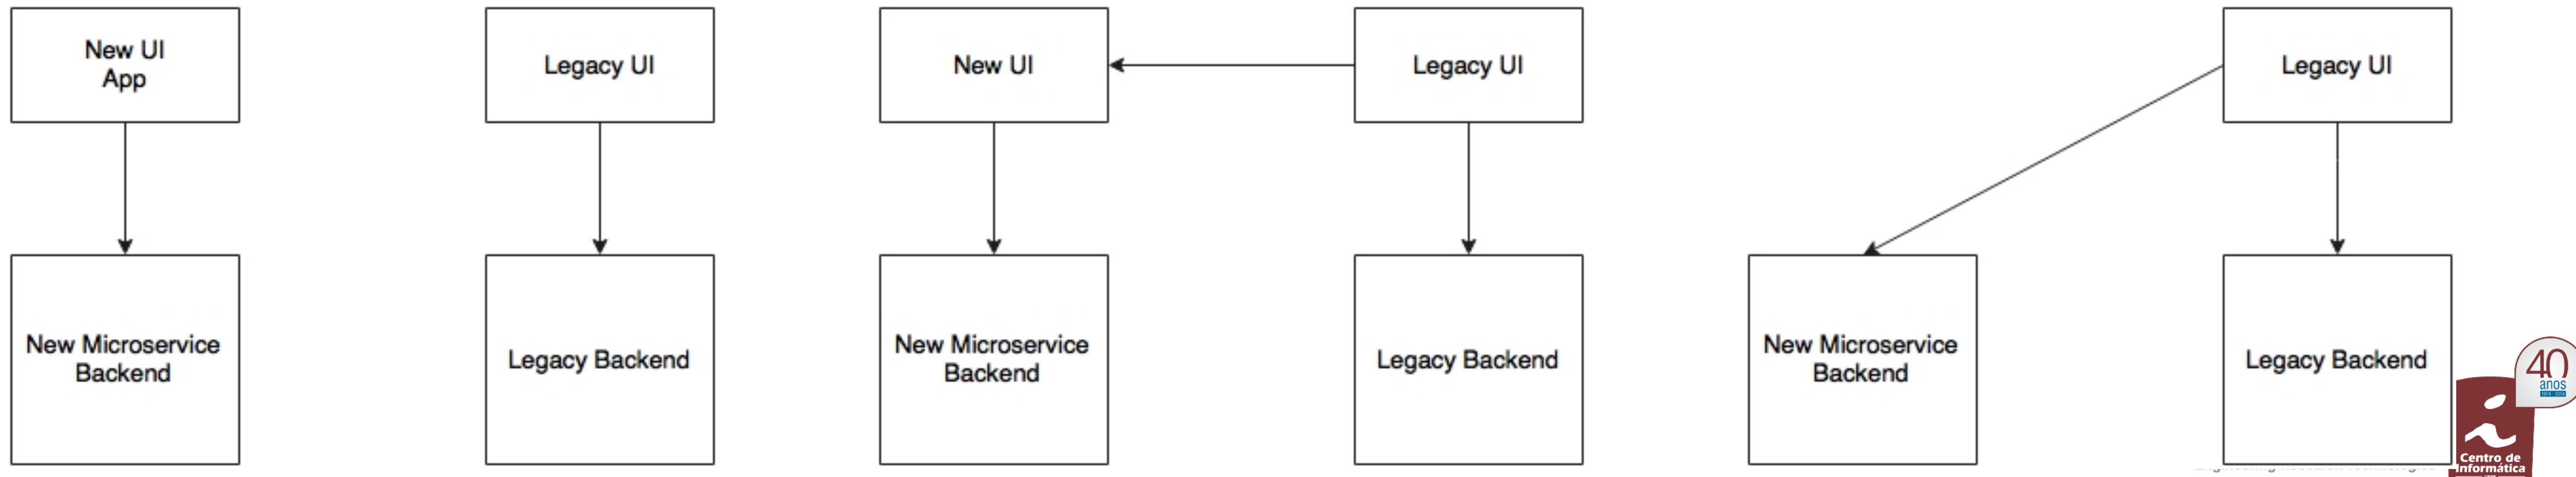
# Managing reference data

- The challenge is to synchronize the data between the master and the slave
- A subscription mechanism is required for those data caches that change frequently
- A better approach is to replace the local cache with an in-memory data grid



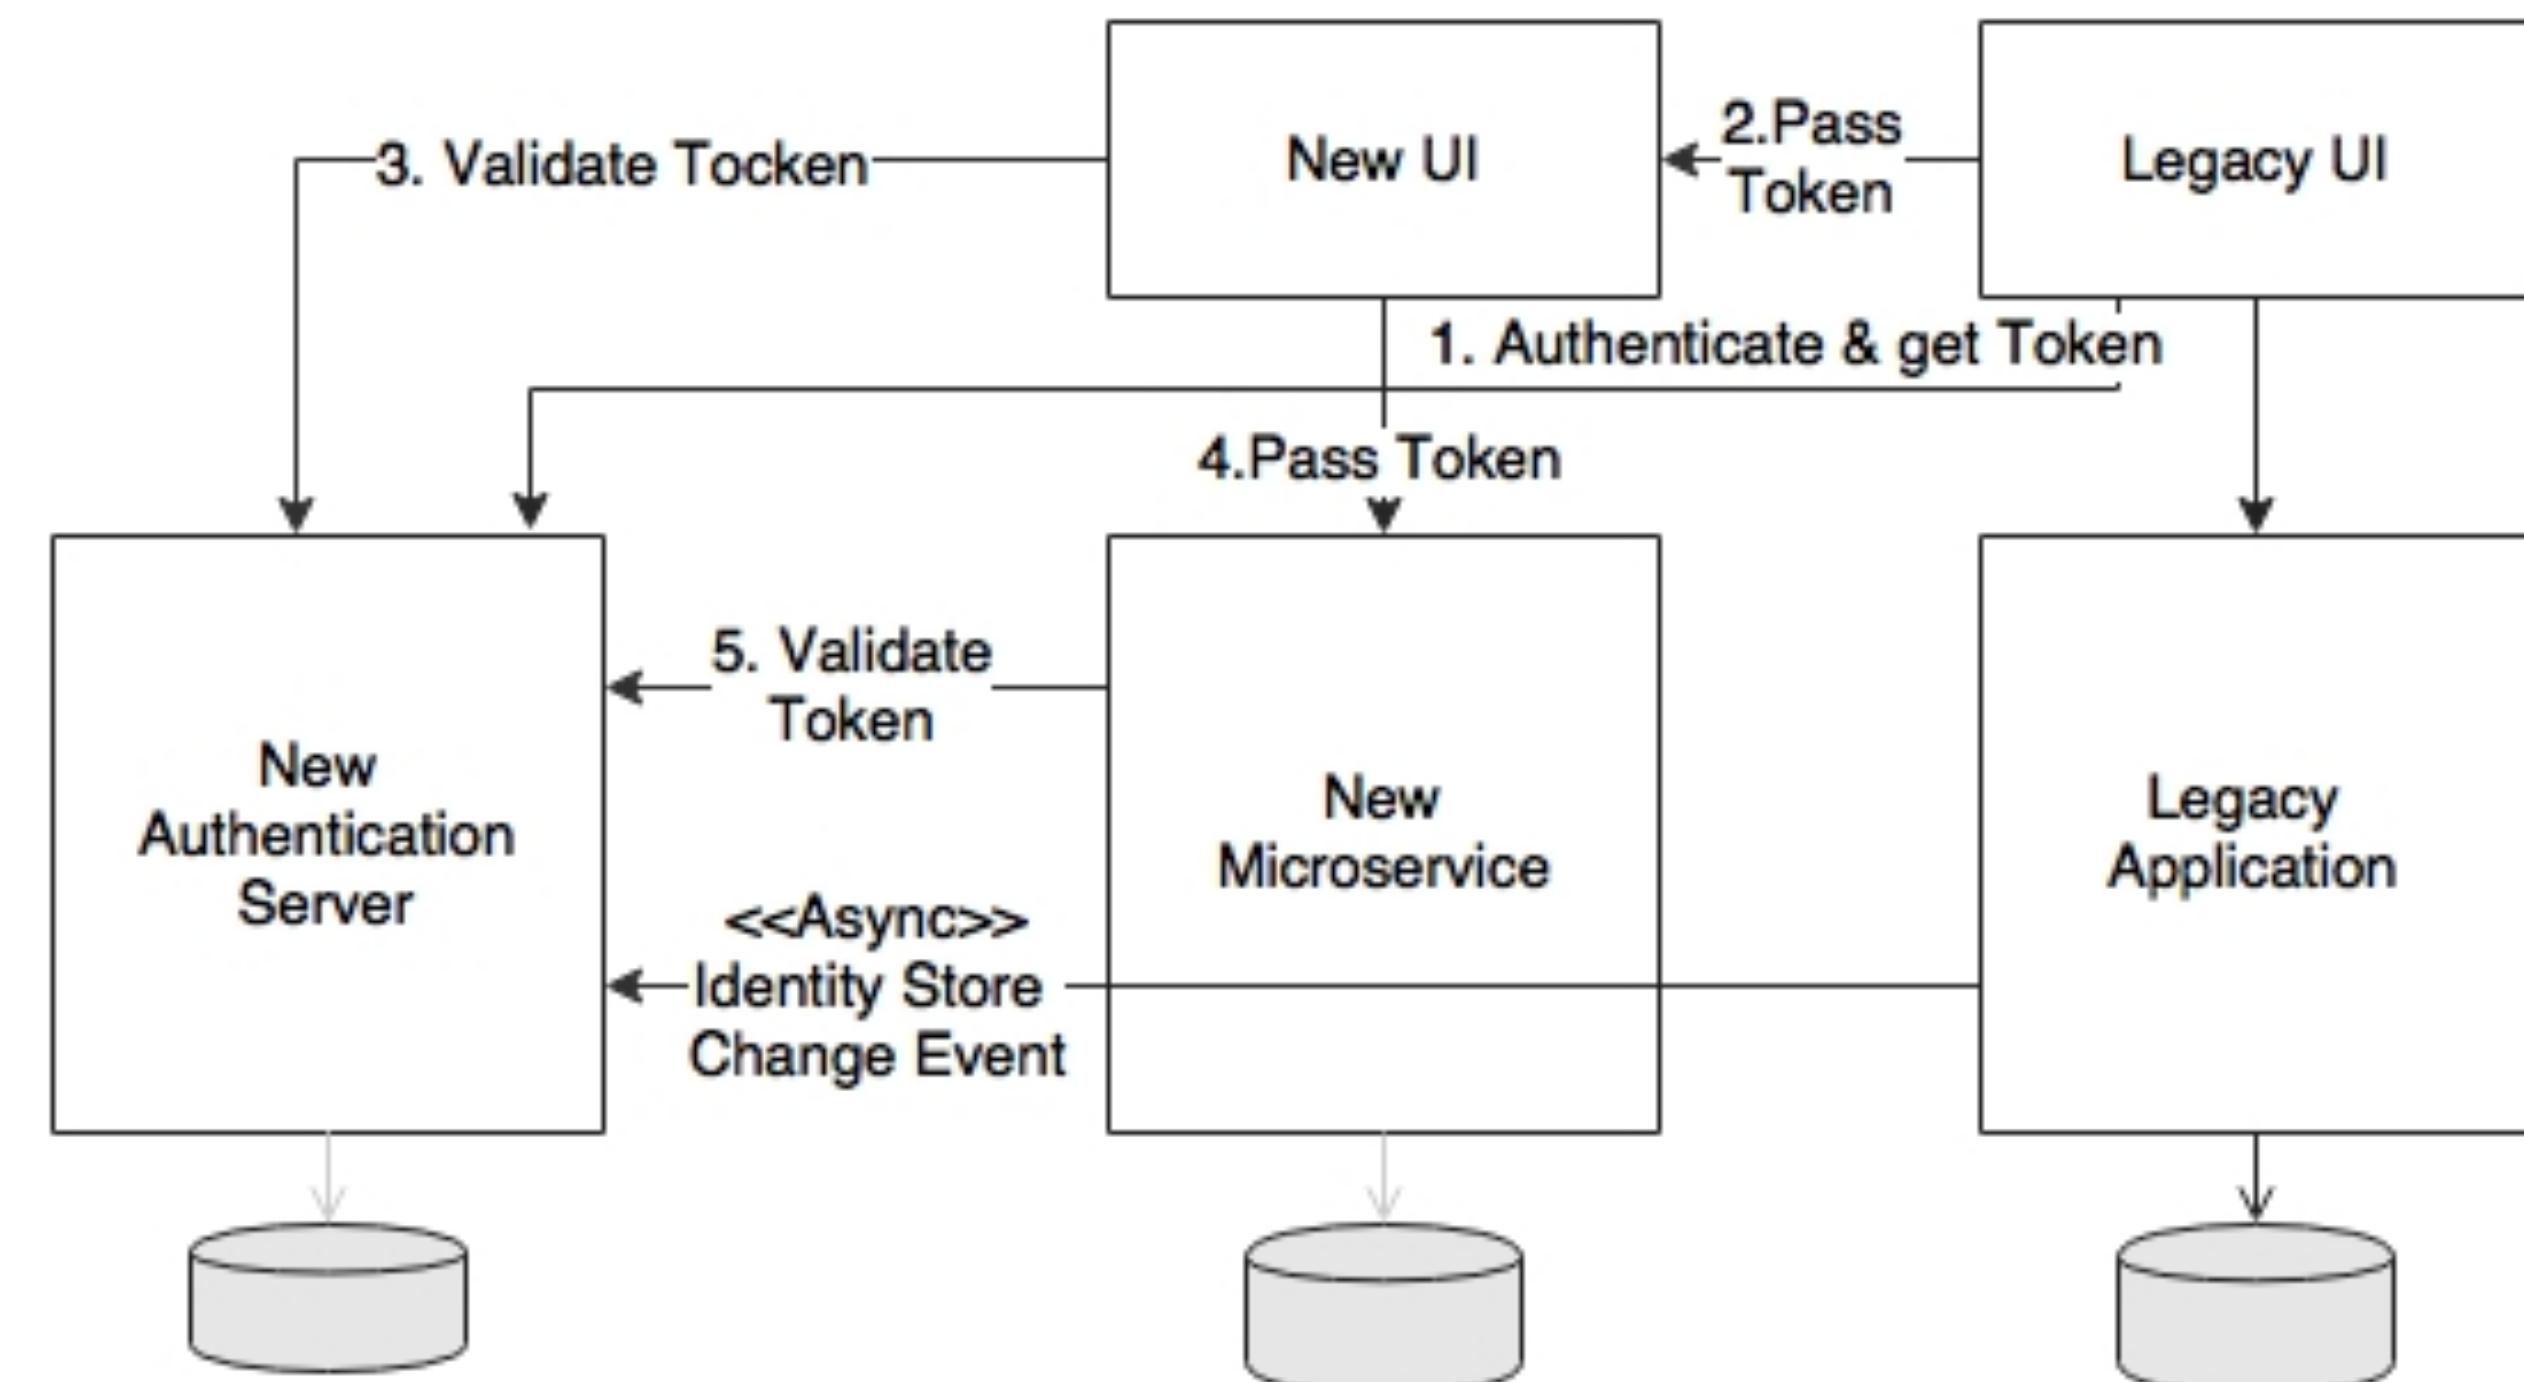
# User interfaces and web applications

- During the transition phase, we have to keep both the old and new user interfaces together.
  - The first approach is to have the old and new user interfaces as separate user applications with no link between them
  - The second approach is to use the legacy user interface as the primary application, and then transfer page controls to the new user interfaces when the user requests pages of the new application
  - The third approach is to integrate the existing legacy user interface directly to the new microservices backend



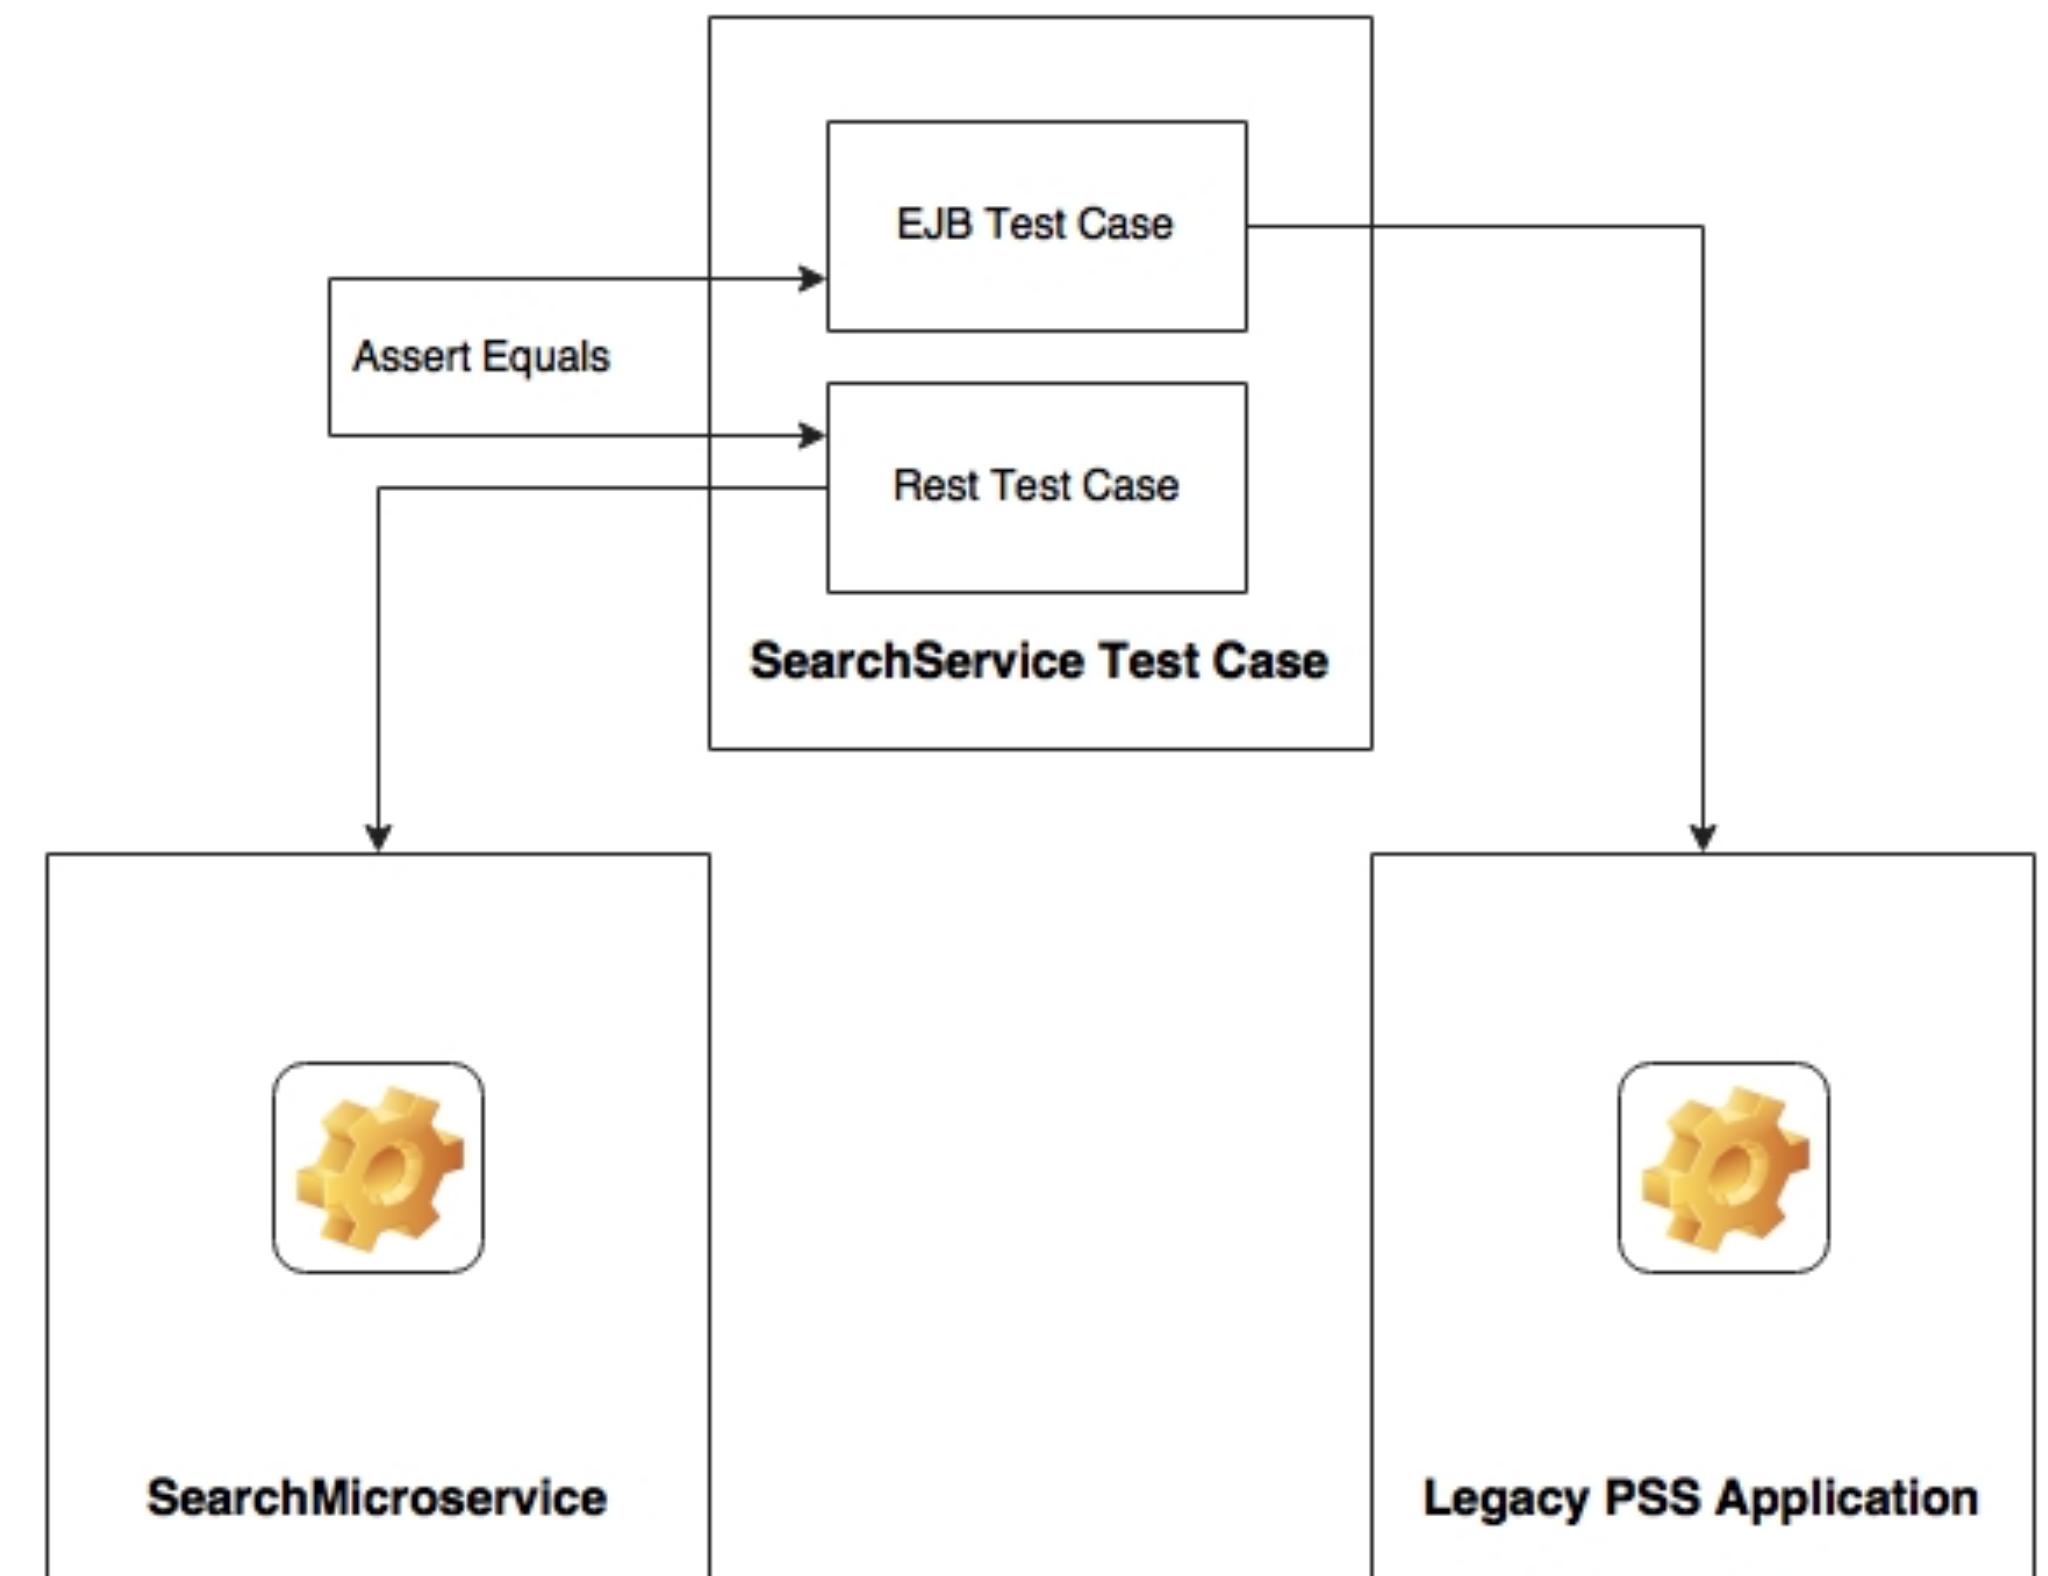
# Session handling and security

- Assume that the new services are written based on **Spring Security** with a **token-based authorization strategy**, whereas the old application uses a **custom-built authentication** with its local identity store
- The simplest approach, as shown in the preceding diagram, is to build a new identity store with an authentication service as a new microservice using Spring Security.
  - This will be used for all our future resource and service protections, for all microservices.



# Test strategy

- One important question to answer from a testing point of view is how can we ensure that all functions work in the same way as before the migration?
- Integration test cases should be written for the services that are getting migrated before the migration or refactoring



# Building ecosystem capabilities

- Before we embark on actual migration, we have to **build all of the microservice's capabilities** mentioned under the **capability model**
- In addition to these capabilities, **certain application functions** are also **required** to be built upfront
  - Reference data, security and SSO, and Customer and Notification
  - Data warehouse or a data lake is also required as a prerequisite
- An effective approach is to build these capabilities in an **incremental fashion**, **delaying development** until it is really required

# Migrate modules only if required

- It is important to understand that **it is not necessary to migrate all modules** to the new microservices architecture, **unless it is really required**
  - A major reason is that these migrations incur **cost**

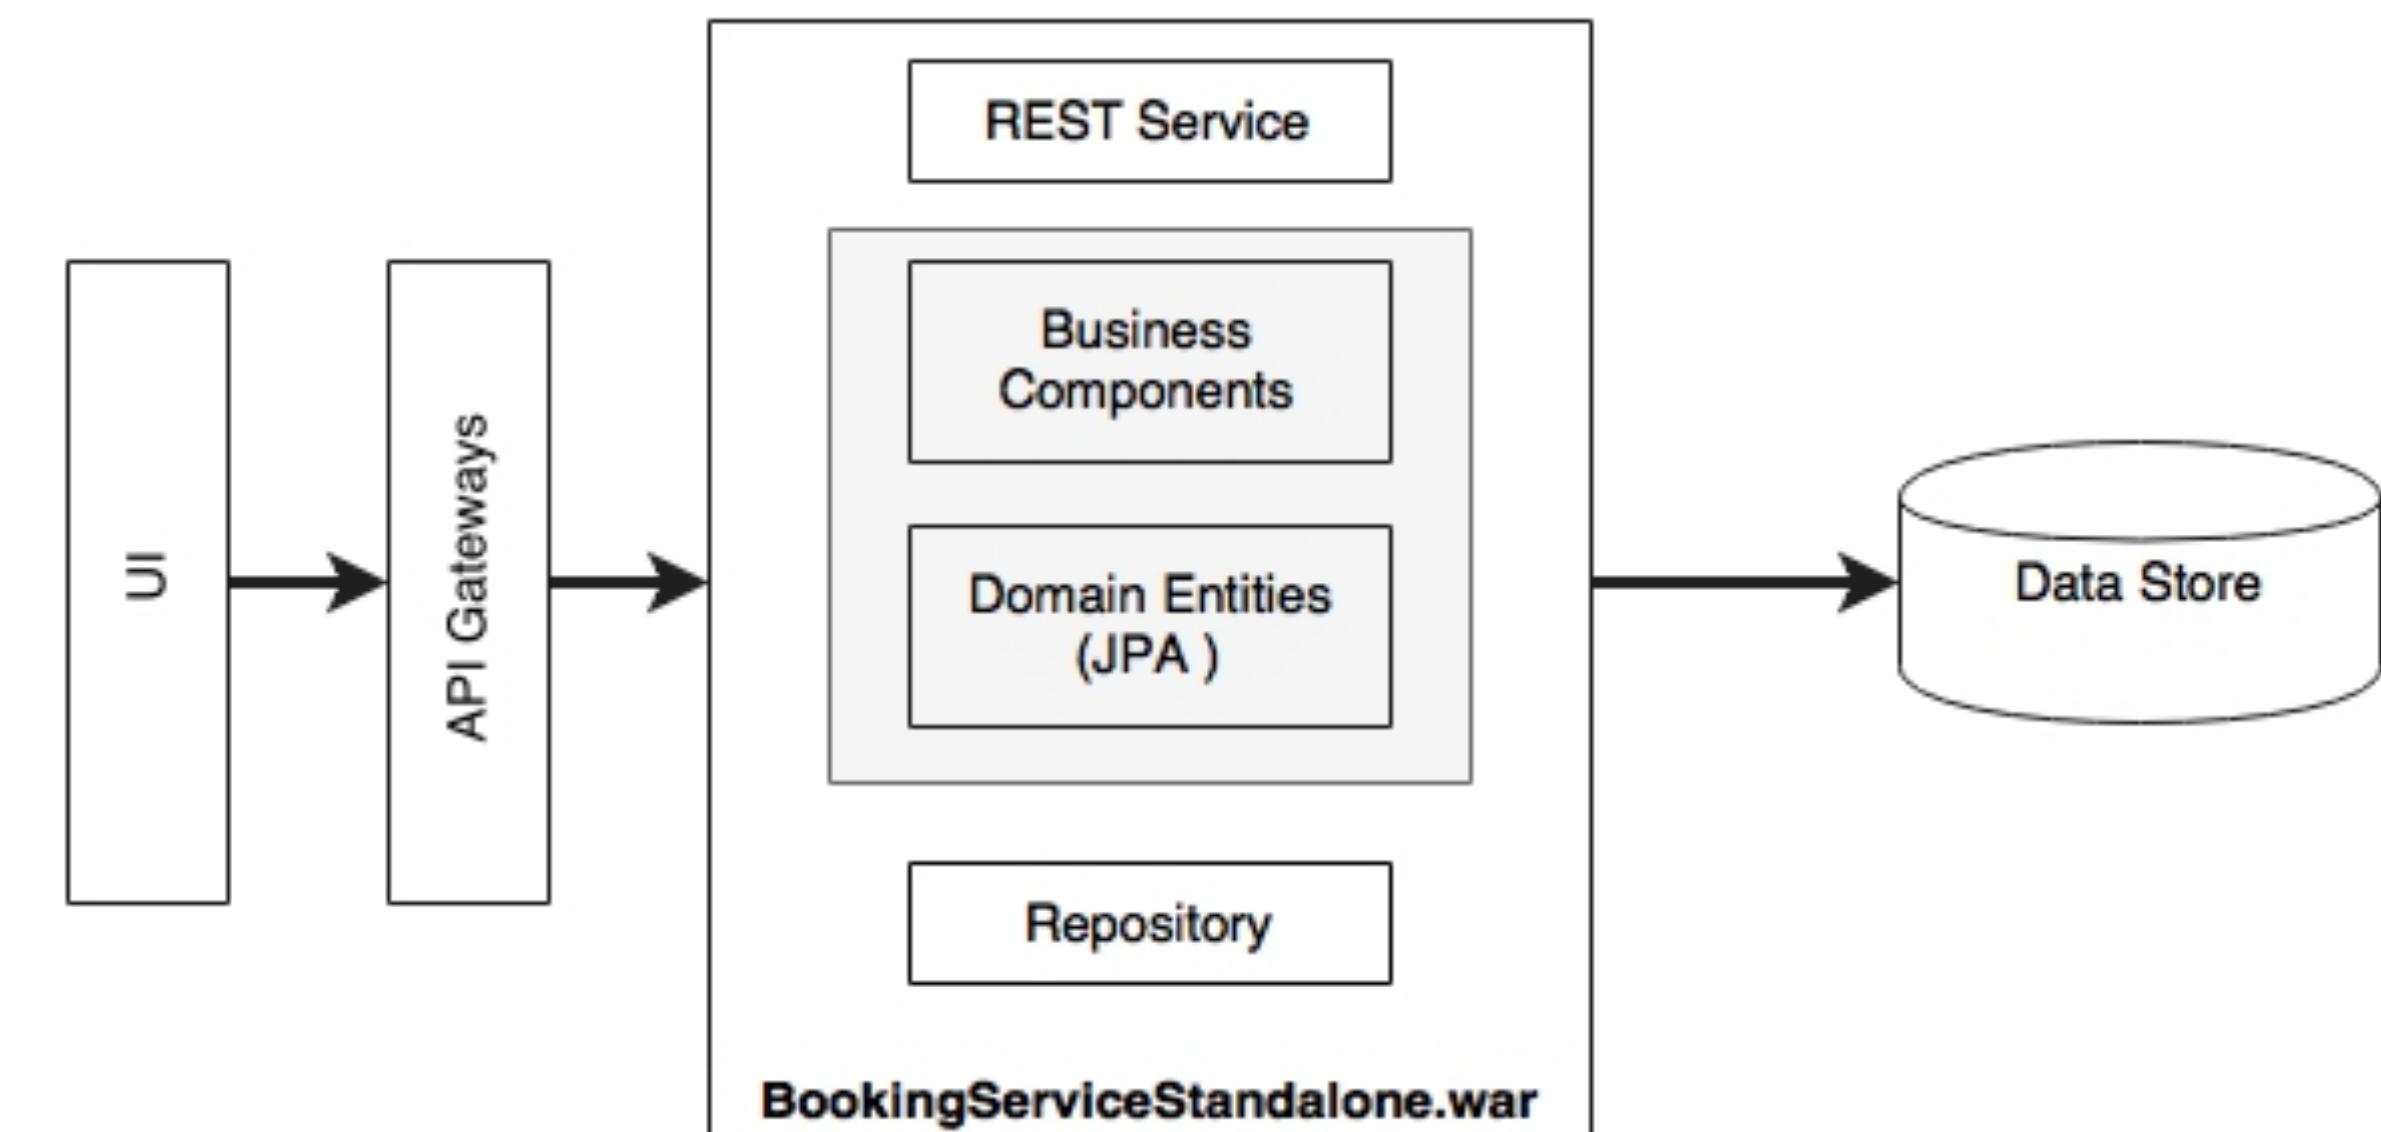
# Target architecture

- Each block in the diagram represents a microservice
- The shaded boxes are core microservices, and the others are supporting microservices
- The diagram also shows the internal capabilities of each microservice



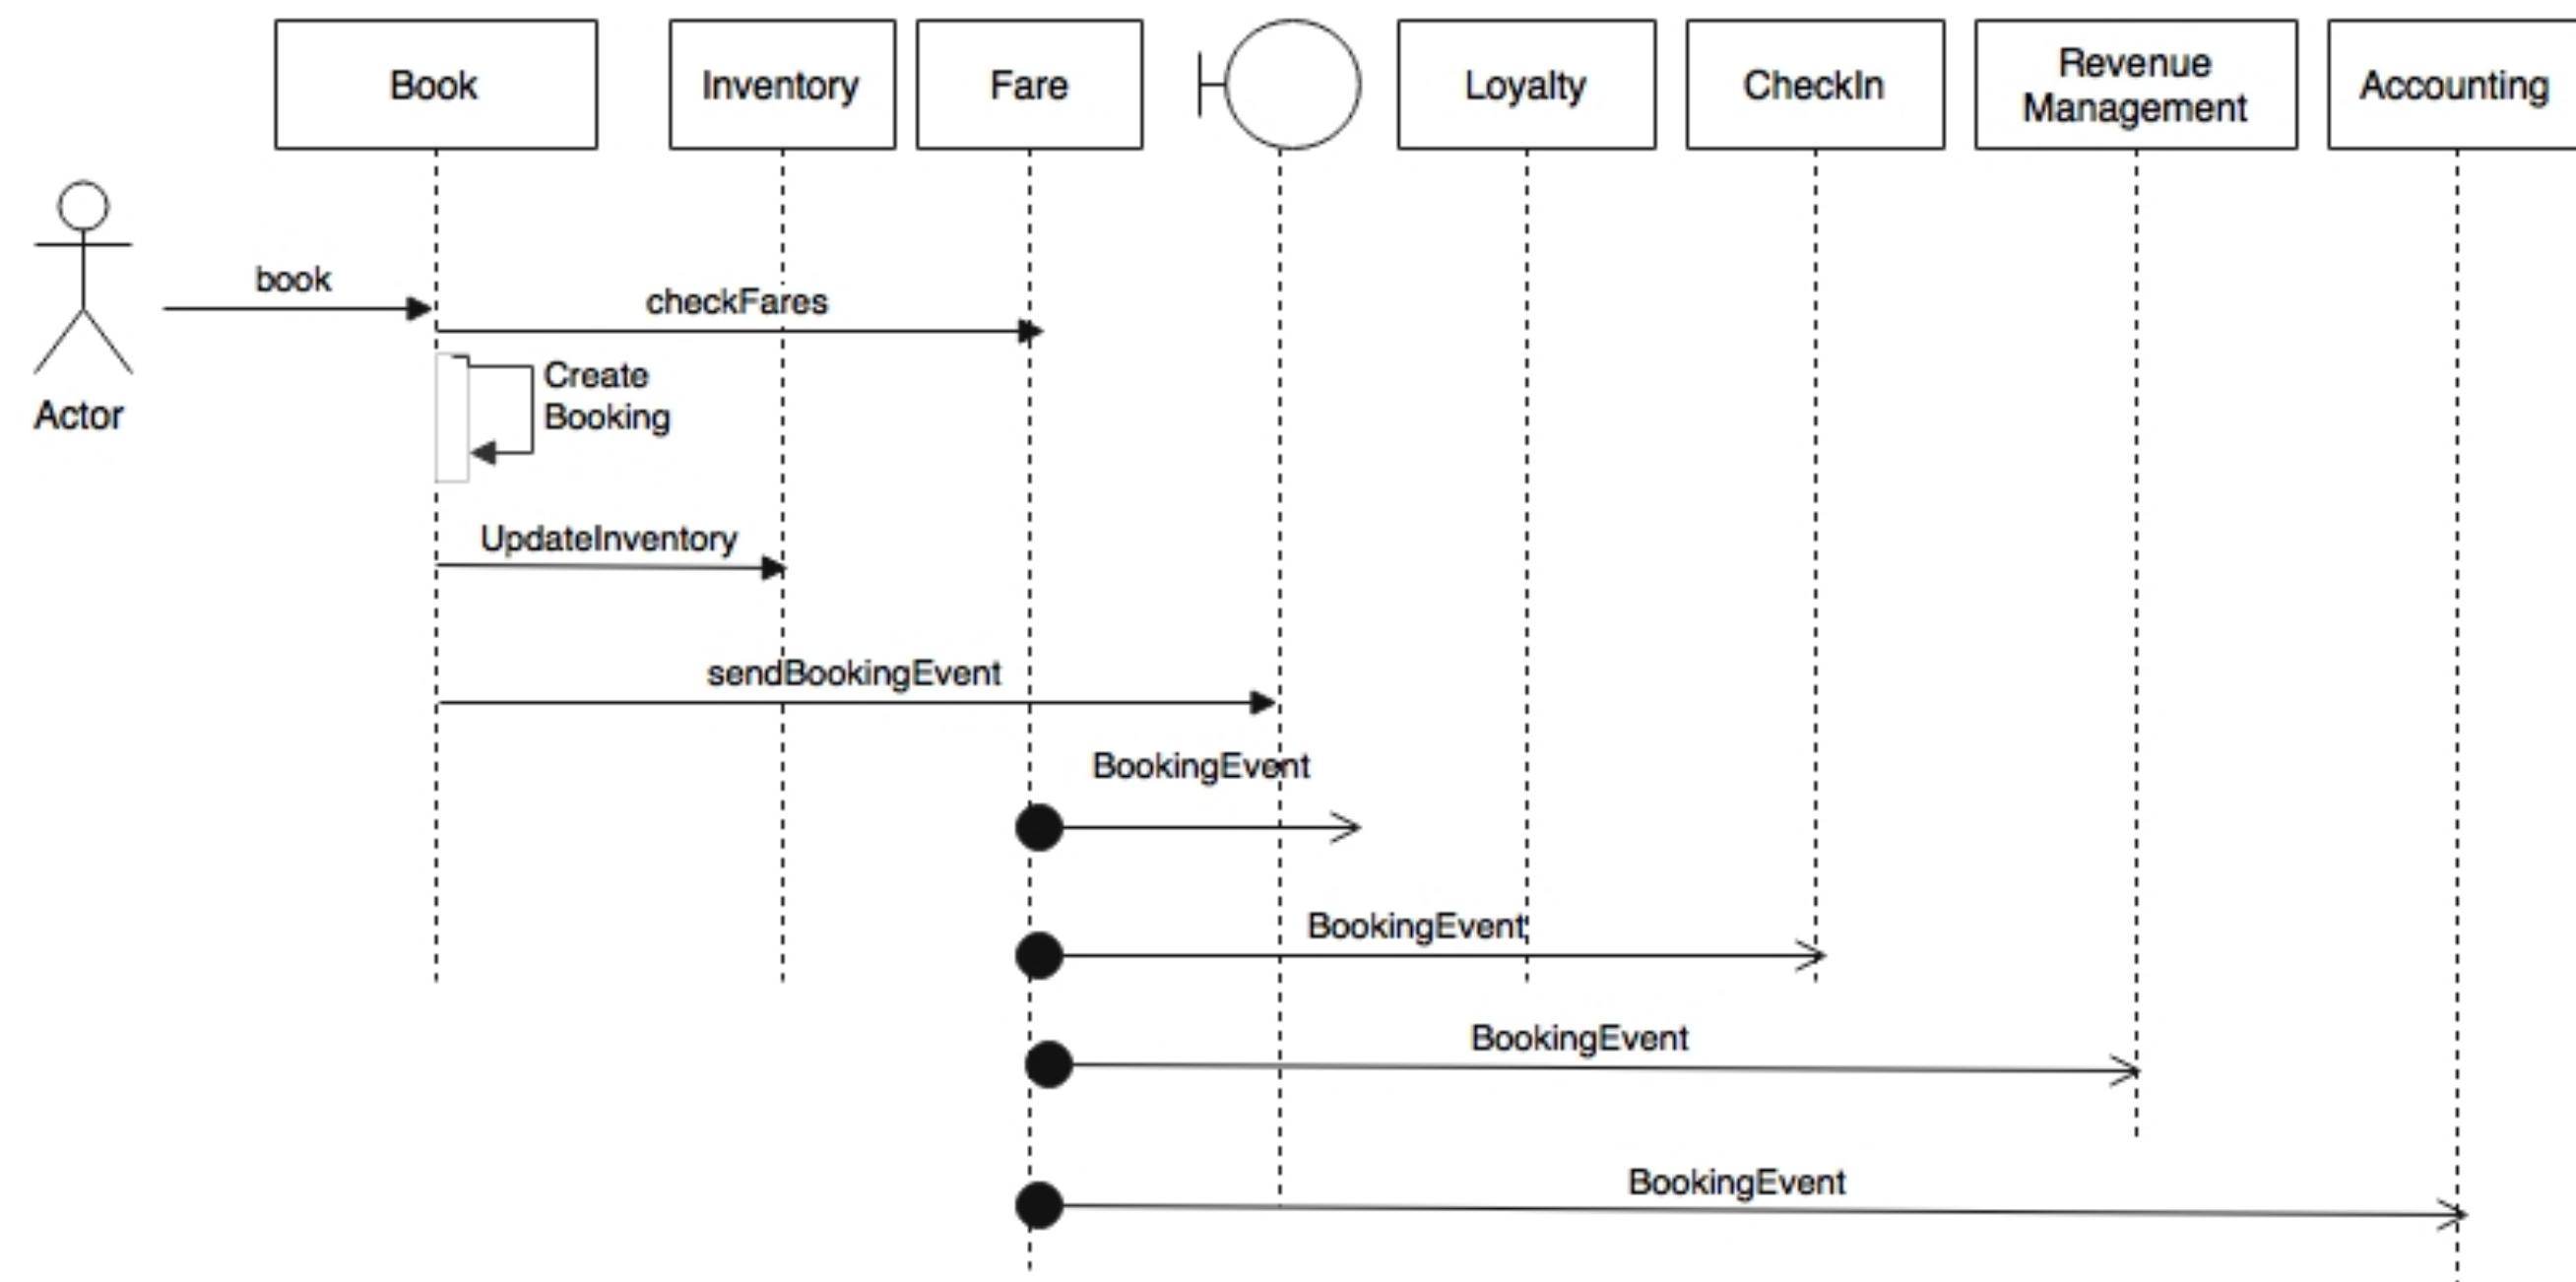
# Internal layering of microservices

- There is no standard to be followed for the internal architecture of a microservice
- The rule of thumb is to abstract realizations behind simple service endpoints



# Orchestrating microservices

- The brain is still inside the Booking service in the form of one or more booking business components
- Internally, business components orchestrate private APIs exposed by other business components or even external services
- The booking service internally calls to update the inventory of its own component other than calling the Fare service



# Integration with other systems

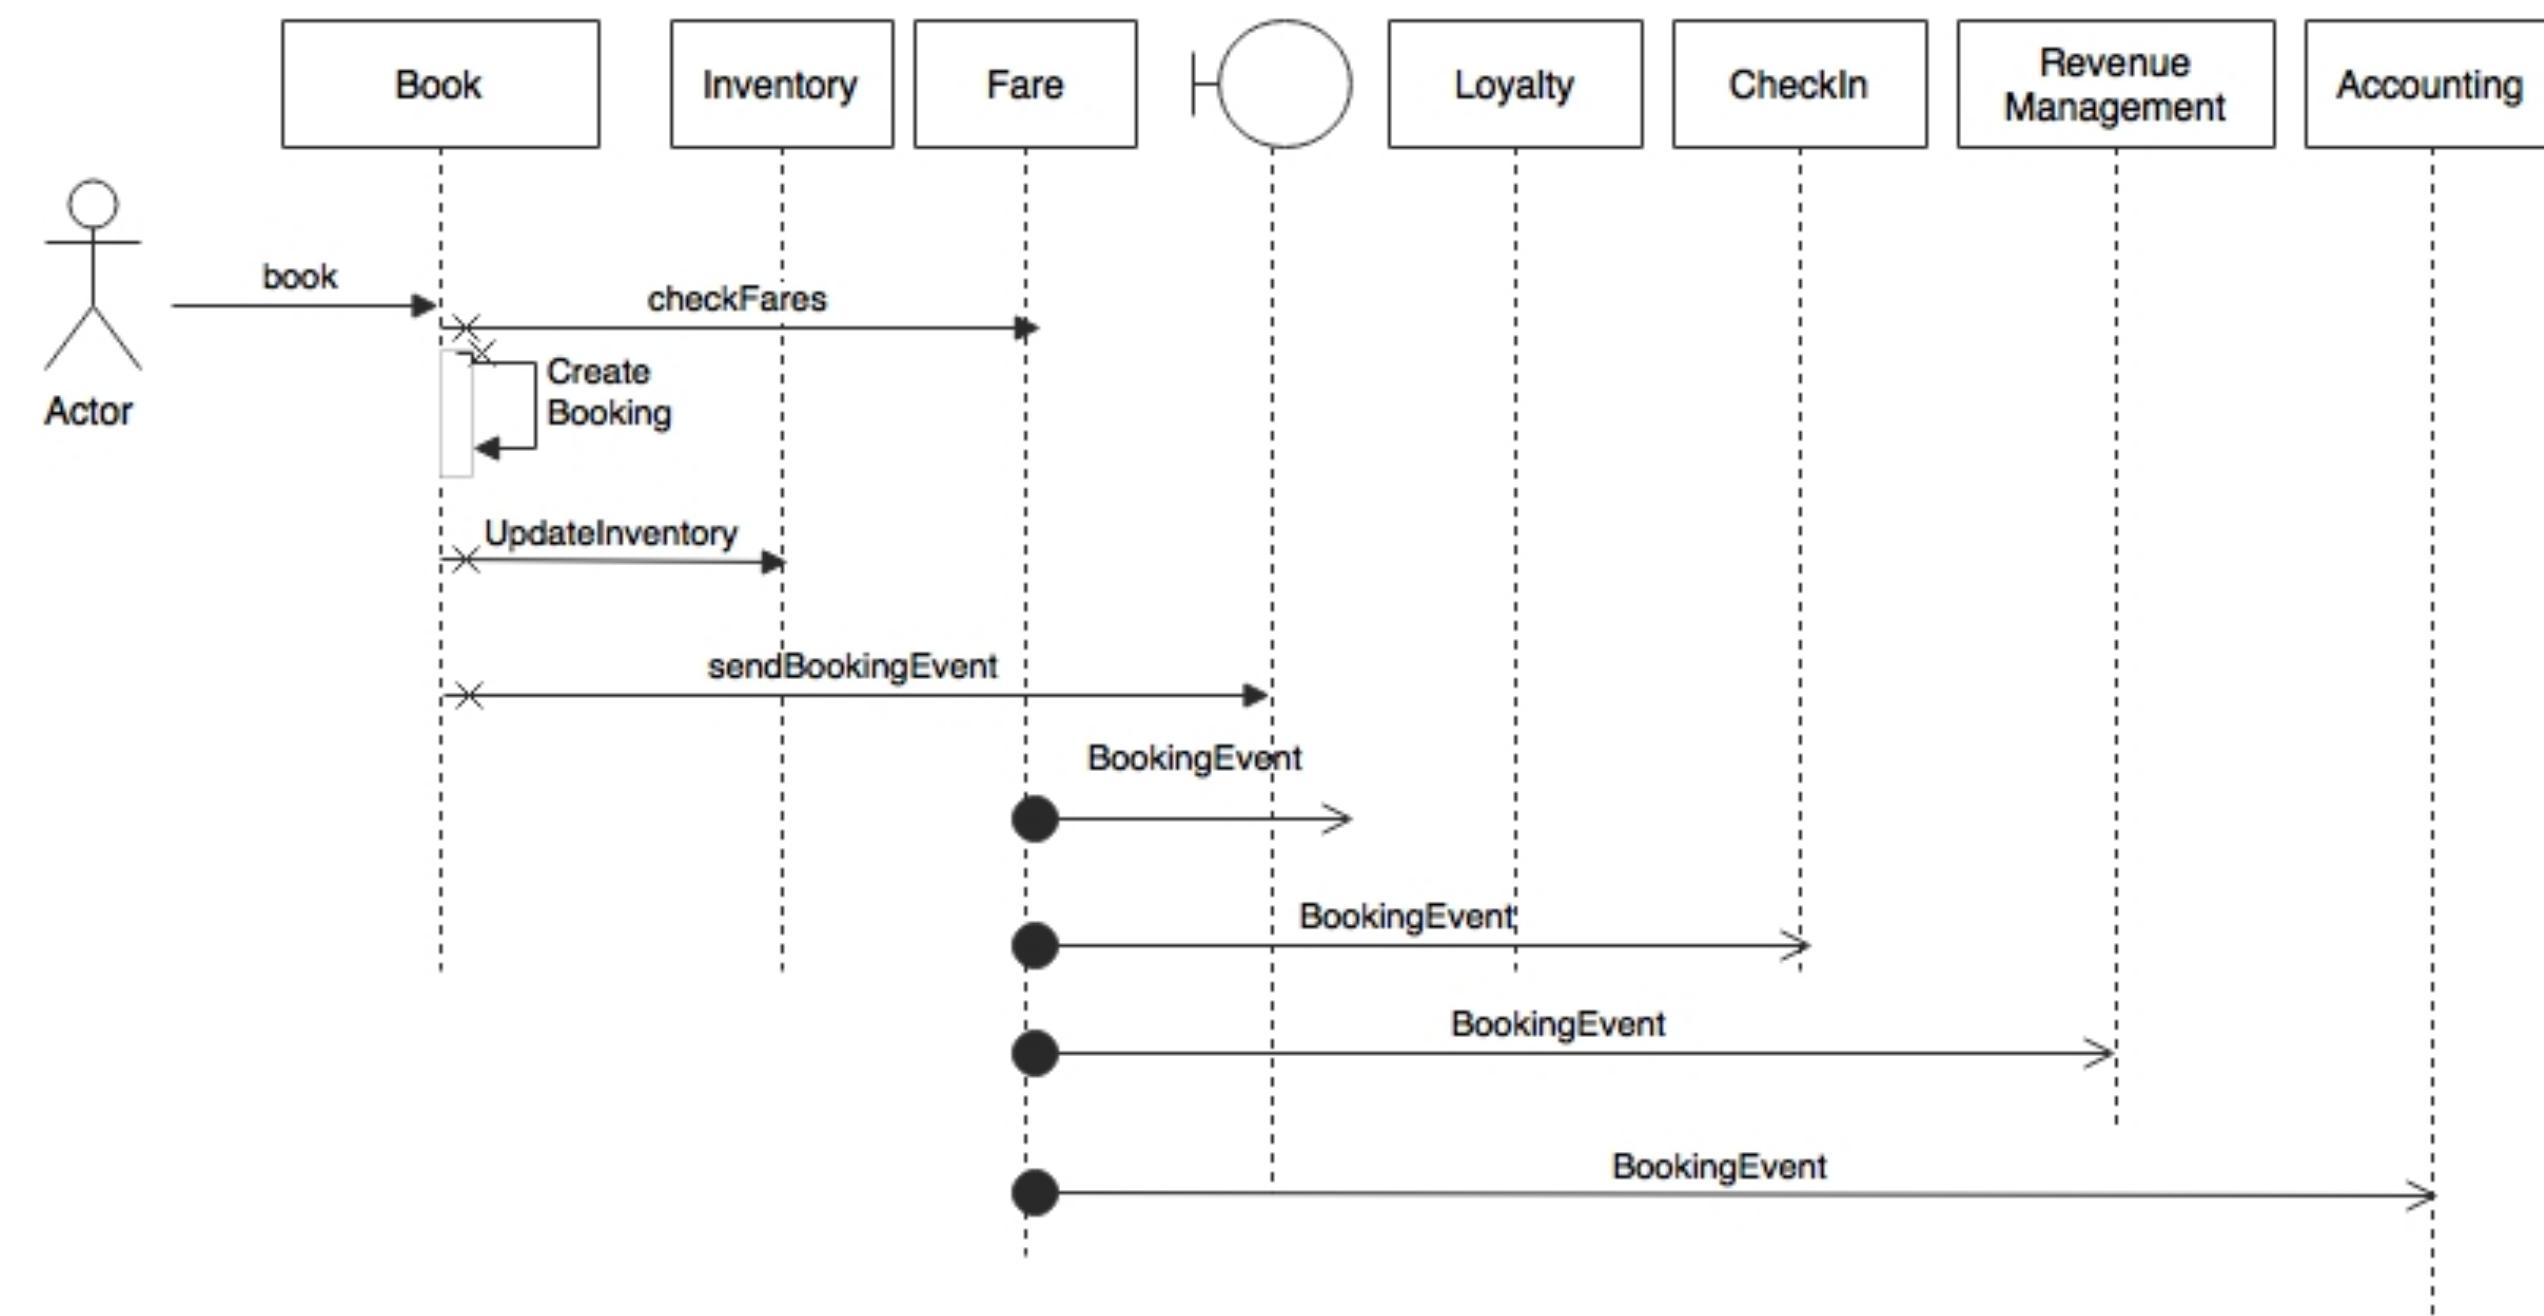
- In the microservices world, we use an API gateway or a reliable message bus for integrating with other non-microservices
- Let us assume that there is another system in BrownField that needs booking data
  - Unfortunately, the system **is not capable of subscribing to the booking events** that the Booking microservice publishes
  - An **Enterprise Application integration (EAI)** solution could be employed, which listens to our booking events, and then uses a native adaptor to update the database

# Managing shared libraries

- Certain business logic is used in more than one microservice
- Search and Reservation, in this case, use inventory rules
- In such cases, these shared libraries will be duplicated in both the microservices

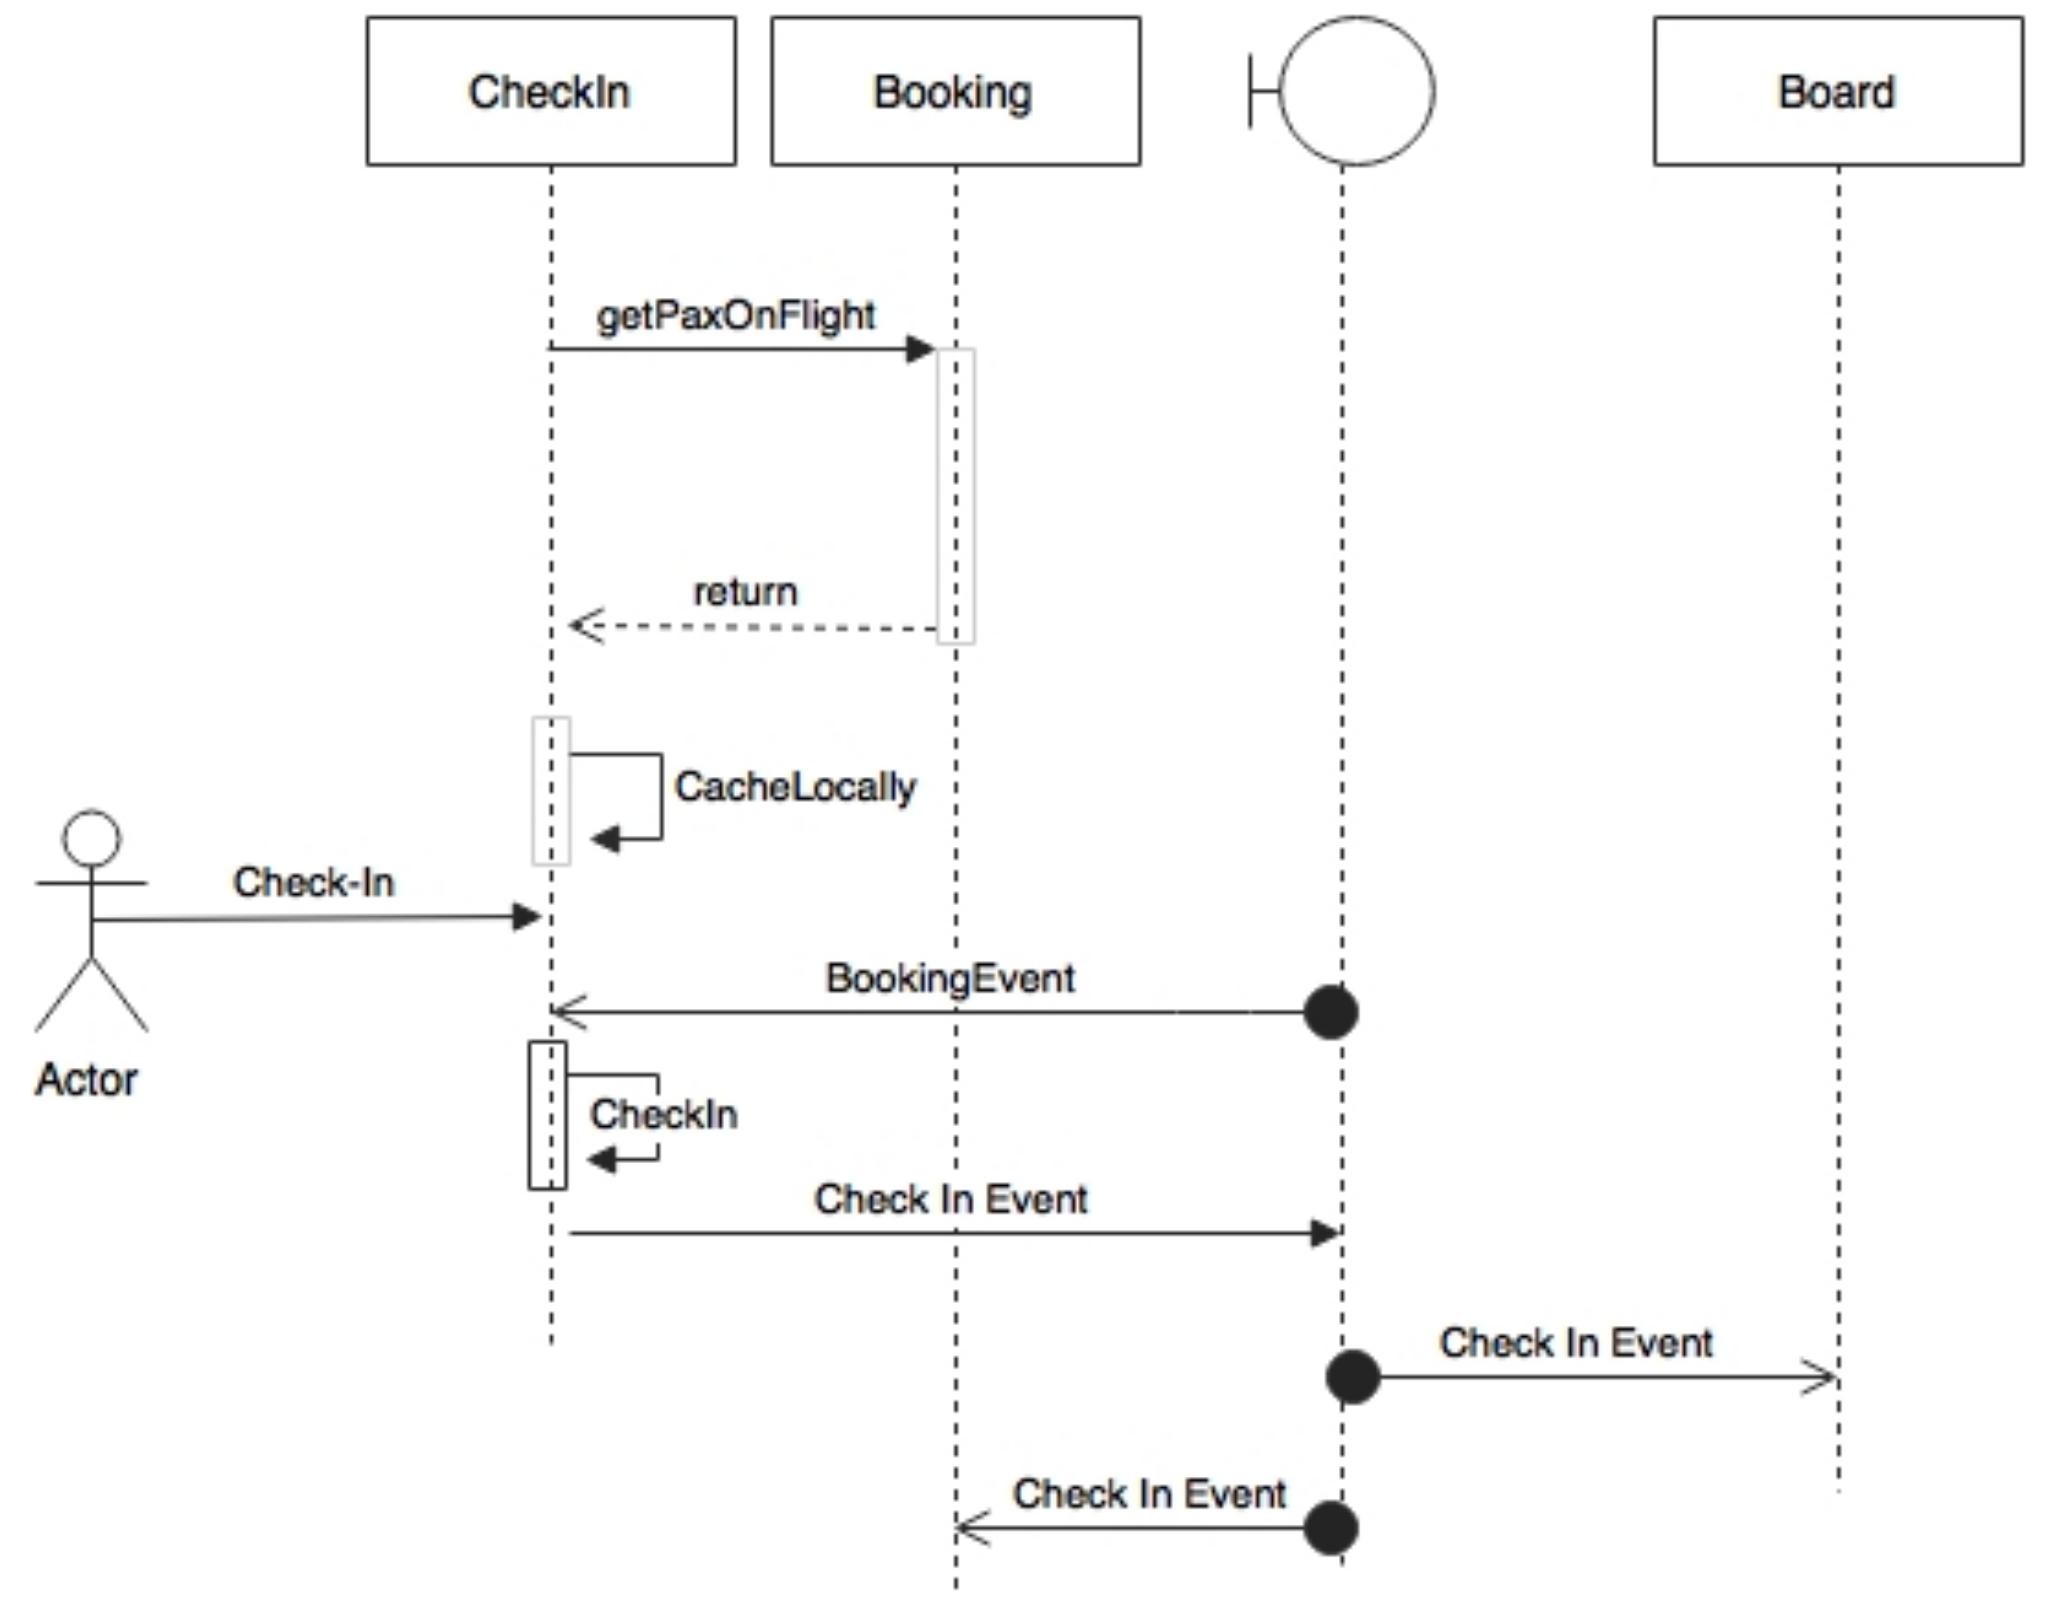
# Handling exceptions

- There is a synchronous communication between Booking and Fare
- In case the Fare service is not available, we trust the incoming request, and accept the Booking
- We will use a **circuit breaker** and a **fallback service** which simply creates the booking with a special status, and **queues** the booking for manual action or a system retry

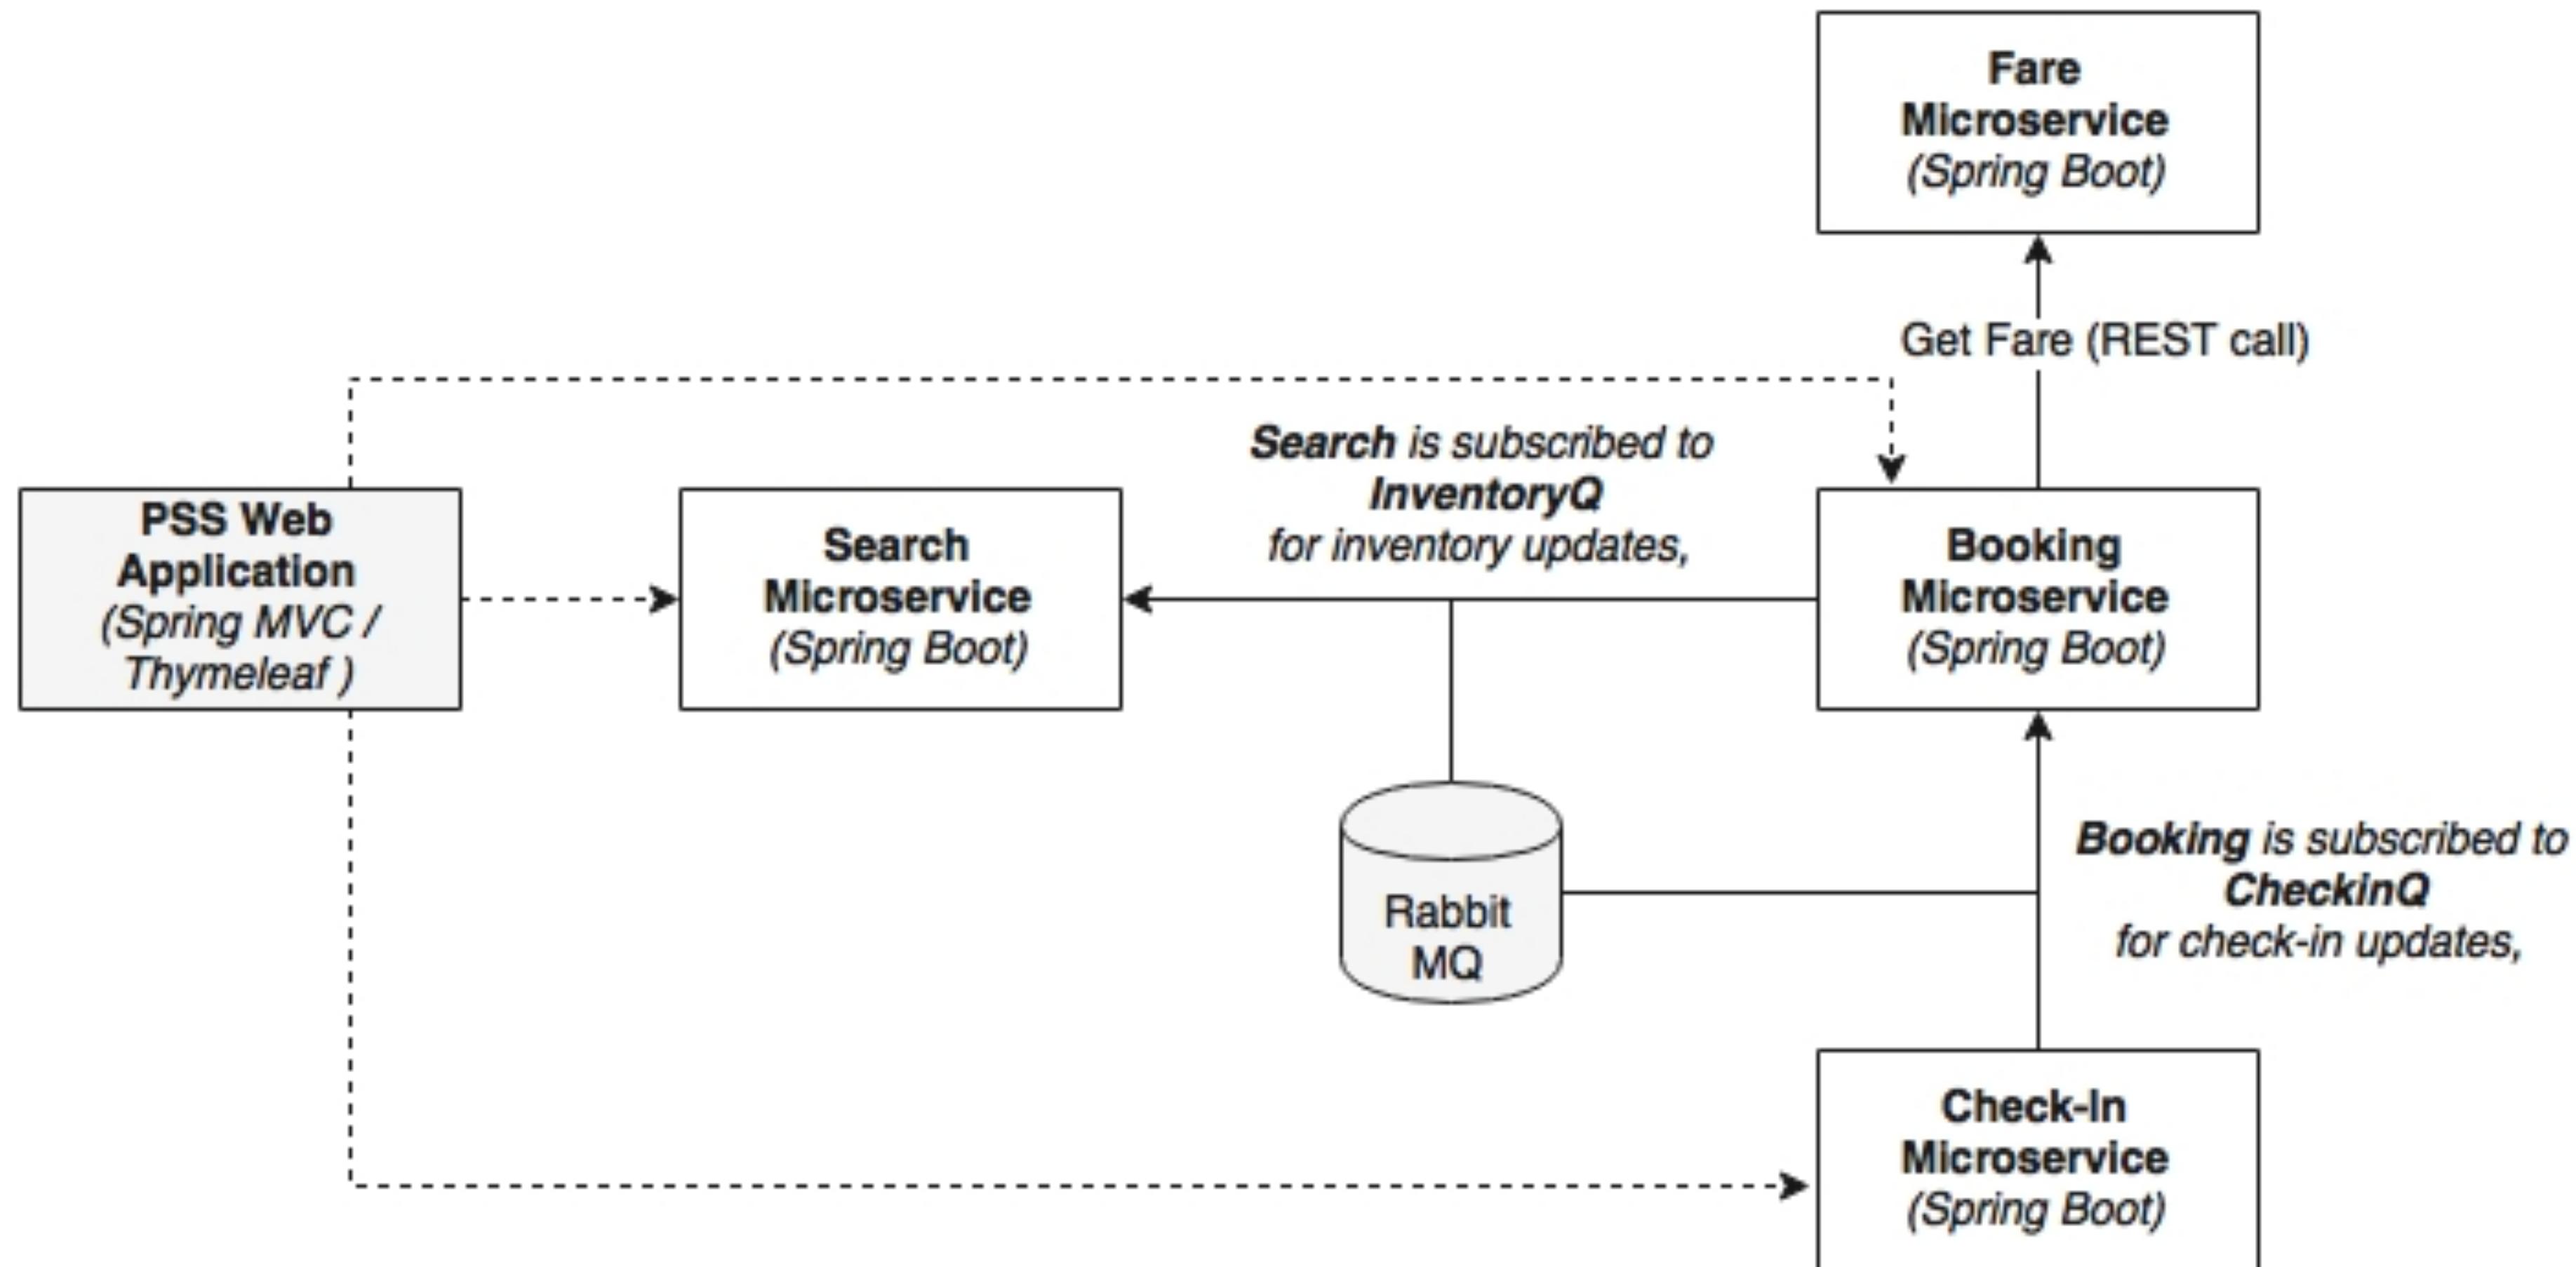


# Handling exceptions

- Consider a scenario where the Check-in services **fail immediately after** the Check-in Complete event is sent out
- The other consumers processed this event, but the actual check-in is rolled back.
  - This is because we are not using a two-phase commit
  - This could be done by catching the exception, and sending another Check-in Cancelled event



# Target implementation view



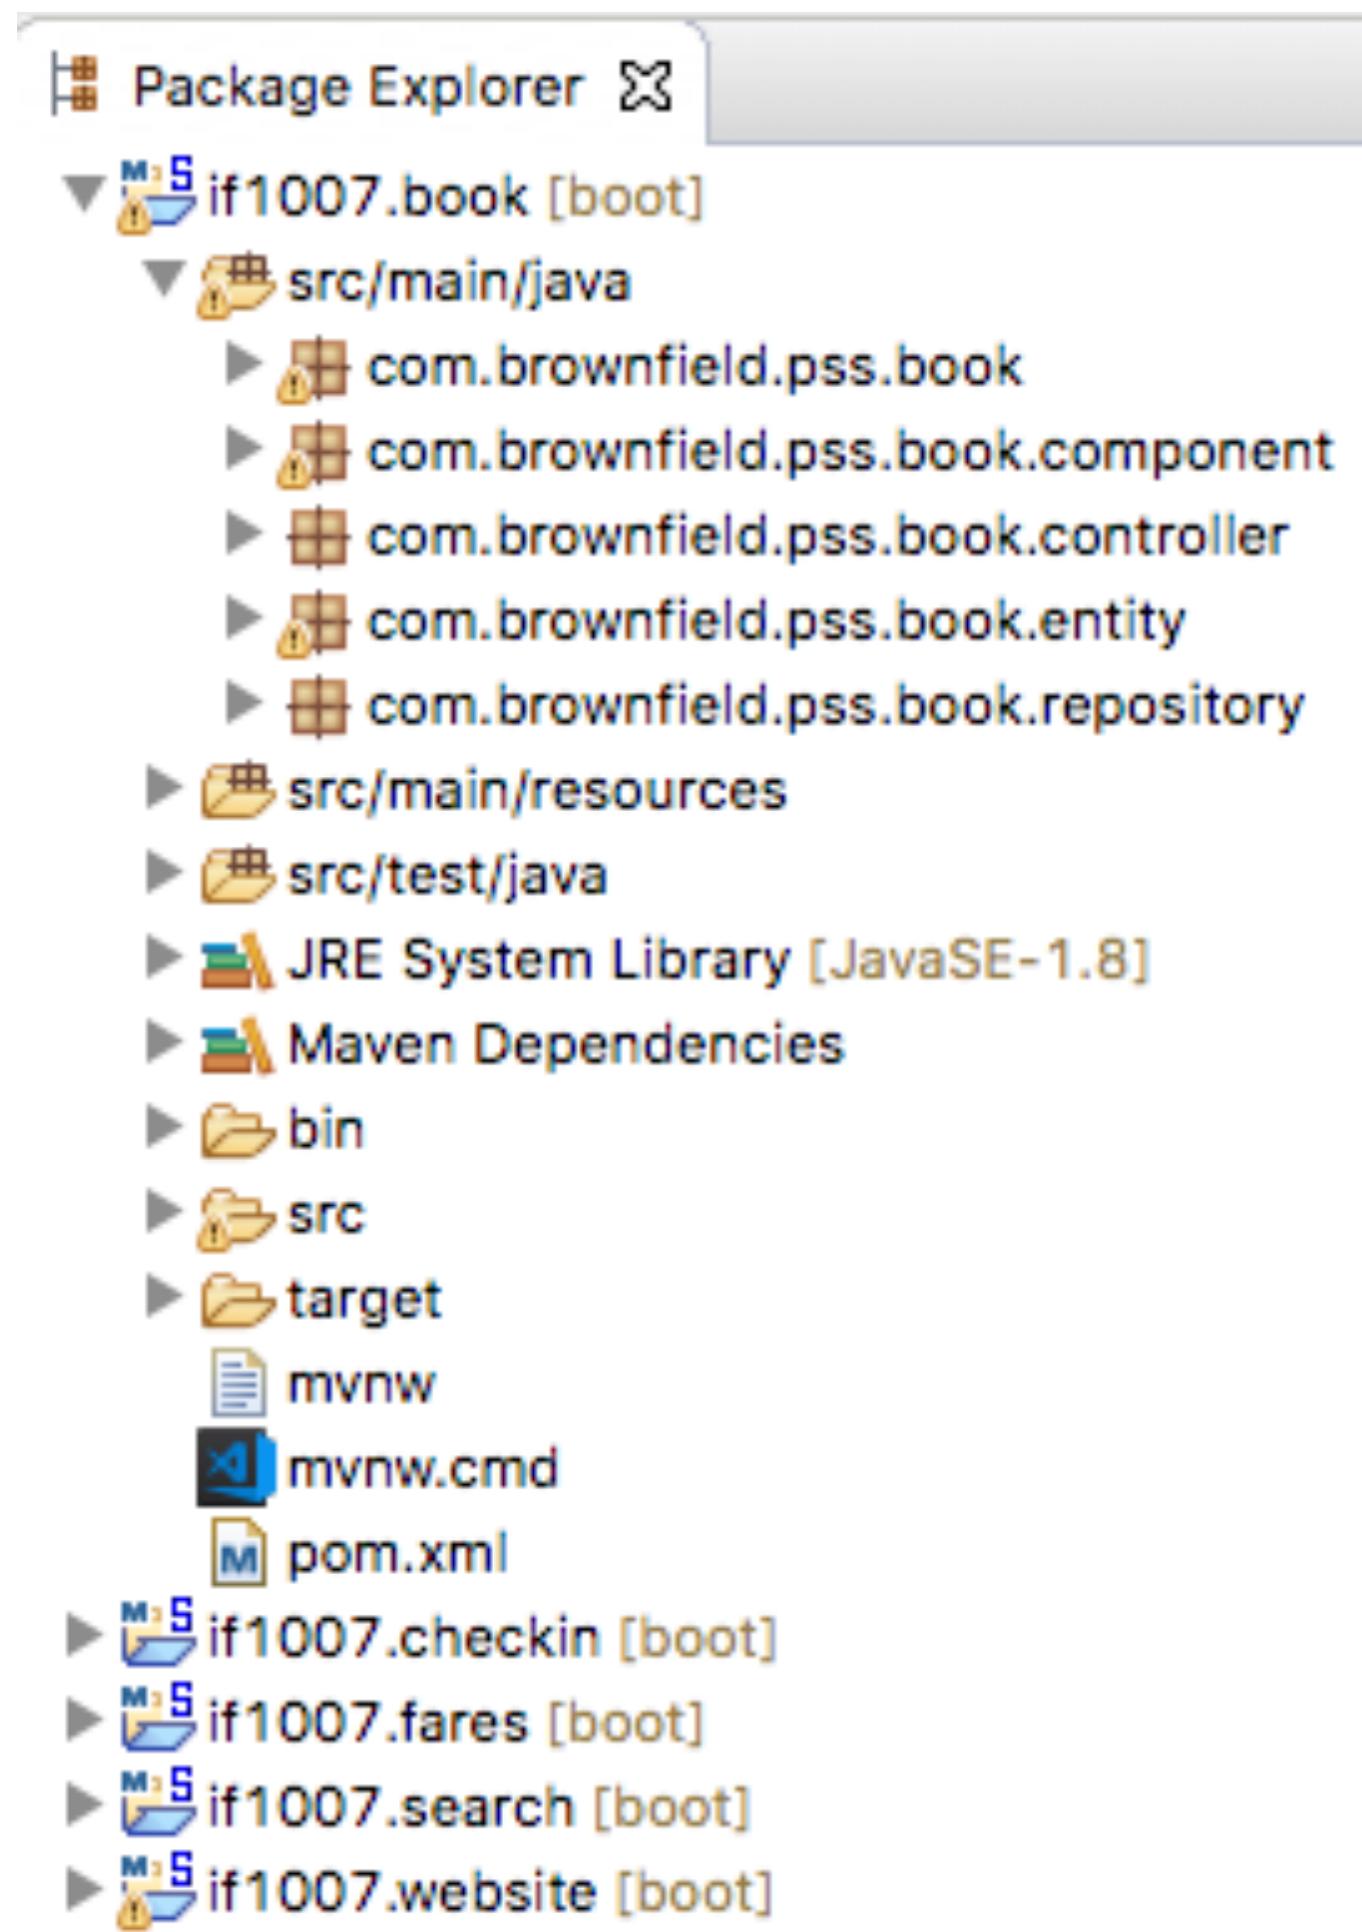
# Implementation projects

- The basic implementation of the BrownField Airline's PSS microservices system has five core projects as summarized in the following table

Microservice	Projects	Port Range
Book msvc	If1007.book	8060-8069
Check-in msvc	If1007.checkin	8070-8079
Fare msvc	If1007.fares	8080-8089
Search msvc	If1007.search	8090-8099
Website	If1007.website	8001

# Implementation projects

- The root folder (`com.brownfield.pss.book`) contains the default Spring Boot application.
- The `component` package hosts all the service components where the business logic is implemented.
- The `controller` package hosts the REST endpoints and the messaging endpoints. Controller classes internally utilize the component classes for execution.
- The `entity` package contains the JPA entity classes for mapping to the database tables.
- Repository classes are packaged inside the `repository` package, and are based on Spring Data JPA.



# Running and testing the project

Follow below steps to run the application.

1. Open terminal Window and move to project home.

```
mvn -Dmaven.test.skip=true install
```

2. Start Rabbit MQ.

```
rabbitmq-server
```

3. Run below commands from the respective project folders, in separate terminal windows.

```
java -jar if1007.fares/target/fares-1.0.jar  
java -jar if1007.search/target/search-1.0.jar  
java -jar if1007.checkin/target/checkin-1.0.jar  
java -jar if1007.book/target/book-1.0.jar  
java -jar if1007.website/target/website-1.0.jar
```

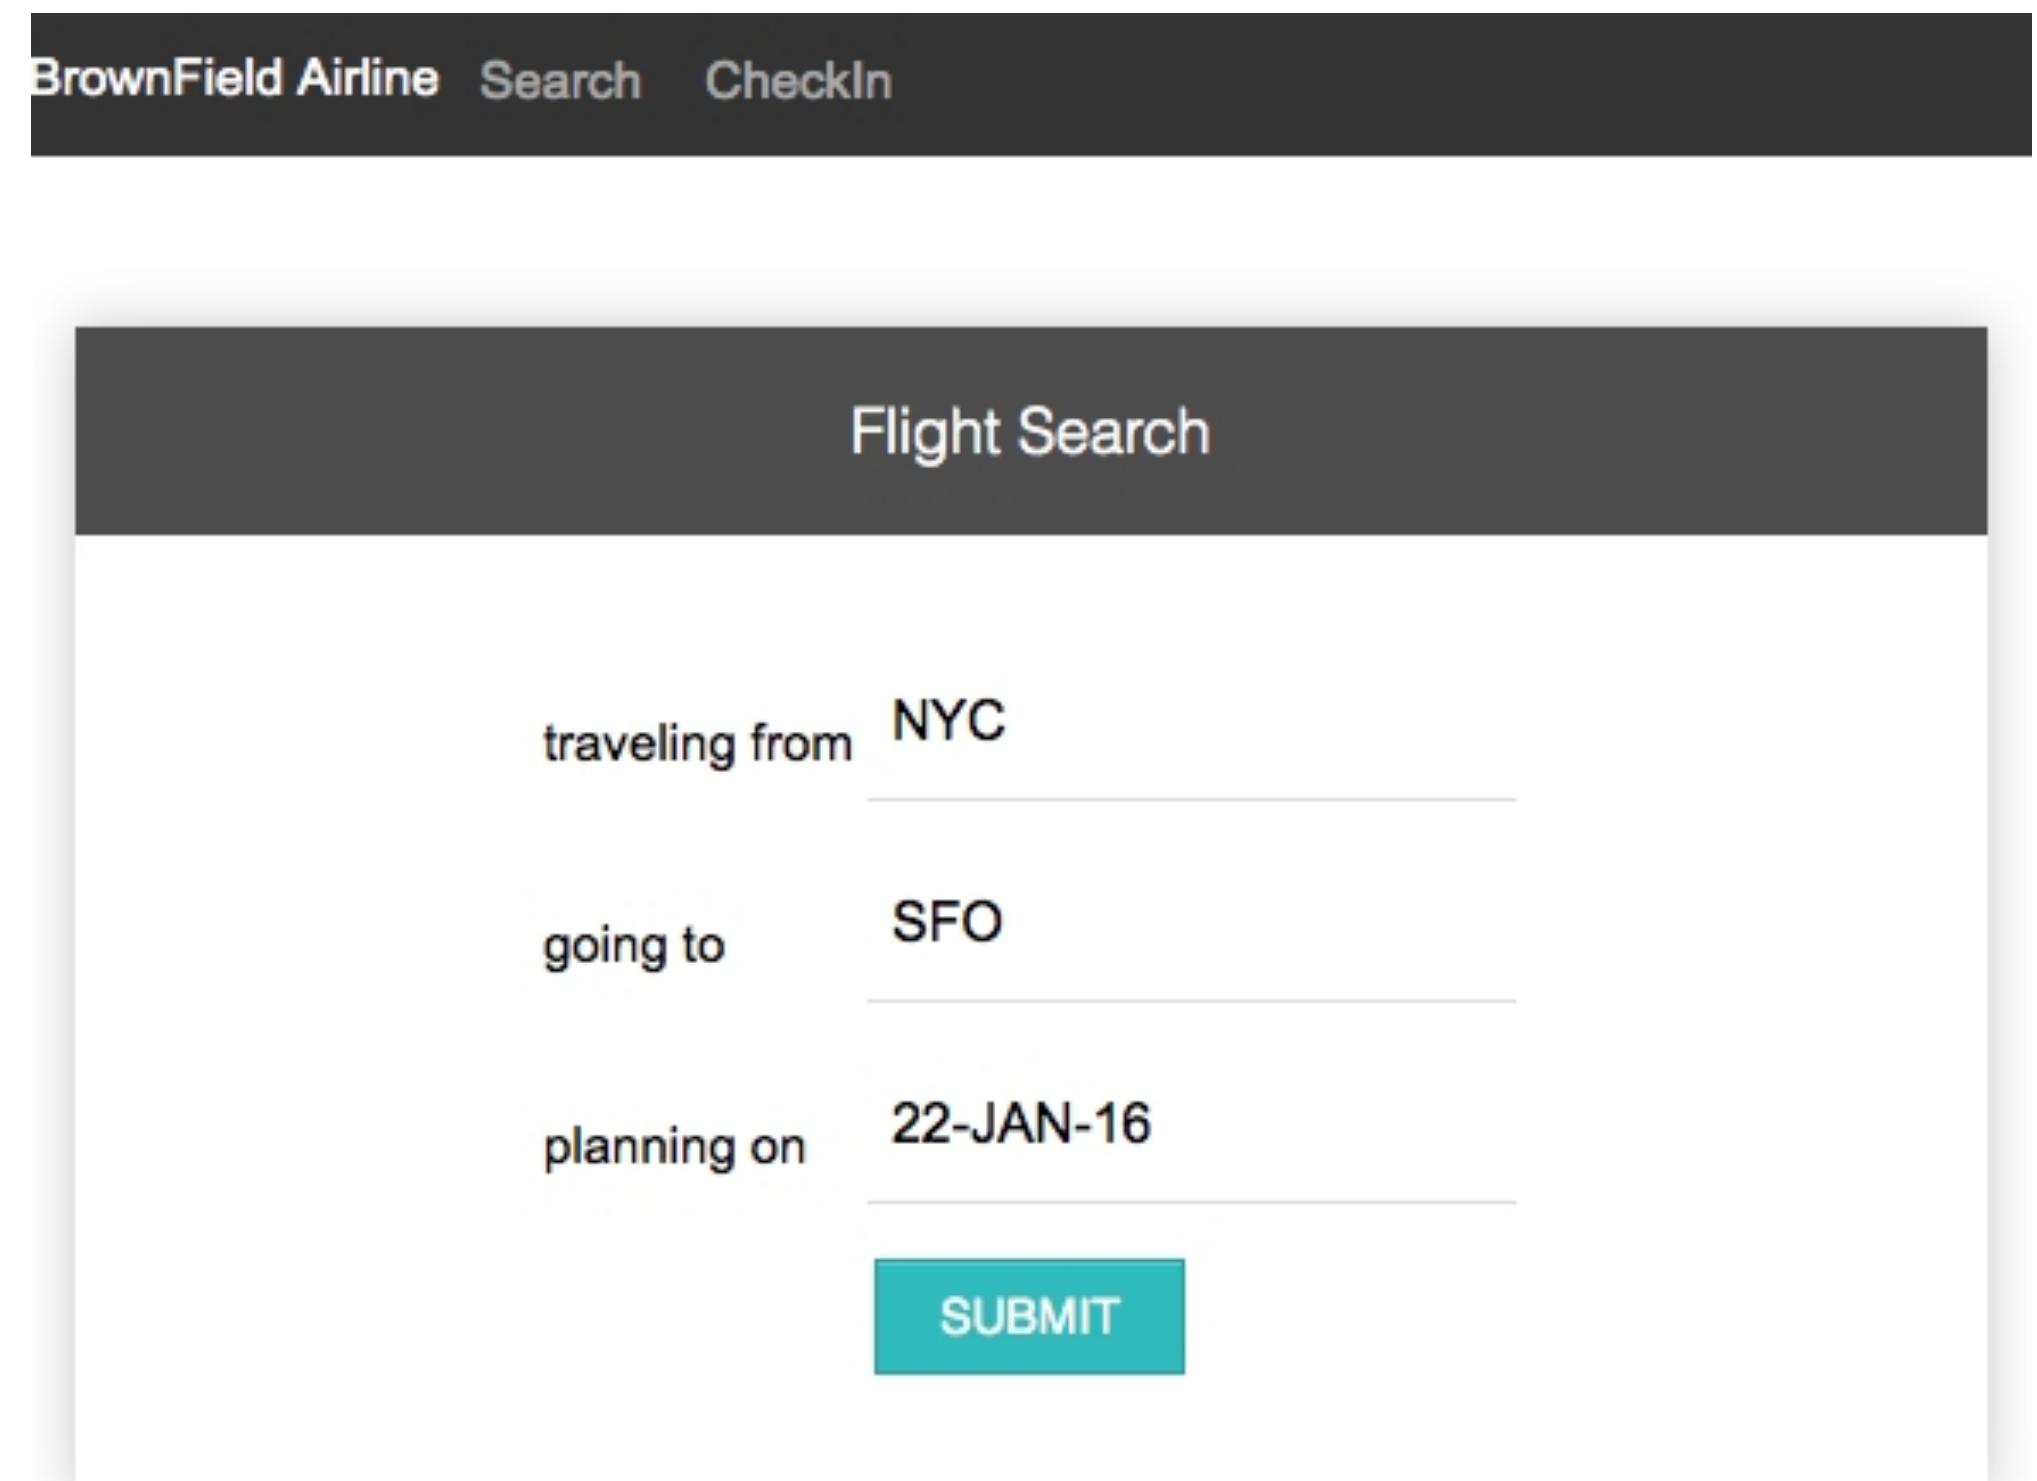
4. Open Browser Window and paste below URL.

```
http://localhost:8001
```

5. When asked for credentials use guest/guest123

6. Click Search Menu for search and Booking

7. Click CheckIn Menu for check-in



# Booking screen

BrownField Airline Search CheckIn

BrownField Airline Search CheckIn

Selected Flight

Available Flights

#	Flight	From	To	Date	Fare
2	BF101	NYC	SFO	22-JAN-16	101
3	BF105	NYC	SFO	22-JAN-16	105
4	BF106	NYC	SFO	22-JAN-16	106

[Book](#)

[Book](#)

[Book](#)

BF101 NYC SFO 22-JAN-16 101

First Name Rajesh

Last Name RV

Gender Male

CONFIRM

The image displays a booking interface for BrownField Airline. On the left, a list of available flights is shown with columns for flight number, route, date, and fare. The flight BF101 is selected and highlighted. On the right, a detailed view for flight BF101 shows fields for first name (Rajesh), last name (RV), and gender (Male). A 'CONFIRM' button is at the bottom.

# Booking screen

BrownField Airline [Search](#) [CheckIn](#)

Booking Confirmation

Your Booking is confirmed. Reference Number is 3

# Check-in microservice

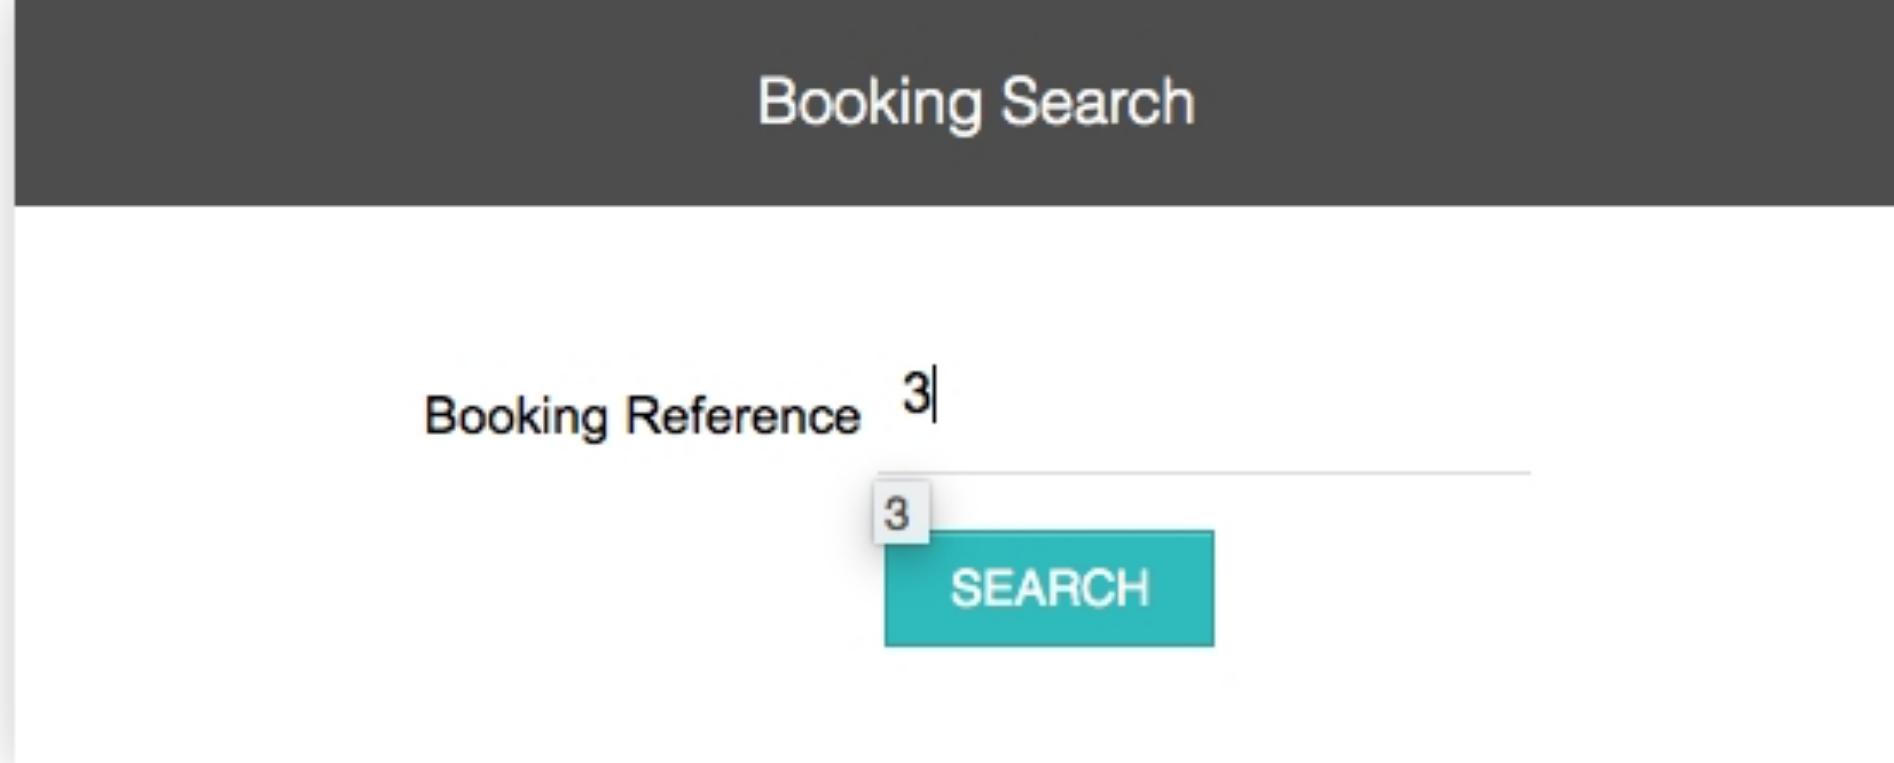
BrownField Airline Search CheckIn

Booking Search

Booking Reference 3

3

SEARCH



BrownField Airline Search CheckIn

Booking Search

Booking Reference 3

SEARCH

BF101 NYC SFO 22-JAN-16 Rajesh RV [CheckIn](#)

